

Marine Race Arena: A Configurable HoloOcean Benchmark for Underwater Gate Racing and Fleet Coordination

Andrea Bedei* [co-author]* [co-author / supervisor]*

*Department of Computer Science and Engineering, University of Bologna, Cesena, Italy

Abstract—Autonomous racing turns perception, planning, and control into a single measurable task under reproducible time pressure, and ground and aerial racing platforms have become effective evaluation substrates for this reason. Underwater robotics, however, lacks a comparable gate-racing benchmark: existing marine simulators provide vehicles and sensors but not a compact racing protocol that couples damped vehicle dynamics, range-limited and low-bandwidth sensing, currents, obstacles, an independent referee, and fleet coordination. This work presents Marine Race Arena, a configurable evaluation layer built around the HoloOcean simulator and a BlueROV2-class underwater vehicle. The central design principle is a strict separation between the autonomy stack and the evaluation: a controller observes only a documented official observation, while an independent referee uses privileged state to validate gate crossings, collisions, out-of-bounds and stuck conditions, timeouts, ranking, and team scoring. The benchmark is declarative—tracks, sensors, currents, obstacles, participants, and scoring are specified in a single configuration—and a custom controller is integrated by implementing three methods. We formalize the benchmark model, the gate-crossing and scoring rules, and the staggered fleet/team aggregation, and we report HoloOcean validation in which a deterministic rule-based baseline completes three official clean tracks. In clean, current-free conditions the baseline is collision-free; physical gate contact emerges under disturbance currents, which we report honestly rather than suppress. A stable two-rover staggered fleet run completes Horseshoe Bay with independent per-rover summaries and a single aggregated team score. The goal of Marine Race Arena is not to provide an optimal underwater racing controller, but a reproducible benchmark and evaluation contract on which such controllers—including learning-based ones—can be built and compared; current compensation, close-proximity fleet racing, and empirically calibrated rover-rover penalties are scoped explicitly as open extensions.

Index Terms—underwater robotics, autonomous racing, HoloOcean, BlueROV2, simulation benchmark, referee, multi-robot systems, fleet coordination

I. INTRODUCTION

Autonomous racing has become a productive research format because it compresses perception, planning, control, and real-time decision-making into a single quantitative task under repeatable constraints [1]. In ground racing, the F1TENTH platform standardizes circuits, a vehicle model, metrics, and baselines for continuous control and reinforcement learning [2], while full-scale leagues such as the Abu Dhabi Autonomous Racing League (A2RL) push planning and control to the limits of vehicle handling in head-to-head competition [3]. In aerial racing, gate-based benchmarks and compe-

titions drove a progression from modular perception-control stacks [4] to learning-based policies that reach champion-level performance [5]. A common thread is that a well-defined racing benchmark, with fixed geometry, an impartial referee, and machine-readable outcomes, accelerates and disciplines the comparison of autonomy methods.

Underwater robotics has mature simulation tools but lacks a comparable racing benchmark. Existing marine simulators emphasize intervention, inspection, navigation, or sensor realism rather than a compact racing protocol with explicit gates, currents, obstacles, independent scoring, and fleet evaluation. Yet underwater racing is not aerial or ground racing relocated to a different simulator. Vehicle dynamics are slow and strongly damped; global localization is unavailable and must be inferred from drifting inertial and Doppler sensing or from acoustic cues; optical perception is range-limited and degrades with turbidity; acoustic guidance is low-bandwidth and noisy; persistent currents act as a structured disturbance; and physical contact with a gate or another vehicle is ambiguous to detect. These properties change both what a controller can sense and how an evaluator must score behavior.

A benchmark for this domain must therefore separate the autonomy stack from the referee and from the simulator internals. The controller should receive only a documented *official observation*—onboard sensing, an acoustic target cue, and its own race progress—while the referee may use privileged ground-truth state to decide gate crossings, collisions, timeouts, and team scoring. Without this separation, it is impossible to claim that a result was obtained under realistic information, and it is easy to inadvertently leak global pose or exact gate geometry into the control loop.

Marine Race Arena is designed around this separation. It is neither a new hydrodynamic simulator nor a new autonomy algorithm. It is an evaluation layer built on HoloOcean [6] and a BlueROV2-class underwater vehicle [7], in which the race—world bounds, ordered gates, start and finish, currents, obstacles, participants, sensors, and referee—is described declaratively and a controller is plugged in by implementing a small interface. The framework supports three usage categories that are kept deliberately distinct so that infrastructure claims are not confused with solved-problem claims:

- **Official benchmark mode:** fixed official tracks, official gate apertures, official observation filtering that excludes ground-truth state, and reproducible referee scoring;

- **Diagnostic mode:** additional logs, participant-state traces, proximity diagnostics, and smoke tests used to validate the benchmark plumbing itself;
- **Experimental mode:** currents, obstacles, current-compensation controllers, and close-proximity fleet interaction, which are available for investigation but are *not* claimed as solved v0.1 results.

The contributions of this release-candidate paper are the following.

- 1) A declarative underwater gate-racing benchmark model: world, ordered gates with $1.5\text{ m} \times 1.5\text{ m}$ apertures, currents, obstacles, participants, sensors, and scoring are specified in a single JSON configuration, with three official clean tracks provided.
- 2) An official controller–observation contract and a HoloOcean race runner that keep ground-truth pose, simulator current vectors, and hidden gate geometry out of participant control, together with a controller interface that a custom policy satisfies by implementing three methods.
- 3) An independent referee that formalizes gate-crossing geometry, missed gates, collisions, out-of-bounds and stuck conditions, timeouts, ranking, and structured event logs.
- 4) A staggered fleet/team mode in which multiple BlueROV2 participants form a single team, maintain independent progress and referee state, and aggregate into one team-level score, with rover–rover proximity diagnostics.
- 5) A reproducible v0.1 validation package with a deterministic rule-based baseline, clean single-rover results on three official tracks, a stable staggered two-rover demonstration, and an explicit, honest scoping of current compensation and close-proximity fleet collisions as open work.

The remainder of the paper is organized as follows. Section II reviews underwater simulators, autonomous-racing benchmarks, and multi-vehicle underwater cooperation. Section III formalizes the benchmark model and the runtime architecture. Section IV describes the vehicle, the HoloOcean adapter, the sensor suite, the current model, and the official observation contract. Section V specifies the controller interface and the baseline controllers. Section VI formalizes gate validation, the referee state machine, and single-rover scoring. Section VII describes staggered fleet execution and team scoring. Section VIII reports the experimental evaluation. Section IX discusses limitations and future work, and Section X concludes.

II. RELATED WORK

We review related work along three axes: the underwater simulators that provide the physical and sensing substrate; the autonomous-racing benchmarks that define the evaluation format we adapt; and the multi-AUV cooperation and acoustic-communication literature that motivates and bounds the fleet

mode. Table I summarizes the positioning of Marine Race Arena against representative systems.

A. Underwater Robotics Simulators

Several simulators provide marine vehicles and sensors. The UUV Simulator offers a Gazebo-based package for underwater intervention and multi-robot scenarios [8]. Stonefish provides custom marine-robotics physics and a wide range of underwater sensors behind a ROS interface [9]. DAVE builds on Gazebo to support acoustic and optical sensing for autonomous underwater vehicles [10]. HoloOcean uses Unreal Engine to render underwater scenes and exposes common marine sensors—sonar, RGB camera, IMU, DVL, and depth—for one or more agents [6], and is the simulation backend we build on. These systems are complementary to Marine Race Arena: they simulate vehicles and sensors, whereas Marine Race Arena adds the race task, the official observation contract, the referee, the scoring, and the output protocol on top of one of them. None of them ships a gate-racing benchmark with an impartial referee, and consequently none enforces an information boundary between the controller and ground-truth state.

B. Autonomous Racing Benchmarks and Competitions

Racing benchmarks are valuable precisely because they expose end-to-end behavior under repeatable, adversarial constraints, as catalogued by the survey of Betz et al. [1]. F1TENTH provides a scaled autonomous car, standardized circuits, and an evaluation environment for continuous control and reinforcement learning [2]; it is the closest in spirit to Marine Race Arena, but it targets ground vehicles with planar LiDAR and does not separate referee logic as a first-class contract. At full scale, the A2RL challenge frames racing as multi-agent interaction at the limits of vehicle handling and reports the algorithms of the winning team in the inaugural Abu Dhabi event [3], while the A2RL league also runs a drone-racing format [12]. In aerial robotics, AlphaPilot established gate sequences as a demanding benchmark for fast perception and control [4], and Swift later demonstrated champion-level drone racing using deep reinforcement learning trained against a simulated opponent [5]. CARLA is not a racing benchmark, but it illustrates how an Unreal-based ecosystem can standardize maps, sensors, and tasks for driving research [11]. More broadly, standardized environment interfaces such as OpenAI Gym [13] and GPU-parallel robot-learning simulators such as Isaac Gym [14] show how a documented observation–action contract enables systematic comparison of learning-based controllers. Marine Race Arena borrows the gate-racing idea and the documented observation contract from these efforts, but adapts the benchmark to underwater sensing, slow and damped vehicle response, acoustic guidance, current disturbances, and a referee whose scoring is computed entirely outside the control loop.

C. Multi-AUV Cooperation and Underwater Communication

Fleet operation underwater is constrained by the physics of the acoustic channel and by uncertain localization. Underwater

TABLE I

POSITIONING OF MARINE RACE ARENA RELATIVE TO REPRESENTATIVE ROBOTICS SIMULATORS AND RACING BENCHMARKS. “RACING TASK” DENOTES A BUILT-IN GATE/LAP PROTOCOL WITH TIMING; “INDEP. REFEREE” DENOTES SCORING LOGIC THAT IS SEPARATED FROM THE CONTROLLER AND USES PRIVILEGED STATE; “OFFICIAL OBS. CONTRACT” DENOTES AN ENFORCED FILTER THAT EXCLUDES GROUND-TRUTH POSE FROM THE CONTROL INPUT.

System	Domain	Physics / sensing	Racing task	Multi-robot	Indep. referee	Official obs. contract
UUV Simulator [8]	Underwater	Gazebo, hydrodynamics, multi-robot	–	✓	–	–
Stonefish [9]	Underwater	Custom physics, marine sensors, ROS	–	–	–	–
DAVE [10]	Underwater	Gazebo, acoustic/optical sensing	–	✓	–	–
HoloOcean [6]	Underwater	Unreal Engine, sonar/camera/DVL/IMU	–	✓	–	–
F1TENTH [2]	Ground	Scaled car, LiDAR, RL/control	✓	partial	partial	✓
AlphaPilot / Swift [4], [5]	Aerial	Quadrotor, vision-based gate racing	✓	–	✓	✓
CARLA [11]	Ground	Unreal Engine, urban sensing, tasks	partial	–	partial	✓
Marine Race Arena (this work)	Underwater	HoloOcean, BlueROV2, acoustic+camera	✓	✓	✓	✓

“–” not provided as a first-class feature; “partial” available but not the primary design contract. Marine Race Arena is the only entry that combines an underwater simulation backend with a gate-racing protocol, an independent referee, multi-vehicle/team scoring, and an enforced official-observation contract.

acoustic links suffer from low bandwidth, long propagation delay, multipath, and time-varying attenuation [15], so the assumption of continuous high-bandwidth, low-latency inter-vehicle communication that underpins many terrestrial multi-robot methods does not transfer. Surveys of AUV formation control document how these communication and navigation constraints shape the achievable coordination: Yang et al. analyze formation performance, control, and communication capability jointly [16], and Yan et al. review formation-control methods and their dependence on relative-state estimation [17]. Consistent with this literature, Marine Race Arena does not assume shared ground-truth state or high-rate communication between vehicles. In v0.1 the fleet mode is deliberately conservative: it validates multiple HoloOcean agents, staggered releases, independent per-rover progress, and team-level scoring, while keeping close-proximity arbitration and empirically calibrated rover-rover penalties in diagnostic or experimental status. This positions the fleet contribution as an evaluation substrate for future cooperative strategies rather than a claim that close-proximity underwater racing is solved.

D. Positioning

In contrast to the systems above, Marine Race Arena is the only one that simultaneously provides an underwater simulation backend, a gate-racing protocol with timing, an independent referee that uses privileged state, multi-vehicle and team-level scoring, and an enforced official-observation contract that excludes ground-truth pose from control (Table I). It does not compete with the underwater simulators on physical fidelity, nor with the racing platforms on autonomy performance; it complements both by supplying the missing evaluation contract for underwater gate racing.

III. BENCHMARK MODEL AND SYSTEM ARCHITECTURE

We specify Marine Race Arena as an *evaluation problem* rather than a controller-design problem: the benchmark fixes the world, the task, and the scoring, and a participant supplies only a policy. This section defines the benchmark instance, the participant interface, and the runtime architecture that enforces the separation between the two.

A. Benchmark Instance

Definition 1 (Benchmark instance). A Marine Race Arena instance is the tuple

$$\mathcal{E} = (\mathcal{W}, \mathcal{G}, \mathcal{S}, \mathcal{F}, \mathcal{C}, \mathcal{O}, \mathcal{P}, \mathcal{R}), \quad (1)$$

where:

- $\mathcal{W}$ is the *world*: an axis-aligned bounds box $\mathcal{B} = [x_{\min}, x_{\max}] \times [y_{\min}, y_{\max}] \times [z_{\min}, z_{\max}] \subset \mathbb{R}^3$ (with z positive-up, so the navigable volume satisfies $z < 0$) and the HoloOcean environment in which it is rendered;
- $\mathcal{G} = (g_1, \dots, g_M)$ is the *ordered* gate sequence; each gate is the tuple $g_k = (c_k, n_k, u_k, v_k, w_k, h_k)$ of center $c_k \in \mathbb{R}^3$, unit normal n_k (the passage direction), in-plane right and up unit axes u_k, v_k , and inner aperture width w_k and height h_k ;
- $\mathcal{S}$ is the start configuration: a spawn pose and an optional release delay $\tau_i \geq 0$ per participant;
- $\mathcal{F}$ is the finish condition: completion of the last gate of the last lap (Section VI);
- $\mathcal{C}$ is the current field $\mathbf{v}_c : \mathbb{R}^3 \times \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}^3$ (Section IV);
- $\mathcal{O}$ is the obstacle set;
- $\mathcal{P} = \{1, \dots, N\}$ is the participant set, each with a vehicle, sensor profile, controller, spawn, and release delay;
- $\mathcal{R}$ is the referee configuration: gate-validation parameters, penalty values, and scoring rules.

The gate basis is right-handed and derived from the passage direction alone, so the orientation of a gate is unambiguous and independent of any visual rotation (Section VI).

B. Participant Interface

Definition 2 (Policy). A participant i implements a policy

$$\pi_i : o_{i,t} \mapsto u_{i,t}, \quad (2)$$

mapping the official observation $o_{i,t}$ (Section IV) to a normalized, high-level, body-frame command

$$u_{i,t} = [u^{\text{surge}}, u^{\text{sway}}, u^{\text{heave}}, u^{\text{yaw}}]^\top \in [-1, 1]^4. \quad (3)$$

Time is discretized at a fixed control timestep Δt , with $t_k = k \Delta t$. The runner may clamp $u_{i,t}$ for actuator and

TABLE II
DESIGN REQUIREMENTS FOR THE V0.1 BENCHMARK LAYER.

ID	Requirement
R1	Race configuration is declarative and reproducible from a single set of track JSON files.
R2	Official controller observations exclude ground-truth pose and simulator-only state.
R3	Referee logic is independent of the participant controller and cannot be influenced by it.
R4	Gate dimensions, referee margins, and track geometry are not changed to hide failures.
R5	Every run emits structured event logs and machine-readable summaries.
R6	Single-rover behavior is unchanged when fleet mode is disabled.
R7	Fleet mode aggregates a team score while preserving per-rover diagnostics.

depth safety, but in official mode the policy must not read referee internals, ground-truth pose, simulator current vectors, or hidden gate geometry. This restriction is what makes a reported result attributable to realistic onboard information.

C. Design Requirements

The release-candidate design follows the seven requirements in Table II. They encode the central commitment of the benchmark—that scoring is independent of, and not improvable by, the participant—and the secondary commitment that enabling fleet mode must not silently change single-rover behavior.

D. Runtime Architecture

Fig. 1 shows the runtime architecture and the modules that realize it. The configuration loader (`config/loader.py`) parses and validates the track JSON into a typed configuration; the arena builder (`arena/arena_builder.py`) instantiates gates, beacons, the current field, obstacles, and bounds. The race runner (`scripts/run_marine_race.py`) drives the discrete-time loop. On the simulation side, the adapter (`adapters/holocean_adapter.py`) is responsible for reset, actions, ticking, and sensor extraction; a kinematic fallback adapter implements the same interface for plumbing tests without the engine. On the control side, the runner assembles the official observation and passes it to the controller, which returns a command. On the scoring side, the referee (`referee/referee.py`) consumes privileged simulator state, validates progress, and emits structured events and summaries (`referee/logger.py`). The dashed boundary in Fig. 1 marks the information contract: nodes inside it see only the official observation, while ground-truth state is confined to the referee. This separation lets a privileged oracle controller exist in the repository for feasibility checks without ever becoming part of the official benchmark.

IV. VEHICLE, SIMULATOR, AND SENSING

This section describes the simulation backend, the vehicle and its sensors, the acoustic and current models, the declarative track configuration, and the official observation contract that the referee enforces.

A. Simulator and Vehicle

The reference adapter uses HoloOcean [6], an Unreal-Engine underwater simulator, with a BlueROV2-class agent [7] in direct eight-thruster mode (HoloOcean `control_scheme` 0). The simulator advances at 30 physics ticks per second; the benchmark runs a control loop at timestep $\Delta t = 0.033$ s, so each control step corresponds to one physics tick and the effective control rate is ≈ 30 Hz. The adapter is intentionally replaceable: a pure-kinematic fallback adapter implements the same interface for unit and plumbing tests, integrating a damped point-vehicle model (maximum linear speed 1.25 m/s, maximum yaw rate $65^\circ/\text{s}$) and adding the current as a world-frame velocity, so that referee, scoring, and fleet logic can be tested deterministically without the engine.

A high-level command $u = [u^{\text{surge}}, u^{\text{sway}}, u^{\text{heave}}, u^{\text{yaw}}]^\top$ is mapped to the eight BlueROV2 thrusters by a fixed mixing matrix. The four vertical thrusters receive u^{heave} ; the four horizontal thrusters receive

$$\tau_{1-4} = s_t \begin{bmatrix} u^{\text{surge}} + u^{\text{sway}} + \kappa u^{\text{yaw}} \\ u^{\text{surge}} - u^{\text{sway}} - \kappa u^{\text{yaw}} \\ u^{\text{surge}} + u^{\text{sway}} - \kappa u^{\text{yaw}} \\ u^{\text{surge}} - u^{\text{sway}} + \kappa u^{\text{yaw}} \end{bmatrix}, \quad (4)$$

with yaw-mixing gain $\kappa = 0.35$ and thruster scale s_t , each component clamped to the thruster limit. A controller may instead return the eight thruster values directly; the high-level interface is the default because it is the one the official baselines and the documented contract use.

B. Sensor Suite

Table III lists the onboard sensors. The official profile combines a forward RGB camera, a depth sensor, an IMU, a DVL and a world-velocity sensor, a contact sensor, and a synthesized acoustic beacon. Ground-truth pose, location, rotation, and dynamics sensors exist in the simulator and are used internally by the referee, but are removed from the controller-visible observation in official mode (Section IV-F).

C. Acoustic Beacon Model

Each gate carries one beacon, placed at $b_k = c_k + \delta_n n_k + \delta_r u_k + \delta_u v_k$ relative to the gate center (default offset $(0, 0, 0.35)$ m along the gate-up axis). For a receiver at position r with heading ψ , let $\delta = b_k - r$ and $\rho = \|\delta\|$. The beacon reports range, a yaw-relative bearing, an elevation, and a signal strength:

$$\rho = \|\delta\|, \quad \beta = \text{wrap}(\text{atan2}(\delta_y, \delta_x) - \psi), \quad (5)$$

$$\varepsilon = \text{atan2}(\delta_z, \sqrt{\delta_x^2 + \delta_y^2}), \quad \text{SS} = \max\left(0, 1 - \frac{\rho}{\rho_{\max}}\right),$$




Fig. 1. Runtime architecture of Marine Race Arena. Declarative configuration (top) is loaded and validated into an arena; the race runner drives the control loop, exchanging state and actions with the simulator adapter and reporting progress to the independent referee, which emits structured outputs. The dashed region is the official-observation boundary: the controller observes only the filtered observation and returns a command, while ground-truth pose and gate geometry remain on the referee side. Single-rover and fleet runs share the same pipeline.

TABLE III

ONBOARD SENSOR SUITE OF THE BLUEROV2-CLASS VEHICLE. RATES ARE THE CONFIGURED HOLOOCEAN UPDATE RATES; THE ACOUSTIC BEACON IS SYNTHESIZED BY THE ARENA RATHER THAN RENDERED BY THE ENGINE. GROUND-TRUTH POSE SENSORS USED INTERNALLY BY THE REFEREE ARE EXCLUDED FROM THE OFFICIAL OBSERVATION.

Sensor	Rate	Output and configuration
Depth	30 Hz	Scalar depth (m); noise $\sigma = 0$.
IMU	30 Hz	Linear acceleration and angular rate, with bias terms.
DVL	15 Hz	Body-frame velocity (m/s); beam elevation 22.5° , max range 50 m.
Front camera	30 Hz	RGB image, 640×480 , 90° field of view.
Velocity	30 Hz	World-frame linear velocity (m/s).
Collision	30 Hz	Boolean contact flag.
Acoustic beacon	10 Hz	Range, bearing, and elevation to the target-gate beacon; range 50–130 m, noise std. 0–0.6 depending on track.

Derived fields: heading ψ (deg) and depth (m) from the pose used by the referee; the environment current vector is exposed as telemetry but is excluded from official no-cheat control.

where $\rho_{\max}$ is the beacon range and $\text{wrap}(\cdot)$ maps to $(-180^\circ, 180^\circ]$. A reading is invalid beyond $\rho_{\max}$ or when a Bernoulli dropout fires. In the noisy acoustic mode, independent Gaussian noise $\mathcal{N}(0, \sigma_b)$ is added to β , ε , and ρ ; the official tracks use $\sigma_b \in \{0.2, 0.6\}$ and dropout probabilities up to 0.04. Only the beacon of the current target gate is active in the default *active-when-target* mode, so the beacon provides a target cue but not a map of the whole course.

D. Current Field Model

The current field is a sum of analytic primitives evaluated at the vehicle position p and time t ,

$$\mathbf{v}_c(p, t) = \sum_j \mathbf{v}_j(p, t), \quad (6)$$

where each primitive is one of the following. A *constant* term contributes a fixed velocity $\mathbf{v}_0$. A *localized jet* centered at c_j with radius R_j contributes, for $\|p - c_j\| \leq R_j$,

$$\mathbf{v}_j = \mathbf{v}_0 w(d), \quad w(d) = \begin{cases} \max(0, 1 - \frac{d}{R_j}) & \text{linear,} \\ \exp(-\frac{d^2}{2\sigma^2}) & \text{Gaussian,} \end{cases} \quad (7)$$

with $d = \|p - c_j\|$ and $\sigma = \max(R_j/2, 10^{-6})$, and zero outside R_j . A *sinusoidal* term modulates one axis, $v_{j,\text{axis}}(t) = v_{0,\text{axis}} + A \sin(2\pi f t + \phi)$. A *vortex* of radius R_j , tangential speed T , and vertical speed V_z contributes, with $r = \sqrt{\Delta x^2 + \Delta y^2}$ and tangent $\hat{t} = (-\Delta y, \Delta x)/r$,

$$\mathbf{v}_j = (\hat{t}_x T w(r), \hat{t}_y T w(r), V_z w(r)). \quad (8)$$

The provided strong profile superposes a constant set-flow of $[0.75, 1.05, 0]$ m/s, two Gaussian jets, a vortex, and a vertical sinusoid; the medium profile scales these magnitudes by 0.5. The current physically perturbs the vehicle when the HoloOcean build exposes ocean-current coupling; otherwise the field is exposed as telemetry only, and the fallback adapter always couples it kinematically.

E. Declarative Track Configuration

A track is a self-contained JSON file. Its principal sections are `race` (name, benchmark task, timing mode, laps, duration), `world` (bounds and environment), `start/finish`, `track/gates` (ordered sequence, centers, passage directions, aperture sizes), `beacon`, `currents/current_profiles`, `obstacle_generation/obstacles`, `participants` (vehicle, spawn, sensors, controller, optional `start_delay_s`), and `referee` (gate validation, penalties, scoring). A named `benchmark_task—clean_gate, obstacle_gate, current_gate`, or `multi_rov`—selects a preset overlay. Listing 1 shows a representative excerpt.

Listing 1. Representative excerpt of a track configuration.

```
"track": { "gate_inner_size_m": [1.5, 1.5],
  "gate_sequence": ["G01", "G02", "..."] },
"gate": { "id": "G01",
  "center": [4.0, -2.0, -3.5],
  "passage_direction": [1.0, 0.0, 0.0] },
"beacon": { "mode": "active_when_target",
  "range_m": 90.0, "noise_std": 0.2,
  "update_rate_hz": 10.0 },
"participants": [ { "id": "bluerov2_01",
  "vehicle": "BlueROV2",
  "controller": "rule_gate_baseline",
  "official_sensor_profile": true,
  "start_delay_s": 0.0 } ],
"referee": { "penalties": { "minor_collision_s": 5.0,
  "missed_gate_dnf": true },
  "gate_validation": {
    "vehicle_clearance_margin_m": 0.1 } }
```

The three official v0.1 tracks are summarized in Table IV; all use $1.5\text{ m} \times 1.5\text{ m}$ apertures that are never enlarged for validation (requirement R4).

F. Official Observation Contract

The controller `step` function receives a dictionary with the fields in Table V. In official mode the runner removes pose, location, rotation, dynamics, exact-position, debug ground-truth, and exact-gate-geometry fields, then keeps only the configured allow-list of sensor names; numeric arrays are made JSON-safe while image arrays are preserved. The environment current vector may appear in diagnostic telemetry, but using it as a control input voids official no-cheat claims, since it constitutes privileged knowledge of the disturbance. The intended information contract is therefore acoustic target guidance, onboard sensing, and race progress—not global state.

V. CONTROLLER INTERFACE AND BASELINES

A central usability goal of Marine Race Arena is that a custom autonomy stack is integrated by implementing a small interface, without touching the runner, the referee, or the adapter. This section specifies that interface, the dynamic loading mechanism, the action mapping, and the reference baseline whose results we report.

A. Controller Interface

A controller implements three methods: `reset(race_info)`, called once before the run with static race metadata (timing mode, laps, official mode, bounds, and the initial target-gate id); `step(observation)`, called every control tick with the official observation of (2) and returning the command of (3); and `close()`, called at teardown. The contract is duck-typed: any object exposing these three callables is accepted, so a participant need not depend on framework base classes. A manual controller may raise a dedicated stop exception from `step` to end a run cleanly, and any other exception is caught, recorded, and treated as a zero command so that one faulty controller cannot crash an evaluation.

Listing 2. Minimal custom controller.

```
class MyController:
  def reset(self, race_info):
    self.target = race_info.get("initial_target_gate_id")
  def step(self, observation):
    # ... map beacon/sensors to a command ...
    return {"surge": 0.3, "sway": 0.0,
      "heave": 0.0, "yaw": 0.0}
  def close(self):
    pass
```

```
def step(self, observation):
  beacon = observation.get("beacon", {})
  # ... map beacon/sensors to a command ...
  return {"surge": 0.3, "sway": 0.0,
    "heave": 0.0, "yaw": 0.0}

def close(self):
  pass
```

The high-level command fields are listed in Table VI; values are normalized to $[-1, 1]$ and clamped by the framework before the action mapping.

B. Dynamic Loading

A controller is selected at run time in one of three ways: a built-in alias (e.g. `rule_gate_baseline`); a dotted module path with an explicit class; or a file path together with `--controller-class`, which is imported dynamically. This lets an external policy be evaluated without modifying the package:

```
python -m marine_race_arena.scripts.run_marine_race \
--track ../marine_race_horseshoe_bay.json \
--benchmark-task clean_gate \
--participant-controller path/to/my_controller.py \
--controller-class MyController \
--adapter holoocean --official --headless
```

C. Action Mapping

The framework clamps the returned command to $[-1, 1]^4$ and applies a depth-safety rule that suppresses heave that would drive the vehicle past $z_{\min}$ or $z_{\max}$, then maps it to thrusters via (4). A controller may alternatively return raw thruster values; the high-level interface is the documented default.

D. Reference Baseline: `rule_gate_baseline`

The reference controller is deliberately interpretable rather than optimal. Its role is to be a deterministic policy that, using only official observations, completes the clean tracks and exposes failures under currents, obstacles, or dense fleet interaction. Let β and ρ be the beacon bearing and range, z the depth, and z^* the target depth implied by the gate. The controller combines four mechanisms.

Align-before-surge. Yaw tracks the beacon bearing through a clamped proportional law with a small deadband,

$$u^{\text{yaw}} = \text{clamp}(k_\psi \beta, -\bar{u}_{\text{yaw}}, \bar{u}_{\text{yaw}}), \quad \bar{u}_{\text{yaw}} = 0.16, \quad (9)$$

and surge is gated by alignment: the vehicle turns in place when $|\beta| \geq 34^\circ$, surges at a reduced rate for $24^\circ \leq |\beta| < 34^\circ$, and otherwise drives forward up to a conservative cap $\bar{u}_{\text{surge}} = 0.46$ that scales with range and with $\cos \beta$.

Depth hold. A PI law on depth error $e_z = z^* - z$ produces a depth-keeping heave,

$$u_{\text{depth}}^{\text{heave}} = \text{clamp}\left(k_p e_z + k_i \int e_z dt\right), \quad (k_p, k_i) = (0.20, 0.08), \quad (10)$$

which is blended with the beacon-elevation heave as $u^{\text{heave}} = 0.70 u_{\text{beacon}}^{\text{heave}} + 0.30 u_{\text{depth}}^{\text{heave}}$.

TABLE IV
OFFICIAL V0.1 TRACK CONFIGURATIONS. GATE APERTURES ARE DECLARED IN THE TRACK FILES AND ARE NOT ENLARGED FOR VALIDATION.

Track file	Track name	Gates	Laps	Total gates	Length (m)	Intended use
marine_race_horseshoe_bay.json	Marine Race Horseshoe Bay	12	1	12	93.8	Clean-gate horseshoe baseline
marine_race_vertical_serpent.json	Marine Race Vertical Serpent	17	1	17	118.3	Vertical/slalom gate sequencing
marine_race_mixed_endurance.json	Marine Race Mixed Endurance	22	1	22	206.3	Longer endurance/current-oriented layout

All three official tracks use `gate_inner_size_m` = [1.5, 1.5]. The first two are configured as `clean_gate`; Mixed Endurance is configured as `current_gate` but can be run with `current-profile` none for clean validation.

TABLE V
TOP-LEVEL FIELDS OF THE OFFICIAL OBSERVATION PASSED TO A CONTROLLER. SENSOR SUBKEYS FOLLOW THE CONFIGURED PARTICIPANT PROFILE (TABLE III).

Field	Meaning
<code>time_s</code>	Race time in seconds.
<code>participant_id</code>	Identifier of the vehicle receiving the observation.
<code>sensors</code>	Filtered onboard sensor dictionary (camera, depth, IMU, DVL/velocity, collision, derived heading/depth).
<code>beacon</code>	Acoustic message: validity, target gate id, bearing, elevation, range, signal strength, and mode.
<code>race</code>	Participant-visible progress: status, lap, completed gates, target gate id, target sequence index, and official-timing status.

TABLE VI
HIGH-LEVEL CONTROLLER COMMAND FIELDS OF (3).

Field	Interpretation
<code>surge</code>	Forward/backward body-frame thrust; positive moves forward.
<code>sway</code>	Lateral body-frame thrust; sign follows the adapter mixing of (4).
<code>heave</code>	Vertical thrust; additionally constrained by depth-safety logic near the bounds.
<code>yaw</code>	Yaw-rate command; sign follows the adapter mixing of (4).

Visual servo. When the front camera detects a gate, normalized image errors $(c_x, c_y) \in [-1, 1]^2$ nudge sway and yaw by $-c_x$ and heave by $-c_y$ above thresholds (0.13, 0.30), refining the acoustic approach in the final meters.

Exit hold and smoothing. After each gate the controller holds an exit heading for a fixed number of steps to clear the structure, and all commands are low-pass filtered, $u_t = (1 - \alpha)u_{t-1} + \alpha u_t^{\text{aw}}$ with $\alpha = 0.45$, and rate-limited per axis to avoid abrupt thrust changes.

E. Additional Baselines

Table VII lists the other shipped controllers, spanning beacon-only, vision-assisted, manual, and a debug-only oracle. The oracle reads privileged gate geometry to verify that a track is feasible; it is marked debug-only and is never used for official results, illustrating how the observation contract keeps such tools out of the benchmark.

TABLE VII
BASELINE CONTROLLERS SHIPPED WITH MARINE RACE ARENA. ALL RETURN THE HIGH-LEVEL COMMAND OF (3) AND USE ONLY OFFICIAL OBSERVATIONS, EXCEPT THE ORACLE, WHICH IS A DEBUG-ONLY FEASIBILITY TOOL THAT READS PRIVILEGED GEOMETRY AND IS EXCLUDED FROM OFFICIAL RUNS.

Controller	Sensing and method
<code>rule_gate_baseline</code>	Beacon + front camera + depth hold; conservative rule-based servo with align-before-surge and post-gate exit hold (reference baseline).
<code>acoustic_baseline</code>	Beacon only; three-phase approach/transit/exit state machine.
<code>acoustic_vision_baseline</code>	Beacon plus front-camera correction blended near the gate.
<code>vision_gate_baseline</code>	Beacon-cued visual servo with multi-target detection and collision recovery.
<code>student_template</code>	Minimal proportional beacon follower provided as a starting point.
<code>oracle</code>	<i>Debug only:</i> reads ground-truth gate geometry; used for feasibility checks, never scored officially.
<code>keyboard / pygame</code>	Manual teleoperation for inspection and data collection.

VI. REFEREE, GATE VALIDATION, AND SCORING

The referee evaluates progress independently of the controller (requirement R3) using privileged ground-truth state. This section formalizes the gate-crossing test, the referee state machine and its arbitration order, the penalty model, and the single-rover score.

A. Gate Frame and Crossing

Each gate's orientation derives solely from its passage direction n_g , giving a right-handed frame

$$\hat{n} = \frac{n_g}{\|n_g\|}, \quad u_g = \frac{\hat{z} \times \hat{n}}{\|\hat{z} \times \hat{n}\|}, \quad v_g = \hat{n} \times u_g, \quad (11)$$

with world-up $\hat{z} = (0, 0, 1)$. For consecutive referee-side positions p_{t-1}, p_t , the signed distances to the gate plane are

$$d_{t-1} = \hat{n}^\top (p_{t-1} - c_g), \quad d_t = \hat{n}^\top (p_t - c_g). \quad (12)$$

A candidate crossing requires a sign change in the *correct direction*,

$$d_{t-1} \leq 0 \leq d_t \wedge (p_t - p_{t-1})^\top \hat{n} > 0, \quad (13)$$

i.e. travel from the negative to the positive side along the normal. The intersection of the segment with the plane is

$$\lambda = \frac{-d_{t-1}}{d_t - d_{t-1}}, \quad q = p_{t-1} + \lambda(p_t - p_{t-1}), \quad \lambda \in [0, 1], \quad (14)$$




Fig. 2. Gate-crossing validation on the aperture face. The segment from the previous to the current referee position must cross the gate plane in the direction of the normal n_g (out of the page here, (13)); the plane intersection q ((14)) must lie inside the shaded acceptance region—the aperture shrunk on every side by the safety margin m_g , i.e. $|u_g^\top(q - c_g)| \leq w_g/2 - m_g$ and $|v_g^\top(q - c_g)| \leq h_g/2 - m_g$ ((15)). The frame (u_g, v_g, n_g) derives from the passage direction alone.




Fig. 3. Participant lifecycle. A rover is released to RUNNING when the clock reaches its start delay τ_i ; before release it receives a zero command, its controller is not stepped, and no stuck or gate-timeout timer accrues. Running rovers reach a terminal state by finishing the last gate, missing a gate (DNF), or exceeding a timeout/stuck condition. A single-rover run is the degenerate case $\tau_i = 0$.

and the crossing is accepted only when q lies inside the aperture, shrunk by the configured safety margin m_g ,

$$|u_g^\top(q - c_g)| \leq \frac{w_g}{2} - m_g, \quad |v_g^\top(q - c_g)| \leq \frac{h_g}{2} - m_g. \quad (15)$$

Gates must be passed *in sequence*: the referee tests only the expected gate g_σ , advancing the index $\sigma \leftarrow \sigma + 1$ on a valid crossing and incrementing the lap when the sequence completes. A crossing of a *non-target* gate aperture is a missed-gate event; a sign change in the wrong direction is logged but neither penalized nor counted as progress. Fig. 2 illustrates the test and Fig. 3 the participant lifecycle.

B. Referee State Machine and Arbitration

Each participant occupies one of the states NOT_STARTED, RUNNING, FINISHED, DNF, TIMEOUT, STUCK, MANUAL_STOP, or CONTROLLER_ERROR; all but the first two are terminal. A participant is released to RUNNING when the race clock reaches its start delay (Section VII). On each tick of a running participant, events are evaluated in a fixed priority order that defines the arbitration: (1) a controller error

TABLE VIII
REFeree EVENTS, THEIR DEFAULT PENALTY, AND WHETHER THEY TERMINATE THE RUN. ALL VALUES ARE CONFIGURABLE IN THE REFeree BLOCK; THE TABLE LISTS THE v0.1 DEFAULTS ACTUALLY APPLIED BY THE REFeree.

Event	Penalty	Terminal?
Gate/bar or generic collision	5.0 s each	no
Obstacle collision	5.0 s each [†]	no
Out of bounds	10.0 s each	no
Stuck (soft)	15.0 s once	no
Missed gate	—	yes (DNF)
Timeout (if enabled)	—	yes
Gate-timeout stuck (external)	—	yes
Inter-vehicle (penalize mode)	5.0 s/pair	no (team-level)

[†]Per-obstacle value if specified, else the collision default. A per-event cooldown (1.0 s) prevents a single sustained contact from being charged repeatedly.

transitions to CONTROLLER_ERROR and returns; (2) out-of-bounds; (3) generic collision; (4) obstacle collision; (5) stuck accumulation; (6) duration timeout; (7) gate progression. Steps (2)–(5) accrue penalties without terminating; only (1), (6), a missed gate in (7), or an external gate-timeout produce a terminal state. To avoid double counting, an obstacle-collision tick suppresses the generic collision flag, and each event class has an independent cooldown.

C. Penalties and Termination

Penalties accumulate into a per-participant total P_i ; the defaults applied in v0.1 are listed in Table VIII. The non-trivial conditions are formalized as follows. The instantaneous speed is $s_t = \|p_t - p_{t-1}\|/\Delta t$; a stuck accumulator integrates low-speed time,

$$T^{\text{stuck}} \leftarrow \begin{cases} T^{\text{stuck}} + \Delta t, & s_t < v_{\text{stuck}} \\ 0, & \text{otherwise,} \end{cases} \quad (16)$$

and a one-time soft penalty is charged when $T^{\text{stuck}} \geq T_{\text{max}}^{\text{stuck}}$, with $v_{\text{stuck}} = 0.02$ m/s and $T_{\text{max}}^{\text{stuck}} = 45$ s. An out-of-bounds event is $p_t \notin \mathcal{B}$. A duration timeout, disabled by default, fires when $t - t_{\text{green}} > T_{\text{max}}$. A missed gate sets DNF when missed_gate_dnf is enabled (the default), but carries no time penalty, so a disqualifying error is never cheaper than a slow-but-legal lap.

D. Single-Rover Score and Ranking

For a finished participant the official time T_i^{official} is measured between the first and last gate (or green-to-finish, per the timing mode), and the score is the penalized time

$$T_i^{\text{pen}} = T_i^{\text{official}} + P_i. \quad (17)$$

Ranking places finished participants before unfinished ones and is defined by the ascending lexicographic key

$$\kappa_i = (\bar{f}_i, T_i^{\text{pen}}, -G_i, n_i^{\text{coll}}, P_i, \text{dist}_i), \quad (18)$$

where $\bar{f}_i = 0$ for finished and 1 otherwise, G_i is the number of completed gates, n_i^{coll} the collision count, and dist_i the distance to the next gate. Thus finished rovers sort by penalized time, while unfinished rovers sort by progress, then fewer collisions, then fewer penalties, then proximity to the next gate, so partial runs remain informative rather than being discarded.

E. Event Logging

Every run emits a line-delimited event log and a machine-readable summary (requirement R5). Events include `gate_passed`, `wrong_direction`, `collision`, `obstacle_collision`, `out_of_bounds`, `stuck`, `race_finish`, and `dnf`, each timestamped and attributed to a participant. The summary records, per participant, the status, official and penalized times, completed gates, and counts of collisions, out-of-bounds, missed gates, and stuck events, together with the final ranking; in fleet mode it additionally carries the team summary of Section VII.

VII. FLEET AND TEAM EVALUATION

Fleet mode evaluates several rovers as one cooperative team rather than as competitors against each other (requirement R7). All rovers share a single simulated world and a single referee, but each keeps an independent progress and scoring state, and enabling fleet mode does not alter single-rover behavior (R6).

A. Staggered Release and Spawn Geometry

Each participant i is assigned a release delay $\tau_i = (i - 1)g$ for a start gap g , so rovers enter the course at $0, g, 2g, \dots$. A waiting rover is sent a zero command, its controller step is *not* called, and neither stuck nor gate-timeout timers accumulate; at release the referee logs a `participant_released` event and resets the rover's progress timers, so a delayed start is never misclassified as stuck or timed out. Spawn poses are laterally separated about the base spawn to reduce launch interference: the k -th rover is offset by

$$o_k = (-1)^k \left\lceil \frac{k}{2} \right\rceil s \quad (19)$$

along the right axis $(-\sin \psi_0, \cos \psi_0)$ of the base yaw ψ_0 , giving the alternating pattern $0, -s, +s, -2s, +2s, \dots$ for spacing s ; a candidate that would fall outside the world bounds $\mathcal{B}$ is scaled inward until admissible.

B. Inter-Vehicle Collision Diagnostics

Rover-rover proximity is detected on the referee side and is never exposed to controllers. For two running rovers a, b the detector applies a separable cylinder test

$$\sqrt{(x_a - x_b)^2 + (y_a - y_b)^2} \leq r_{xy} \wedge |z_a - z_b| \leq r_z, \quad (20)$$

with hysteresis (a pair is released only when the horizontal distance exceeds a release threshold r_{rel} or the vertical distance exceeds r_z) and a refractory cooldown to debounce sustained proximity. The v0.1 conservative defaults are $r_{xy} = 0.8\text{m}$, $r_z = 0.75\text{m}$, $r_{\text{rel}} = 1.05\text{m}$, and a 1.0 s cooldown. Three modes are provided: `off` disables detection;




Fig. 4. Team-level evaluation of a two-rover staggered fleet. Each rover is active from its release to its finish; the team elapsed time spans the earliest release to the latest finish ((23)), and the team summary aggregates gates, penalties, and inter-vehicle events ((21)–(24)). Per-rover rows remain diagnostics; the fleet competes as one team rather than against itself.

diagnostic logs proximity events without changing any score; and `penalize` additionally charges one team-level penalty per pair event. Because the geometric thresholds have not yet been calibrated against measured BlueROV2–BlueROV2 contact distances (Section IX), `diagnostic` is the recommended default.

C. Team-Level Score

Only the `team_summary` is the official fleet result; per-rover rows are diagnostics. For a team $\mathcal{P}$ of N rovers with G_{exp} expected gates each,

$$G_{\text{team}}^{\text{exp}} = N G_{\text{exp}}, \quad G_{\text{team}}^{\text{done}} = \sum_{i \in \mathcal{P}} G_i^{\text{done}}. \quad (21)$$

The team finishes only if every rover finishes,

$$\text{finished}_{\text{team}} = \bigwedge_{i \in \mathcal{P}} \text{finished}_i, \quad (22)$$

in which case the team elapsed time spans the first release to the last finish,

$$T_{\text{team}} = \max_i T_i^{\text{finish}} - \min_i T_i^{\text{release}}, \quad (23)$$

and the penalized team score aggregates all penalties,

$$T_{\text{team}}^{\text{pen}} = T_{\text{team}} + P_{\text{gate}} + P_{\text{obstacle}} + P_{\text{ivc}}, \quad (24)$$

where P_{ivc} counts each inter-vehicle event once per pair, not once per rover. Fig. 4 illustrates the aggregation.

VIII. EXPERIMENTAL EVALUATION

The goal of this evaluation is not to establish an optimal underwater racing controller, but to demonstrate that the benchmark, the official observation contract, the referee, and the team scoring behave as specified, and to characterize the reference baseline honestly—including where it fails.

A. Protocol and Metrics

All runs use the HoloOcean adapter with the BlueROV2 agent, seed 0, control timestep $\Delta t = 0.033$ s, official mode (so the controller receives no ground-truth state), and the unmodified official gate apertures (requirement R4). The reported metrics are: course completion (all expected gates passed in sequence and the finish reached), completed gates, official time, penalized time of (17), and the counts of collisions, out-of-bounds, and stuck events. Because HoloOcean execution is not bit-deterministic across runs, single-seed results should be read as representative rather than as exact expectations; multi-seed statistics over the three tracks are deferred to future work (Section IX), and we are explicit about this limitation rather than reporting a single number as if it were a population mean.

B. Single-Rover Clean Tracks

Table IX reports the baseline on the three official tracks under clean, current-free conditions. The baseline completes all three courses—12/12, 17/17, and 22/22 gates—without collisions, out-of-bounds, or stuck events, with official times of 228.4 s, 315.5 s, and 542.4 s on Horseshoe Bay, Vertical Serpent, and Mixed Endurance respectively. The Vertical Serpent result was re-verified by an independent run for this paper. These results establish that the official clean tasks are feasible under the official observation contract and that the referee accepts well-formed sequential crossings; they are infrastructure-level evidence, not a claim of optimality.

C. Staggered Fleet and Team Scoring

The bottom block of Table IX reports the stable two-rover staggered demonstration on Horseshoe Bay. Both rovers finish 12/12 gates; the leading rover is collision-free, while the trailing rover incurs a single gate-bar contact, charged at 5 s, so its penalized time exceeds its official time by 5 s. The team summary aggregates to 24/24 gates, an elapsed time of 339.7 s from first release to last finish, and a penalized team time of 344.7 s. No inter-vehicle proximity events are recorded, which is expected: the 90 s start gap deliberately separates the rovers in time so that the run exercises multi-agent spawning, staggered release, independent per-rover referee state, ranking, and team aggregation *without* physical interaction, until rover-rover arbitration is calibrated. We therefore present this as a valid low-contact smoke demonstration, not as evidence that close-proximity fleet racing is solved.

D. Behavior Under Currents

The clean-track performance does not transfer to disturbance conditions. In experimental current-enabled runs the rule baseline, which uses no current compensation, fails to complete Horseshoe Bay under the medium current profile in our trials (zero gates completed), and current-compensation modes are not yet able to recover clean-track performance. We report this negative result deliberately: it quantifies the gap that the current-compensation track family and future controllers must close, and it follows the benchmark principle that failures are reported honestly rather than hidden by enlarging gates,

relaxing referee margins, or weakening the current profile (requirement R4).

E. Reproducibility

Every run emits a structured event log and a summary, and the repository provides the exact commands used here [18]. The single-entry configuration-driven launch and the per-track JSON files make the official runs reproducible from a single command; the fleet demonstration and the unit-test suite are likewise scripted, so the results in Table IX can be regenerated end to end.

IX. DISCUSSION, LIMITATIONS, AND FUTURE WORK

A. What v0.1 Claims

The release candidate demonstrates that a HoloOcean-backed underwater gate-racing benchmark can be configured declaratively, executed, scored by an independent referee, and reproduced under a clean official observation contract. It provides three official tracks, a deterministic rule baseline, single-rover clean validation, staggered multi-rover execution, team-level scoring, and inter-vehicle proximity diagnostics. These are infrastructure contributions: the artifact is the benchmark and its evaluation contract, not a claim of optimal underwater racing control.

B. What v0.1 Does Not Claim

Several components are intentionally outside the official v0.1 claims. Current compensation is not solved: as Section VIII shows, the baseline fails under disturbance currents, and the experimental DVL-based and observer-based compensation modes do not yet recover clean-track performance. Close-proximity fleet racing is not validated: the inter-vehicle detector uses conservative geometric thresholds ((20)) whose calibration against measured BlueROV2–BlueROV2 contact distances is incomplete—an earlier calibration sweep terminated before collecting contact samples—so diagnostic mode remains the recommended default and penalized mode is experimental. Consistent with requirement R4, these limitations are reported rather than masked by changing gates, margins, or current profiles.

C. Future Work

We see five concrete next steps.

(i) *Inter-vehicle collision calibration.* Sweep two BlueROV2 agents over a grid of relative longitudinal, lateral, vertical, and yaw offsets, record the engine contact sensor against the geometric separation, and fit the thresholds (r_{xy}, r_z, r_{rel}) of (20) to the measured contact envelope, replacing the current conservative defaults.

(ii) *Penalized fleet mode.* With calibrated thresholds, validate the `penalize` mode under controlled close-proximity scenarios in which rovers are deliberately brought within contact range, confirming that team penalties are charged once per pair event and that scoring remains stable.

(iii) *Cooperative fleet strategies.* Move beyond staggered, non-interacting starts to deliberate coordination—formation

TABLE IX

HOLOOCEAN VALIDATION OF THE `RULE_GATE_BASELINE` ON THE THREE OFFICIAL TRACKS. SINGLE-ROVER RUNS USE SEED 0, $\Delta t = 0.033$ s, OFFICIAL MODE, AND NO OBSTACLES, MOTION COMPENSATION, OR ACTIVE CURRENTS. THE FLEET SMOKE USES TWO ROVERS ON HORSESHOE BAY WITH A 90 s START GAP, A 3.0 m LATERAL OFFSET, AND INTER-VEHICLE DETECTION IN DIAGNOSTIC MODE.

Mode	Track / participant	Gates	Status	Official time (s)	Penalized time (s)	Collisions	OOB/stuck
Single rover	Horseshoe Bay	12/12	FINISHED	228.36	228.36	0	0/0
Single rover	Vertical Serpent	17/17	FINISHED	315.55	315.55	0	0/0
Single rover	Mixed Endurance (clean profile)	22/22	FINISHED	542.42	542.42	0	0/0
Fleet smoke	bluerov2_01, Horseshoe Bay	12/12	FINISHED	228.86	228.86	0	0/0
Fleet smoke	bluerov2_02, Horseshoe Bay	12/12	FINISHED	235.49	240.49	1	0/0
Fleet team	fleet_01 aggregate	24/24	FINISHED	339.67 (elapsed)	344.67	1 (total)	0/0

keeping and leader–follower passing of gate sequences—under the realistic constraints of low-bandwidth, high-latency acoustic communication [15] and uncertain relative localization, as analyzed in the AUV-formation literature [16], [17].

(iv) *Current-compensation track family.* Promote the analytic current field of (6)–(8) into a graded benchmark family (medium and strong profiles) and develop controllers that reject it, for example a DVL-based feedforward that estimates and cancels the body-frame drift, or an observer that fuses DVL, IMU, and beacon residuals into a current estimate, with honest failure reporting at each level.

(v) *Learning-based controllers.* The official observation contract is directly usable as a reinforcement-learning interface: the observation (beacon range, bearing, and elevation; depth; IMU; DVL/velocity; and camera) is the state, the normalized command of (3) is the action, and a natural reward combines sequential gate progress with time and collision penalties mirroring (17). A model-free policy-gradient or actor–critic method could be trained at scale against the kinematic fallback adapter and fine-tuned in HoloOcean, in the spirit of champion-level drone racing trained against a simulated opponent [5] and of GPU-parallel robot-learning environments [13], [14], all without weakening the official observation contract.

X. CONCLUSION

This paper presented Marine Race Arena, a configurable benchmark layer for underwater gate racing and fleet/team evaluation in HoloOcean. Its defining commitment is a strict separation between the autonomy stack and the evaluation: a controller is integrated by implementing three methods and sees only a documented official observation, while an independent referee uses privileged state to validate gate crossings, penalties, timeouts, ranking, and team scoring. We formalized the benchmark instance, the gate-crossing geometry, the single-rover and team-level scoring, and the staggered fleet semantics, and we grounded each in the implementation. In HoloOcean validation, a deterministic rule baseline completes the three official clean tracks collision-free under the official observation contract, and a stable two-rover staggered run produces independent per-rover summaries and a single aggregated team score; under disturbance currents the same baseline fails, a gap we report rather than hide. The release candidate is therefore a reproducible benchmark foundation,

with current compensation, close-proximity fleet arbitration, calibrated rover–rover penalties, and learning-based controllers identified as concrete and open next steps.

REFERENCES

- [1] J. Betz, H. Zheng, A. Liniger, U. Rosolia, P. Karle, M. Behl, V. Krovi, and R. Mangharam, “Autonomous vehicles on the edge: A survey on autonomous vehicle racing,” *IEEE Open Journal of Intelligent Transportation Systems*, vol. 3, pp. 458–488, 2022.
- [2] M. O’Kelly, H. Zheng, D. Karthik, and R. Mangharam, “F1TENTH: An open-source evaluation environment for continuous control and reinforcement learning,” in *Proceedings of the NeurIPS 2019 Competition and Demonstration Track*, ser. Proceedings of Machine Learning Research, vol. 123, 2020, pp. 77–89. [Online]. Available: <https://proceedings.mlr.press/v123/o-kelly20a.html>
- [3] S. Hoffmann, S. Sagmeister, T. Betz, J. Bongard, S. Büttner, D. Ebner, D. Esser, G. Jank, S. Goblirsch, A. Langmann, M. Leitenstern, L. Ögretmen, P. Pitschi, A.-K. Schwehn, C. Schröder, M. Weinmann, F. Werner, B. Lohmann, J. Betz, and M. Lienkamp, “Head-to-head autonomous racing at the limits of handling in the A2RL challenge,” *arXiv preprint arXiv:2602.08571*, 2026, abu Dhabi Autonomous Racing League; preprint, not peer reviewed. [Online]. Available: <https://arxiv.org/abs/2602.08571>
- [4] P. Foehn, D. Brescianini, E. Kaufmann, T. Cieslewski, M. Gehrig, M. Muglikar, and D. Scaramuzza, “AlphaPilot: Autonomous drone racing,” in *Robotics: Science and Systems (RSS)*, 2020. [Online]. Available: <https://www.roboticsproceedings.org/rss16/p081.html>
- [5] E. Kaufmann, L. Bauersfeld, A. Loquercio, M. Müller, V. Koltun, and D. Scaramuzza, “Champion-level drone racing using deep reinforcement learning,” *Nature*, vol. 620, no. 7976, pp. 982–987, 2023.
- [6] E. Potokar, S. Ashford, M. Kaess, and J. G. Mangelson, “HoloOcean: An underwater robotics simulator,” in *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, 2022, pp. 3040–3046. [Online]. Available: <https://www.ri.cmu.edu/publications/holocean-an-underwater-robotics-simulator/>
- [7] M. von Benzon, F. F. Sørensen, E. Uth, J. Jouffroy, J. Liniger, and S. Pedersen, “An open-source benchmark simulator: Control of a BlueROV2 underwater robot,” *Journal of Marine Science and Engineering*, vol. 10, no. 12, p. 1898, 2022.
- [8] M. M. M. Manhães, S. A. Scherer, M. Voss, L. R. Douat, and T. Rauschenbach, “UUV simulator: A Gazebo-based package for underwater intervention and multi-robot simulation,” in *OCEANS 2016 MTS/IEEE Monterey*, 2016, pp. 1–8.
- [9] P. Cieřlak, “Stonefish: An advanced open-source simulation tool designed for marine robotics, with a ROS interface,” in *OCEANS 2019 - Marseille*, 2019, pp. 1–6.
- [10] M. M. Zhang, W.-S. Choi, J. Herman, D. Davis, C. Vogt, M. McCarrin, Y. Vijay, D. Dutia, W. Lew, S. Peters, and B. Bingham, “DAVE aquatic virtual environment: Toward a general underwater robotics simulator,” in *Proceedings of the IEEE/OES Autonomous Underwater Vehicles Symposium (AUV)*, 2022, pp. 1–8. [Online]. Available: <https://arxiv.org/abs/2209.02862>

- [11] A. Dosovitskiy, G. Ros, F. Codevilla, A. Lopez, and V. Koltun, "CARLA: An open urban driving simulator," in *Proceedings of the Conference on Robot Learning (CoRL)*, ser. Proceedings of Machine Learning Research, vol. 78, 2017, pp. 1–16. [Online]. Available: <https://proceedings.mlr.press/v78/dosovitskiy17a.html>
- [12] Autonomous Robotics Research Center, Technology Innovation Institute, "A2RL: Abu Dhabi autonomous racing league," <https://a2rl.io>, 2024, autonomous car and drone racing championship; accessed 2026-06-27.
- [13] G. Brockman, V. Cheung, L. Pettersson, J. Schneider, J. Schulman, J. Tang, and W. Zaremba, "OpenAI gym," *arXiv preprint arXiv:1606.01540*, 2016. [Online]. Available: <https://arxiv.org/abs/1606.01540>
- [14] V. Makoviychuk, L. Wawrzyniak, Y. Guo, M. Lu, K. Storey, M. Macklin, D. Hoeller, N. Rudin, A. Allshire, A. Handa, and G. State, "Isaac gym: High performance GPU-based physics simulation for robot learning," in *Proceedings of the NeurIPS 2021 Track on Datasets and Benchmarks*, 2021. [Online]. Available: <https://arxiv.org/abs/2108.10470>
- [15] M. Stojanovic and J. Preisig, "Underwater acoustic communication channels: Propagation models and statistical characterization," *IEEE Communications Magazine*, vol. 47, no. 1, pp. 84–89, 2009.
- [16] Y. Yang, Y. Xiao, and T. Li, "A survey of autonomous underwater vehicle formation: Performance, formation control, and communication capability," *IEEE Communications Surveys & Tutorials*, vol. 23, no. 2, pp. 815–841, 2021.
- [17] T. Yan, Z. Xu, and S. A. Gadsden, "Formation control of multiple autonomous underwater vehicles: A review," *Intelligence & Robotics*, vol. 3, no. 1, pp. 1–22, 2023.
- [18] A. Bedei, "HoloDroneCompetition: Marine race arena repository," <https://github.com/AndreaBedei/HoloDroneCompetition>, 2026, v0.1 release candidate, accessed 2026-06-27.