

Marine Race Arena: A Configurable HoloOcean Benchmark for Underwater Gate Racing and Team-Level Fleet Evaluation

Andrea Bedei^a, Lorenzo Bacchiani^a, Giovanni Pau^a, Roberto Girau^a

^a*Department of Computer Science and Engineering, University of Bologna, Italy*

Abstract

We present Marine Race Arena, a configurable and reproducible framework for developing and evaluating autonomous systems in underwater gate-racing tasks. Complete race scenarios can be specified through declarative configurations, including tracks, vehicles, onboard sensors, acoustic beacons, environmental currents, obstacles, starting conditions, race rules and penalties. A central design principle is the strict separation between autonomy and evaluation: controllers operate exclusively from onboard observations, acoustic measurements and optional inter-vehicle communication, while an independent referee accesses privileged simulator state only to validate ordered gate crossings, enforce the rules and compute the official results. This common interface enables different autonomy strategies to be evaluated under the same sensing and scoring conditions and supports both single-vehicle trials and cooperative multi-vehicle scenarios. The reference implementation includes onboard visual-acoustic gate perception, two interpretable rule-based controllers, leader-follower coordination mechanisms and reinforcement-learning controllers integrated through the same participant-level information boundary and action interface. Three ready-to-use race circuits with different geometric characteristics are provided together with structured logging and automated validation tools for reproducible experimentation. Experiments conducted in HoloOcean with a BlueROV2-class vehicle assess course completion, gate progression, timing, collisions, robustness to environmental currents and fleet coordination. The results show that Marine Race Arena can expose meaningful differences between autonomy strategies and operating conditions while maintaining a common evaluation protocol, providing a unified testbed for studying underwater perception, navigation, control, learning and cooperation under repeatable and progressively challenging racing scenarios.

Keywords: underwater robotics, autonomous racing, robotics benchmark, reproducible evaluation, onboard perception, multi-robot systems, reinforcement learning, HoloOcean

1. Introduction

Autonomous racing has emerged as an effective research format for evaluating perception, planning, control and real-time decision-making within a single quantitative task under repeatable conditions [1]. In ground robotics, F1TENTH provides standardized vehicles, circuits, metrics and baselines for continuous control and reinforcement learning [2], while full-scale competitions such as the Abu Dhabi Autonomous Racing League (A2RL) extend the same principle to head-to-head racing near the limits of vehicle handling [3]. In aerial robotics, gate-based competitions have driven the development of high-speed perception and control pipelines [4], and learning-based systems have more recently demonstrated autonomous drone-racing performance at champion level [5]. Across these domains, the value of racing is not limited to speed: a fixed task, explicit course geometry, common rules and measurable outcomes provide a practical basis for comparing, reproducing and improving complete autonomous systems. Underwater robotics provides a natural but substantially different setting for this type of evaluation. Satellite-based global positioning is unavailable underwater, inertial and Doppler-based navigation accumulates uncertainty over time, and visual navigation is affected by limited visibility, illumination variations and water conditions [6]. Acoustic sensing and communication can extend the observable range but introduce limited bandwidth, propagation delay and measurement uncertainty [7], while persistent currents and strongly damped vehicle dynamics further affect how a vehicle approaches, aligns with and crosses a gate. Underwater racing is therefore not simply a ground or aerial racing task transferred to a marine simulator.




Figure 1: Underwater gate racing in Marine Race Arena. A BlueROV2-class vehicle approaches the ordered, beacon-marked gate line of an official circuit in HoloOcean. The controller sees this scene only through its forward camera, its depth, inertial and Doppler sensors and the acoustic packets it physically receives; the gate positions visible to the reader are never available to it.

A growing ecosystem of underwater simulators already provides the physical and sensor models required to develop autonomous systems. UUV Simulator introduced a Gazebo-based environment for underwater vehicle dynamics and multi-robot experimentation [8], while HoloOcean provides Unreal Engine environments, simulated underwater sensing and multi-agent support [9]. More recently, MarineGym has emphasized high-throughput simulation for learning-based underwater control [10]. These platforms provide increasingly capable simulation backends, but a simulator alone does not define a racing benchmark. It does not necessarily specify the ordered task to be completed, which information a controller is allowed to use, how gate crossings are validated, how penalties and timeouts are applied, or how different participants are evaluated under a common set of rules. This distinction is particularly important in simulation, where global vehicle pose, exact gate coordinates, the true current field or official course progress can simplify navigation substantially despite being unavailable to a deployed underwater vehicle. At the same time, relying only on the controller’s internal estimate to determine task completion can hide perception, localization or sequencing failures. A meaningful benchmark should therefore separate the information available to autonomy from the information required for evaluation: controllers should operate from observations representative of onboard sensing, whereas privileged simulator state should remain confined to an independent evaluation mechanism.

This paper presents Marine Race Arena, a configurable HoloOcean benchmark for autonomous underwater gate racing and team-level fleet evaluation (Fig. 1). Marine Race Arena is implemented as a benchmark and race-management layer above the simulation backend. Complete race scenarios are defined through declarative configurations specifying the environment, ordered gates, vehicles, sensor profiles, acoustic beacons, environmental currents, obstacles, participants, starting conditions, timing rules and penalties. The reference implementation and the experiments reported in this work use HoloOcean with a BlueROV2-class vehicle [11], while simulator-specific sensing and actuation are isolated behind an adapter layer. A central feature of the framework is its controller-agnostic autonomy interface: the benchmark does not prescribe how a vehicle should navigate the course, but only defines the observations available to the participant and the commands it must provide. A controller can therefore be rule-based, optimization-based, learning-based or follow another autonomy paradigm without requiring changes to the race runner or evaluation logic. In parallel, a configurable automated referee uses privileged simulator state exclusively to validate ordered gate crossings, detect collisions and other rule violations, apply penalties and timeouts, determine race termination and compute the official results. Referee-side state and official course progress are never returned

to the controller. The same architecture extends to multi-vehicle experiments, where each vehicle retains independent autonomy and local state while the framework provides simultaneous or staggered starts, optional inter-vehicle communication, proximity monitoring, distributed leader–follower coordination and team-level scoring.

The release includes three ready-to-use gate-racing circuits with different geometric characteristics, together with configurable current profiles, obstacles and participant configurations. The common controller interface is demonstrated with two interpretable rule-based controllers and a recurrent reinforcement-learning controller, showing that substantially different autonomy approaches can be integrated without modifying the race-management or scoring logic. Structured event logs and machine-readable race summaries record both participant execution and official referee outcomes. The experimental evaluation exercises the framework across different circuits, environmental conditions and multi-vehicle configurations, with the purpose of assessing whether heterogeneous autonomy strategies and race scenarios can be executed and compared within the same participant-level information boundary, action interface and independent evaluation mechanism.

The main contributions of this work are as follows:

1. **A configurable benchmark for autonomous underwater gate racing.** Marine Race Arena defines complete race scenarios through declarative configurations covering tracks, ordered gates, vehicles, onboard sensors, acoustic beacons, environmental currents, obstacles, participants, starting conditions, timing, penalties and scoring. Three ready-to-use circuits are provided, while the same configuration model supports custom racing scenarios.
2. **A controller-agnostic autonomy interface.** Controllers interact with the benchmark through a common observation–action contract and can implement rule-based, optimization-based, learning-based or other autonomy strategies without modifying the race execution or evaluation logic. The release provides rule-based and reinforcement-learning controllers as reference implementations of this interface.
3. **A configurable race manager with an independent automated referee.** Race conditions and rules are defined independently from the controller, while privileged simulator state is restricted to the evaluation layer. The referee validates ordered gate crossings, detects violations, applies penalties and timeouts and computes official race results without exposing privileged progress information to autonomy.
4. **Single-vehicle and team-level multi-vehicle evaluation within the same framework.** Marine Race Arena supports independent per-vehicle autonomy, staggered or simultaneous starts, optional acoustic-inspired communication, proximity monitoring, distributed leader–follower coordination and team-level scoring while maintaining independent official progress for each participant.
5. **An experimental validation of the common benchmark interface.** The framework is exercised in HoloOcean across different race circuits, current conditions and multi-vehicle configurations, including both rule-based reference controllers and a reinforcement-learning controller integrated through the same participant-level information boundary and action interface. Structured outputs and automated validation procedures support repeatable evaluation and extension with new controllers and racing scenarios.

The remainder of the paper is organized as follows. Section 2 reviews underwater simulation, autonomous-racing benchmarks, underwater perception, multi-vehicle cooperation and learning-based marine control. Sections 3 and 4 introduce the benchmark model, runtime architecture, HoloOcean environment and observation interface. Section 5 describes onboard gate perception, followed by the reference controllers in Section 6 and the independent referee in Section 7. Section 8 presents multi-vehicle execution and coordination, while Section 9 describes the integration of learning-based control. Section 10 reports the experimental evaluation. The final sections discuss the results, limitations and reproducibility before concluding the paper.

2. Related Work

Underwater robotics has benefited from increasingly capable simulation environments. UUV Simulator extends Gazebo with underwater vehicle dynamics, sensors and multi-robot support [8]. Stonefish provides a marine-oriented simulator with hydrodynamics, underwater rendering and ROS integration [12], while DAVE combines underwater vehicle simulation with optical and acoustic sensing [13]. HoloOcean uses Unreal Engine to provide visually rich

environments, onboard sensor models and multi-agent execution [14]. UNav-Sim focuses on vision-based underwater navigation in photorealistic environments [15], whereas MarineGym targets high-throughput reinforcement-learning experiments using GPU-accelerated underwater simulation [10]. These platforms differ in physical fidelity, sensing, throughput and intended tasks, but primarily provide the simulation substrate on which autonomy is developed. UroBench addresses reproducible comparison at a different level by evaluating underwater simulators through common navigation and learning tasks [16]. Marine Race Arena is complementary to both directions: the underlying simulator determines how vehicles, sensors and the environment evolve, while the benchmark defines the race, the participant interface and how the result is officially evaluated.

Autonomous racing provides the closest methodological reference. F1TENTH combines a standardized scaled vehicle with racing environments, interchangeable controllers and support for both conventional and learning-based methods [2]. Its simulator can also execute multiple vehicles in the same environment. At full scale, A2RL extends autonomous racing to head-to-head operation near the handling limits of the vehicle [3]. In aerial robotics, AlphaPilot established ordered gate racing as a system-level task for perception, planning and control [4], while Swift demonstrated the potential of deep reinforcement learning in high-performance drone racing [5]. CARLA offers a related example from autonomous driving: the simulator supports multiple simultaneously active vehicles [17], and its Autonomous Driving Leaderboard defines a participant interface, restricts access to privileged simulator information and evaluates route completion and infractions outside the submitted agent [18]. Although these systems address different domains and tasks, they show the value of separating the autonomy method from a common task definition and evaluation procedure.

Underwater gate traversal has also been investigated directly, although mainly as part of specific missions rather than as a reusable racing protocol. FeelHippo demonstrates underwater gate navigation through integrated perception and control [19], while Harada et al. consider vision-based navigation for an underwater Gate Pass task [20]. The broader underwater perception problem is shaped by the absence of satellite positioning, uncertain localization and limited optical visibility. Vision-based localization and control have therefore been studied on real underwater vehicles [6], and perception-aware navigation has been used to explicitly account for the visibility of relevant targets during motion [21]. Acoustic sensing provides complementary range and bearing information but introduces limited bandwidth, latency, multipath and time-varying uncertainty [7]. In Marine Race Arena these sensing problems remain part of the participant autonomy stack: visual detections and acoustic measurements may be combined by the controller, but privileged gate geometry is not used to resolve the target.

The same distinction becomes important in multi-vehicle operation. Cooperation between AUVs depends strongly on communication availability and relative-state estimation [22], while formation-control literature highlights the close coupling between localization, communication and coordinated motion [23]. Marine Race Arena does not propose a new formation or acoustic communication method. Instead, it provides the race-level mechanisms needed to test such strategies under a common protocol. Each vehicle retains its own controller and local state, optional information is exchanged through an acoustic-inspired communication channel and the independent race manager tracks per-vehicle progress, proximity events and team-level results. The included leader-follower strategy is used only as a reference implementation for exercising these capabilities.

Learning-based control provides a further case in which a common interface is useful. MarineGym exposes observation and action spaces designed for underwater reinforcement-learning tasks [10], while general environment interfaces such as OpenAI Gym demonstrate the broader value of separating an algorithm from the environment with which it interacts [24]. Marine Race Arena follows the same interface principle but is intended primarily for race execution and evaluation rather than high-throughput training. A participant is required only to consume the allowed observations and produce the expected vehicle command. Its internal implementation may therefore be rule-based, optimization-based, learned or hybrid. The recurrent reinforcement-learning controller used in this work serves as one example of this extensibility and is evaluated within the same participant-level information boundary and action interface, using the same referee as the rule-based controllers.

Table 1 summarizes the resulting positioning. The comparison considers only properties relevant to the benchmark contribution: an explicit racing task, a participant-facing controller contract, restriction of privileged simulator state, evaluation independent from the participant, multiple controllable vehicles in the same scenario and an aggregate team result. These criteria do not rank simulator quality, and multi-vehicle support alone does not imply team-level evaluation. Marine Race Arena combines these elements around a configurable underwater ordered-gate task, allowing heterogeneous controllers to be evaluated by the same configurable race manager and independent referee.

Table 1: Benchmark-level capabilities of representative systems.

System	Domain	Race task	Agent contract	Privileged-state restriction	Independent evaluator	Multi-vehicle	Team score
UUV Simulator [8]	Underwater	–	–	–	–	✓	–
Stonefish [12]	Underwater	–	–	–	–	✓	–
DAVE [13]	Underwater	–	–	–	–	✓	–
HoloOcean [14]	Underwater	–	–	–	–	✓	–
UNav-Sim [15]	Underwater	–	–	–	–	✓	–
URoBench [16]	Underwater	–	✓	–	–	–	–
MarineGym [10]	Underwater	–	✓	–	–	–	–
F1TENTH [2]	Ground	✓	✓	–	–	✓	–
CARLA [17] / Leaderboard [18]	Ground	–	✓	✓	✓	✓	–
Marine Race Arena (this work)	Underwater	✓	✓	✓	✓	✓	✓

✓: explicitly supported; –: not established as a benchmark-level property in the cited system.

3. Benchmark Model and System Architecture

Marine Race Arena combines four architectural elements: a declarative race configuration, a set of participants and their autonomy interfaces, a race-management and evaluation layer, and a simulator adapter. This separation defines the benchmark independently of both the controller implementation and the simulation backend. A race can therefore be reconfigured without changing the autonomy code, while a new controller or vehicle can be connected without changing the race definition or adjudication logic.

3.1. Benchmark Instance

Definition 1 (Benchmark instance). A Marine Race Arena instance is defined by the tuple

$$\mathcal{E} = (\mathcal{W}, \mathcal{G}, \mathcal{S}, \mathcal{F}, \mathcal{C}, \mathcal{O}, \mathcal{P}, \mathcal{R}), \quad (1)$$

where:

- $\mathcal{W}$ is the simulation world and navigable volume;
- $\mathcal{G} = (g_1, \dots, g_M)$ is the ordered gate sequence;
- $\mathcal{S}$ defines the starting conditions, including participant spawn and release settings;
- $\mathcal{F}$ defines the race finish condition;
- $\mathcal{C}$ represents the environmental current field;
- $\mathcal{O}$ is the set of obstacles in the scenario;
- $\mathcal{P} = \{1, \dots, N\}$ is the participant set, with each participant associated with a vehicle, sensor profile and controller;
- $\mathcal{R}$ contains the race and referee configuration, including validation criteria, penalties, timeouts and scoring rules.

The tuple is a conceptual description of the benchmark instance. Its concrete values are supplied by the declarative track and race configuration and are interpreted by the configuration loader, arena builder and simulator adapter at run time.

3.2. Architectural Requirements

The requirements in Table 2 express the properties that make the benchmark configurable, controller-agnostic and independently evaluable. They describe architectural constraints rather than a particular autonomy algorithm or vehicle implementation.

Table 2: Design requirements.

ID	Requirement
R1	The race is defined through a declarative configuration covering the scenario, gates, participants, environmental conditions and evaluation rules.
R2	Autonomy uses only the information exposed through the participant interface and receives no information produced by the official evaluation system.
R3	Progression, violations and the final result are determined independently of the controller implementation.
R4	Controllers based on different paradigms can be connected through the same participant/controller interface.
R5	Simulator- and vehicle-specific operations are isolated through an adapter.
R6	Every execution produces structured events and machine-readable results, including per-participant and, when applicable, team-level results.

3.3. Runtime Architecture

The complete execution flow is shown in Fig. 2. A race starts from the declarative configuration, which is checked by the configuration loader and used by the arena builder to instantiate the scenario through the selected simulator adapter. Once the arena has been created, the race runner becomes the central runtime component: it advances the simulation, coordinates the participants and mediates their interaction with the vehicle backend. At each control step, the runner obtains the information made available to each participant, passes it to the corresponding controller and forwards the returned vehicle command to the adapter for execution. This cycle is repeated throughout the race and applies independently to every participating vehicle.




Figure 2: Marine Race Arena runtime architecture.

At the same time, the referee receives the information required to determine official course progression, rule violations, penalties and completion. Its decisions are used for adjudication and reporting but are not returned to the controllers, so the evaluation path does not become part of the autonomy loop. The framework records the resulting events and per-participant results and, in multi-vehicle configurations, can additionally produce the configured team-level result. Fig. 3 provides the complementary conceptual view of this architecture, emphasizing the two information flows: one connects participant observations to vehicle commands, while the other supports race evaluation and scoring.

4. Vehicle, Simulator, Sensing and the Observation Contract

4.1. Simulator and Vehicle

The reference adapter used in the HoloOcean experiments reported here drives HoloOcean [9] with a BlueROV2-class vehicle [11]. The simulator advances at 30 physics ticks per second.




Figure 3: Controller and referee information flows.

Table 3: Onboard sensor suite.

Sensor	Rate	Output and configuration
Depth	30 Hz	Scalar depth (m); noise $\sigma = 0$.
IMU	30 Hz	Linear acceleration and angular rate, with bias terms.
DVL	15 Hz	Body-frame velocity (m/s); beam elevation 22.5° , max range 50 m.
Front camera	30 Hz	RGB image, 640×480 , 90° field of view.
Collision	optional	Boolean onboard contact flag when configured.

The experiments use two control rates over this common physics rate: $\Delta t = 0.033$ s (≈ 30 Hz), with one control step per physics tick, and $\Delta t = 0.1$ s (10 Hz). The evaluation protocol identifies the rate used by each campaign; the two rates are not pooled.

The portable participant command uses four normalized body-frame channels,

$$u_{i,t} = [u^{\text{surge}}, u^{\text{sway}}, u^{\text{heave}}, u^{\text{yaw}}]^T \in [-1, 1]^4. \quad (2)$$

The reference HoloOcean adapter maps this abstract command to the eight BlueROV2 thrusters by a fixed mixing matrix. The four vertical thrusters receive u^{heave} ; the four horizontal thrusters receive

$$\tau_{1-4} = s_\tau \begin{bmatrix} u^{\text{surge}} + u^{\text{sway}} + \kappa u^{\text{yaw}} \\ u^{\text{surge}} - u^{\text{sway}} - \kappa u^{\text{yaw}} \\ u^{\text{surge}} + u^{\text{sway}} - \kappa u^{\text{yaw}} \\ u^{\text{surge}} - u^{\text{sway}} + \kappa u^{\text{yaw}} \end{bmatrix}, \quad (3)$$

where $\kappa = 0.35$ is the yaw-mixing gain and s_τ the thruster scale; each component is clamped to the thruster limit. A direct-thruster interface remains an adapter-specific capability, but all results reported here use the abstract four-degree-of-freedom command.

4.2. Sensor Suite

Table 3 lists the onboard sensors available to a participant. The official profile combines a forward RGB camera, a depth sensor, an IMU and a DVL; an onboard contact flag is also permitted when configured.

Before specifying beacon placement, we define the logical frame associated with each gate. For gate g with passage direction n_g , let

$$\hat{n} = \frac{n_g}{\|n_g\|}, \quad \hat{r}_g = \frac{\hat{z} \times \hat{n}}{\|\hat{z} \times \hat{n}\|}, \quad \hat{s}_g = \hat{n} \times \hat{r}_g, \quad (4)$$

with world-up $\hat{z} = (0, 0, 1)$. When the passage direction is parallel to the vertical axis, $\hat{r}_g$ is defined with respect to the global $\hat{x}$ axis. This passage direction defines a right-handed gate frame, reused by both the beacon model and the referee crossing test.

4.3. Acoustic Beacon Model

Using the right-handed frame of (4), each gate carries one beacon, placed at $b_g = c_g + \delta_n \hat{n} + \delta_r \hat{r}_g + \delta_s \hat{s}_g$ relative to the gate centre c_g (default offset $(0, 0, 0.35)$ m along $\hat{s}_g$). For a receiver at position p_r with heading ψ , let $\delta = b_g - p_r$, let $\rho = \|\delta\|$ be the geometric range and let $\tilde{\rho}$ be its noisy delivered measurement. The beacon reports $\tilde{\rho}$, a yaw-relative bearing, an elevation and a signal strength:

$$\rho = \|\delta\|, \quad \tilde{\rho} = \max(0, \rho + \eta_\rho), \quad \eta_\rho \sim \mathcal{N}(0, \sigma_\rho^2), \quad (5)$$

$$\beta = \text{wrap}(\text{atan2}(\delta_y, \delta_x) - \psi) + \eta_\beta, \quad (6)$$

$$\varepsilon = \text{atan2}(\delta_z, \sqrt{\delta_x^2 + \delta_y^2}) + \eta_\varepsilon, \quad (7)$$

$$\text{SS} = \max\left(0, 1 - \frac{\tilde{\rho}}{\rho_{\max}}\right),$$

where $\rho_{\max}$ is the beacon range and $\text{wrap}(\cdot)$ maps to $(-180^\circ, 180^\circ]$. Bearing and elevation have independent angular perturbations $\eta_\beta, \eta_\varepsilon \sim \mathcal{N}(0, \sigma_\theta^2)$, while range has perturbation η_ρ with standard deviation σ_ρ . The signal strength uses the same noisy range. Packets are absent beyond $\rho_{\max}$ or when a Bernoulli dropout occurs. The official tracks use a common per-track value in $\{0.2, 0.45, 0.6\}$ for the angular and range noise standard deviations (degrees and metres, respectively), with dropout probabilities up to 0.04.

Every ordered gate has one independent periodic transmitter with a unique contiguous identifier $B01, \dots, BM$. All beacons transmit independently of referee state, and each vehicle receives every in-range, non-dropped packet. The received packets do not indicate which beacon is currently relevant; the controller maintains its expected beacon from its own local course state and selects the corresponding packet (Sections 5 and 6). Packet randomness is seeded by the run, transmitter, receiver and transmission index, keeping fleet reception reproducible and independent of participant count and call order.

4.4. Current Field Model

The current field is the sum of analytic primitives evaluated at vehicle position p and time t ,

$$\mathbf{v}_c(p, t) = \sum_j \mathbf{v}_j(p, t), \quad (8)$$

The available primitives are a constant flow, a localized jet, a sinusoidal component and a vortex. A localized jet centred at c_j with radius R_j contributes, for $\|p - c_j\| \leq R_j$,

$$\mathbf{v}_j = \mathbf{v}_0 \omega(d_j), \quad \omega(s) = \begin{cases} \max(0, 1 - \frac{s}{R_j}) & \text{linear,} \\ \exp(-\frac{s^2}{2\sigma_j^2}) & \text{Gaussian,} \end{cases} \quad (9)$$

where $d_j = \|p - c_j\|$ is the jet distance and $\sigma_j = \max(R_j/2, 10^{-6})$ the jet width; the contribution is zero outside R_j . The constant term contributes a fixed velocity $\mathbf{v}_0$, while a sinusoidal term modulates one axis, $v_{j,\text{axis}}(t) = v_{0,\text{axis}} + A \sin(2\pi f t + \phi)$. A *vortex* of radius R_j , tangential speed V_t and vertical speed V_z contributes, with $r = \sqrt{\Delta x^2 + \Delta y^2}$ and tangent $\hat{\theta} = (-\Delta y, \Delta x)/r$,

$$\mathbf{v}_j = (\hat{\theta}_x V_t \omega(r), \hat{\theta}_y V_t \omega(r), V_z \omega(r)). \quad (10)$$

The provided strong profile combines a constant set-flow of $[0.75, 1.05, 0]$ m/s, two Gaussian jets, a vortex and a vertical sinusoid; the medium profile scales these magnitudes by 0.5. Both profiles are defined relative to the gate

Table 4: Official track configurations.

Track file	Track name	Gates	Laps	Total gates	Length (m)	Intended use
marine_race_horseshoe_bay.json	Marine Race Horseshoe Bay	12	1	12	93.8	Clean-gate horseshoe baseline
marine_race_vertical_serpent.json	Marine Race Vertical Serpent	17	1	17	118.3	Vertical and slalom gate sequencing
marine_race_mixed_endurance.json	Marine Race Mixed Endurance	22	1	22	206.3	Longer endurance and current-oriented layout

layout and are available on every official circuit. The adapter applies the field physically before each control step but does not report it directly to the controller, which must infer the disturbance from its onboard measurements, including body-frame DVL velocity.

4.5. Declarative Track Configuration

A track is a self-contained JSON file. Its sections define the race timing, benchmark task, world, ordered gates, sensors, beacons, currents, obstacles, participants and referee. The selected task supplies the corresponding validation contract; Listing 1 shows a representative excerpt.

Listing 1: Representative excerpt of a track configuration.

```

"track": {
  "gate_inner_size_m": [1.5, 1.5],
  "gate_sequence": ["G01", "G02", "..."]
},
"gates": [
  { "id": "G01", "position": [4.0, -2.0, -3.5],
    "rotation_rpy_deg": [0.0, 0.0, 0.0],
    "passage_direction": [1.0, 0.0, 0.0] }
],
"beacon": {
  "range_m": 90.0,
  "angular_noise_std_deg": 0.2,
  "range_noise_std_m": 0.2,
  "dropout_probability": 0.0,
  "update_rate_hz": 10.0
},
"participants": [
  { "id": "bluerov2_01", "vehicle": "BlueROV2",
    "controller": "rule_gate_baseline",
    "start_delay_s": 0.0 }
],
"referee": {
  "penalties": { "minor_collision_s": 5.0 },
  "gate_validation": {
    "vehicle_clearance_margin_m": 0.1 }
}

```

The three official circuits are summarized in Table 4 and their layouts shown in Fig. 4. They are deliberately heterogeneous: Horseshoe Bay is a short, largely planar baseline; Vertical Serpent forces repeated vertical transitions and slalom sequencing; Mixed Endurance is more than twice as long as Horseshoe Bay and combines both. All three use 1.5×1.5 m apertures, and neither the aperture size nor the referee margin was changed at any point of this study.

The obstacle_gate task accepts explicit fixed obstacles or deterministic random boxes generated between consecutive gates with a configured density and minimum clearance. The HoloOcean adapter records the requested and physically spawned counts. Obstacle avoidance is supported by the task configuration but is not evaluated in the results reported here.




Figure 4: Official race-circuit layouts.

Table 5: Participant observation fields.

Field	Meaning
<code>local_time_s</code>	Elapsed time since this vehicle’s release; it is not official race time.
<code>sensors</code>	Filtered onboard dictionary containing only the configured camera, depth, IMU, DVL and optional collision measurements.
<code>beacons</code>	List of packets delivered by the simulated acoustic channel from independent transmitters; it may be empty or contain several identifiers, and no field marks which one is the next gate.
<code>comms</code>	Present only when the optional inter-vehicle acoustic channel is enabled: an <code>inbox</code> of messages received from teammates, each with the sender identifier, the controller-authored payload and a receiver-local reception timestamp (Section 8).

4.6. Official Observation Contract

At each control step, the participant receives an observation dictionary whose fields are listed in Table 5. It contains participant-local time, the configured onboard sensors, packets delivered by the acoustic beacon channel and, when enabled, optional inter-vehicle communication. The observation builder is separated from the referee and uses only the participant state and allow-listed sensing interfaces; evaluation data therefore cannot enter the observation through this path. Communication payloads are authored by the corresponding controllers from their own legal information. World-frame velocity, pose-derived heading and depth, and environmental-current telemetry may exist internally for simulation or adjudication but are excluded from the official observation.

The absence of three classes of information defines the task. The controller never receives the global ground-truth vehicle pose; any localization estimate must therefore be derived from participant-accessible onboard information. It does not receive ground-truth gate positions or orientations from the simulator, so the spatial relation to the course must be inferred from onboard sensing and acoustic measurements. It also receives neither the referee’s notion of the next gate, nor whether a crossing was valid, nor the number of remaining gates; course progression is consequently a controller-side estimate that can be wrong.

5. Onboard Gate Perception and Target Association

A racing benchmark is only meaningful if target selection remains part of the participant’s autonomy problem. The observation contract of Section 4.6 does not expose ground-truth gate geometry or referee progress, while the forward camera may contain several gates and the acoustic channel may simultaneously deliver packets from several transmitters. The received packets do not indicate which beacon is currently relevant; that state is maintained locally by the controller. Selecting which visible gate corresponds to the locally expected beacon must therefore be performed onboard rather than by the benchmark.

The reference perception front end combines geometric detection in the forward camera with acoustic target association, as summarized in Fig. 5. The visual stage produces multiple candidate gates rather than a single preselected target, while the acoustic stage provides bearing and range information for the beacon expected by the controller. Their combination produces either one visual target or no visual target for the current frame.




Figure 5: Onboard gate-perception pipeline.

5.1. Gate Detection from the Forward Camera

The gates are rendered as square frames formed by light-coloured bars. Direct colour or brightness segmentation is unreliable because underwater rendering can substantially alter their apparent intensity relative to the water and seabed. The detector therefore relies primarily on geometric structure: edges are extracted from the grayscale image, closed contours are converted into bounding boxes, and candidates are filtered according to their apparent size, aspect ratio and consistency with a rectangular frame. These criteria are fixed for the evaluation and are applied unchanged across the three official circuits.

For an image of width W and height H , the centre of a surviving candidate is expressed in normalized image coordinates as

$$c_x = \frac{(x + \frac{1}{2}w_{px}) - \frac{1}{2}W}{\frac{1}{2}W}, \quad c_y = \frac{(y + \frac{1}{2}h_{py}) - \frac{1}{2}H}{\frac{1}{2}H}. \quad (11)$$

Here (c_x, c_y) is the displacement of the candidate centre from the optical axis, while its apparent size is represented by the image-area fraction α . Each candidate is also assigned a confidence γ that combines shape consistency, aspect-ratio plausibility, apparent size and image centredness. A single fixed parameter set is used for all official circuits; the exact thresholds and weights are provided with the released implementation.

The detector retains multiple plausible candidates and removes near-duplicate detections before passing the resulting ranked list to the association stage. Retaining several candidates is important in the racing setting because consecutive gates can be simultaneously visible; a visual detector that selected a single rectangle by itself would implicitly solve part of the course-association problem.

5.2. Acoustic–Visual Target Association

The controller maintains its own locally expected beacon (Section 6). When a packet from that beacon is available, it provides a noisy bearing β and measured range $\tilde{\rho}$ according to the acoustic model of Section 4. The perception problem is then to determine which, if any, of the visual candidates is geometrically consistent with that acoustic observation.

The conceptual association path is shown in Fig. 6.

Association first rejects candidates that are inconsistent with the camera field of view and the acoustic bearing. The horizontal image position of each candidate provides an approximate visual bearing, which can be compared with β . Apparent target size is also used as a consistency check at short range: a nearby expected gate should occupy a substantial portion of the image, whereas a small complete rectangle is more likely to belong to a farther gate.

When the acoustic bearing provides a clear directional constraint, candidates that strongly disagree with it are rejected. The remaining detections are ranked using visual confidence, image centredness, relative apparent size and




Figure 6: Conceptual onboard visual-acoustic perception path.

acoustic-visual bearing consistency. Relative rather than purely absolute size is used so that a farther but visually centred gate does not systematically dominate a closer candidate. The same association parameters are used throughout the official circuits.

If no packet from the locally expected beacon is currently available, the association falls back to visual evidence and favours large, confident and well-centred candidates. If no candidate remains plausible, no visual target is returned and the controller continues using the other onboard measurements.

This design preserves the benchmark information boundary. Target association depends only on the controller’s local course state, camera detections and received acoustic measurements. It does not use simulator gate poses, the global vehicle pose or the referee’s expected-gate index. An incorrect association is therefore an autonomy error rather than an error hidden by benchmark-side target selection.

5.3. Exploratory Pose-Aware Extension

The centre-based representation does not describe the orientation of the gate plane. We therefore also implemented a pose-aware variant that uses a detected aperture quadrilateral together with the known nominal gate aperture to estimate relative plane orientation and distance.

Its behaviour is strongly dependent on viewing geometry. In a dedicated set of oblique views, a valid pose estimate was obtained in only 11.7% of frames and the median plane-yaw error was 48.2° . The estimator is therefore unreliable under large oblique viewing angles, motivating the separate characterization under the viewing conditions encountered during normal racing.

5.4. Perception Audit on the Official Circuits

To characterize perception under the conditions encountered by the reference controller, we recorded 60 frames per circuit along its trajectories on the three official circuits. Ground truth was used only offline to evaluate the perception output and never entered the control loop. Table 6 summarizes detection availability, quadrilateral recovery, pose availability and multi-candidate target association, while Fig. 7 shows representative states.

Detection was available in approximately 88–92% of the recorded frames, while a validated four-corner aperture was recovered in approximately 78–83%. A metric pose estimate was available in roughly half of the frames, confirming that pose recovery is substantially less reliable than centre-based gate detection. The reference controllers are consequently designed to tolerate intermittent visual measurements rather than assume that a complete gate geometry is available at every control step.

Table 6: Perception audit on the official circuits.

Circuit	Detection	4-corner	Pose avail.	Multi-candidate association
Horseshoe Bay	90.0%	81.7%	51.7%	7/7
Vertical Serpent	88.3%	78.3%	48.3%	4/4
Mixed Endurance	91.7%	83.3%	53.3%	11/11




Figure 7: Representative onboard gate-perception states.

Target association was considerably more stable. In every recorded frame containing multiple visual candidates, the online association selected the same gate as the offline reference: 7/7 cases on Horseshoe Bay, 4/4 on Vertical Serpent and 11/11 on Mixed Endurance. No unintended online target switch was observed. This supports the use of acoustic bearing primarily as an identity cue when several geometrically plausible gates are visible.

The pose results also clarify the difference between the dedicated oblique-view test and normal racing. During the reference trajectories, the controller tends to approach the gate approximately frontally. Under these conditions, the recovered plane orientation was stable, with median yaw error below 1.2° on all three circuits and no gross orientation or sign errors observed. Under strongly oblique views, by contrast, pose recovery became sparse and inaccurate. The relevant conclusion is therefore not that the pose estimator is generally reliable, but that its accuracy depends strongly on viewing geometry.

For the racing task considered here, reliable gate-centre detection and acoustic-visual association are sufficient for the reference controllers, while full pose recovery is substantially more sensitive to viewpoint. The perception audit characterizes these components independently from the course-completion results reported in Section 10.

6. Controller Interface and Rule-Based Reference Controllers

A central usability goal of Marine Race Arena is that an autonomy stack can be integrated through a small interface without modifying the runner, the referee or the adapter. This section specifies that interface, the dynamic loading mechanism, the controller-local course tracker and the two rule-based reference controllers whose comparison anchors the evaluation.

6.1. Controller Interface

A controller implements the participant policy

$$\pi_i : o_{i,t} \mapsto u_{i,t}, \quad (12)$$

mapping the official observation $o_{i,t}$ to the vehicle command $u_{i,t}$. The software contract that realizes this mapping consists of three methods: `reset(mission_info)`, called once before the run with a minimal static assignment;

Table 7: High-level controller command fields of (2).

Field	Interpretation
surge	Longitudinal body-frame thrust; positive moves forward.
sway	Lateral body-frame thrust; sign follows the adapter mixing of (3).
heave	Vertical thrust; additionally constrained by depth-safety logic near the world bounds.
yaw	Yaw-rate command; sign follows the adapter mixing of (3).

`step(observation)`, called at every control tick to receive the official observation of (12) and return the command of (2); and `close()`, called at teardown. The contract is duck-typed, so any object exposing the three callables is accepted and a participant need not depend on framework base classes. A manual controller may raise a dedicated stop exception from `step` to end a run cleanly; any other exception is recorded as a controller error and terminates the participant. Listing 2 shows a minimal example.

Listing 2: Minimal custom controller.

```
class MyController:
    def reset(self, mission_info):
        self.tracker = LocalCourseTracker(
            initial_beacon_id=mission_info["initial_beacon_id"],
            total_beacons=mission_info["total_beacons"],
            laps=mission_info["laps"])

    def step(self, observation):
        sensors = observation["sensors"]
        state = self.tracker.update(
            local_time_s=observation["local_time_s"],
            beacons=observation["beacons"],
            camera_image=sensors.get("FrontCamera"),
            dvl_velocity=sensors.get("DVLSensor"))
        if state.finished:
            return {"surge": 0.0, "sway": 0.0, "heave": 0.0, "yaw": 0.0}
        # ... map state.beacon and camera to a command ...
        return {"surge": 0.3, "sway": 0.0, "heave": 0.0, "yaw": 0.0}

    def close(self):
        pass
```

The `mission_info` passed to `reset` is deliberately minimal: the initial beacon identifier, the number of beacons and the lap count. The high-level command fields are listed in Table 7; values are normalized to $[-1, 1]$ and clamped by the framework before the action mapping of (3). In fleet mode a controller may additionally return a small message payload alongside its command and read incoming teammate messages from the observation.

6.2. Dynamic Loading

A controller is selected at run time by a built-in alias, a dotted module path with an explicit class, or a file path and class name imported dynamically. The runner also provides manual controllers for inspection and data collection. An external policy is therefore evaluated without modifying the package:

```
python -m marine_race_arena.scripts.run_marine_race \
--track ../marine_race_horseshoe_bay.json \
--benchmark-task clean_gate \
--controller path/to/my_controller.py \
--controller-class MyController \
--adapter holoocean --official --headless
```




Figure 8: The LOCALCOURSETRACKER state machine shared by both reference controllers. ADVANCE loops back to SEARCH for the next gate and reaches FINISHED after the last locally confirmed passage. No transition consumes referee state.

The same mechanism loads the learned policies of Section 9: a thin wrapper exposes the three interface methods, encodes the official observation into the policy’s input vector and returns the policy’s output as the command of (2). No part of the runner, referee or adapter is aware that the controller is neural.

6.3. Controller-Local Course Progression

Because the official observation does not contain referee progress (Section 4.6), each controller maintains its own estimate of which gate it is currently approaching. The reference controllers use the reusable LocalCourseTracker, initialized from the mission assignment and updated exclusively from participant-accessible onboard information.

The tracker maintains the locally expected beacon and progresses through the states

$$\text{SEARCH} \rightarrow \text{APPROACH} \rightarrow \text{VISUAL_ALIGN} \rightarrow \text{COMMIT} \rightarrow \text{VERIFY_EXIT} \rightarrow \text{ADVANCE},$$

before returning to SEARCH for the next gate or entering FINISHED after the final locally confirmed passage.

At each control step, the tracker selects the packet corresponding to its locally expected beacon and combines acoustic bearing and range with the associated visual target and body-frame DVL motion. A passage is not inferred from a single measurement. The vehicle first establishes a stable visual–acoustic alignment, then commits to the aperture, and finally checks whether its subsequent onboard measurements are consistent with having crossed and moved away from the expected gate. Forward DVL displacement, the evolution of beacon range and bearing, and the persistence or disappearance of the visual target jointly provide this evidence.

Temporary loss of a camera detection or an acoustic packet is therefore insufficient to advance the local course state, as is an isolated change in range without supporting motion evidence. When the available measurements do not consistently support a completed passage, the tracker remains on, or returns to, the current gate rather than advancing to the next one. This conservative progression logic keeps gate sequencing inside the participant while avoiding any dependence on referee feedback.

Fig. 8 summarizes this controller-local state machine. Controller-local progression and referee decisions are logged separately for evaluation and are never fed back into the controller.

The two controllers described below use the same tracker and the same onboard information. They are intended as interpretable reference baselines rather than optimized control policies.




Figure 9: Center-then-commit passage and local recovery.

6.4. Continuous Servo

Both reference controllers combine acoustic guidance, visual centring, depth hold and DVL-based passage verification through the shared local tracker of Fig. 8. They differ only in how visual feedback is used during the final gate passage.

Continuous Servo steers toward the locally expected beacon using its bearing and elevation, and refines that alignment with the associated visual target as the gate becomes visible. Depth feedback maintains the vertical reference, while DVL motion contributes to passage verification. During `COMMIT` and `VERIFY_EXIT`, the controller continues applying bounded visual corrections while moving through the aperture, continuously servoing toward the observed gate centre.

6.5. Center-then-Commit

Continuous visual correction is useful during approach but can become undesirable very close to the gate. As the vehicle enters the aperture, parts of the frame leave the camera field of view and the detected gate geometry may become partial. The resulting image centre can move even when the vehicle is already correctly aligned, causing unnecessary late corrections when clearance is smallest.

Center-then-Commit uses the same observations, acoustic-visual guidance, local tracker and passage-verification logic as *Continuous Servo*. The difference is restricted to the final traversal. Once a stable alignment has been established, the controller stops reacting to subsequent image-centre changes during `COMMIT` and `VERIFY_EXIT`. It instead proceeds through the aperture while maintaining the previously established depth and attitude from onboard feedback.

This matched design isolates the passage strategy: continuous visual servoing versus a stabilized committed traversal. Their comparative behaviour is evaluated in Section 10.3.

The same distinction is summarized in Fig. 10, which also contrasts continuous servoing with center-then-commit passage behaviour.

7. Independent Referee and the Information Boundary

The referee is the component that makes results comparable, and it is the component a participant must not be able to influence. This section formalizes the two-sided information boundary, the gate-crossing test, the referee state machine and arbitration order, the penalty model and the scoring rules, and it closes by treating the boundary as something to be measured rather than asserted.

7.1. The Two-Sided Boundary

The benchmark partitions the run state into two disjoint views. Writing x_t for the full simulator state at step t , the *controller view* of participant i is

$$o_{i,t} = \Phi_i(x_t) = (t_i^{\text{loc}}, \mathcal{Z}_{i,t}, \mathcal{A}_{i,t}, \mathcal{M}_{i,t}), \quad (13)$$

where t_i^{loc} is time since this participant's release, $\mathcal{Z}_{i,t}$ the allow-listed onboard measurements, $\mathcal{A}_{i,t}$ the set of acoustic packets delivered by the simulated channel at step t , and $\mathcal{M}_{i,t}$ the optional teammate inbox. The *evaluator view* is

$$e_t = \Psi(x_t) = (\{p_{i,t}, \dot{p}_{i,t}\}_{i \in \mathcal{P}}, \mathcal{G}, \mathcal{O}, \mathcal{B}, \mathbf{v}_c), \quad (14)$$




Figure 10: LocalCourseTracker states and passage styles.

comprising exact vehicle poses and velocities, the ordered gate geometry, the obstacle set, the world bounds and the true current field. The benchmark enforces

$$\Phi_i \perp \sigma(\ell_t, G_{i,t}^{\text{done}}, \text{status}_{i,t}, P_{i,t}) \quad \text{for all } i, t, \quad (15)$$

that is, the observation map is independent of the referee's expected-gate index ℓ_t , of official completed-gate counts, of participant status and of accrued penalties. In the implementation this is a structural property rather than a convention: the observation builder's signature admits the track configuration, the arena, the adapter and the participant state, and does not admit a referee object, so referee-derived quantities are unreachable from the code path that constructs $o_{i,t}$.

Information flows in one direction only. The referee reads x_t and the controller's issued commands; the controller reads neither e_t nor any function of it. In particular, a controller is never told that a crossing was accepted, never told which gate is expected next, and never told that it has been penalized. Where the controller's own estimate of its progress disagrees with the referee's, the disagreement is logged and left standing (Section 7.7).

7.2. Gate Frame and Crossing Test

The referee reuses the right-handed gate frame $(\hat{n}, \hat{r}_g, \hat{s}_g)$ of (4), defined from the passage direction n_g with the vertical-normal fallback to the world $\hat{x}$ axis. For consecutive referee-side positions p_{t-1}, p_t , the signed distances to the gate plane are

$$d_{t-1} = \hat{n}^\top (p_{t-1} - c_g), \quad d_t = \hat{n}^\top (p_t - c_g). \quad (16)$$

A candidate crossing requires a sign change in the *correct direction*,

$$d_{t-1} \leq 0 \leq d_t \wedge (p_t - p_{t-1})^\top \hat{n} > 0, \quad (17)$$

i.e. travel from the negative to the positive side along the normal. The intersection of the segment with the plane is

$$\lambda = \frac{-d_{t-1}}{d_t - d_{t-1}}, \quad q = p_{t-1} + \lambda(p_t - p_{t-1}), \quad \lambda \in [0, 1], \quad (18)$$

and the crossing is accepted only when q lies inside the aperture, shrunk by the configured safety margin m_g ,

$$|\hat{r}_g^\top (q - c_g)| \leq \frac{w_g}{2} - m_g, \quad |\hat{s}_g^\top (q - c_g)| \leq \frac{h_g}{2} - m_g. \quad (19)$$




Figure 11: Gate-crossing geometry and acceptance region.




Figure 12: Automatic participant lifecycle.

Gates must be passed *in sequence*: the referee tests only the expected gate g_ℓ , advancing the index $\ell \leftarrow \ell + 1$ on a valid crossing and incrementing the lap when the sequence completes. Crossing the aperture of a *non-target* gate is a missed-gate event; a sign change in the wrong direction is logged but neither penalized nor counted as progress. Fig. 11 illustrates the test and Fig. 12 the principal automatic lifecycle.

Two design choices in this test are worth naming. Testing only the expected gate, rather than all gates, is what turns the course into an ordered task instead of a collection of independent hoops. Shrinking the aperture by m_g before the containment test makes the pass criterion strictly conservative: a trajectory that clips the frame is not credited.

7.3. Referee State Machine and Arbitration

Each participant occupies one of the states NOT_STARTED, RUNNING, FINISHED, DNF, TIMEOUT, STUCK, DSQ, MANUAL_STOP or CONTROLLER_ERROR; all but the first two are terminal. A participant is released to RUNNING when the race clock reaches its start delay (Section 8). On each tick of a running participant, events are evaluated in a fixed priority order that defines the arbitration: (1) a controller error transitions to CONTROLLER_ERROR and returns; (2) out-of-bounds; (3) generic collision; (4) obstacle collision; (5) stuck accumulation; (6) duration timeout; (7) gate progression. Steps (2) to (5) accrue penalties without terminating; only (1), (6), a missed gate in (7) or an external gate-timeout produce a terminal state.

The order is not arbitrary. Placing controller errors first prevents a crashed participant from also accruing collision or stuck charges from its final frozen command. Placing gate progression last ensures that a tick in which the vehicle both crosses a gate and touches its frame is charged for the contact and credited for the crossing, rather than either event masking the other. To avoid double counting, an obstacle-collision tick suppresses the generic collision flag.

Table 8: Default referee events, penalties and terminal conditions.

Event	Penalty	Terminal
Gate, bar or generic collision	5.0 s each	no
Obstacle collision	5.0 s each [†]	no
Out of bounds	10.0 s each	no
Stuck (soft)	15.0 s per episode	no
Missed gate	N/A	yes (DNF)
Timeout	N/A	yes
Gate-timeout stuck	N/A	yes
Inter-vehicle proximity	0 or 5.0 s per pair	no (team-level)

[†]Per-obstacle value if specified, else the collision default. Collision, obstacle, out-of-bounds and inter-vehicle proximity charges use independent configured cooldowns (default 1.0 s); stuck uses one latch per episode and terminal events have no cooldown.

Repeated collision, obstacle and out-of-bounds charges use independent cooldowns; inter-vehicle proximity uses its own team-level cooldown; stuck is latched once per episode; and terminal events are not cooldown-limited.

7.4. Penalties and Termination

Penalties accumulate into a per-participant total P_i ; the applied defaults are listed in Table 8. The non-trivial conditions are formalized as follows. The instantaneous speed is $s_t = \|p_t - p_{t-1}\|/\Delta t$, and a stuck accumulator integrates low-speed time,

$$T^{\text{stuck}} \leftarrow \begin{cases} T^{\text{stuck}} + \Delta t, & s_t < v_{\text{stuck}} \\ 0, & \text{otherwise,} \end{cases} \quad (20)$$

with a soft penalty charged once per continuous low-speed episode when $T^{\text{stuck}} \geq T_{\text{max}}^{\text{stuck}}$ and $v_{\text{stuck}} = 0.02$ m/s. The configured $T_{\text{max}}^{\text{stuck}}$ is 45 s on Horseshoe Bay and 65 s on Vertical Serpent and Mixed Endurance. An out-of-bounds event is $p_t \notin \mathcal{B}$. A duration timeout, disabled by default, fires when $t - t_{\text{green}} > T_{\text{max}}$.

A missed gate sets DNF when `missed_gate_dnf` is enabled, which is the default, but carries no time penalty. The combination is deliberate: making a disqualifying error terminal but free of accumulated time ensures that losing the course is never cheaper than completing a slow but legal lap, which would otherwise be an exploitable scoring artefact.

7.5. Single-Vehicle Score and Ranking

For a finished participant the official time T_i^{official} is measured between the first and last gate (or green-to-finish, according to the timing mode) and the score is the penalized time

$$T_i^{\text{pen}} = T_i^{\text{official}} + P_i. \quad (21)$$

Ranking places finished participants before unfinished ones and is defined by the ascending lexicographic key

$$\kappa_i = \begin{cases} (0, T_i^{\text{pen}}, 0, 0, 0, \text{id}_i), & f_i = 1, \\ (1, -G_i^{\text{done}}, n_i^{\text{coll}}, P_i, \text{dist}_i, \text{id}_i), & f_i = 0, \end{cases} \quad (22)$$

where $f_i = 1$ denotes a finished vehicle, G_i^{done} the number of completed gates, n_i^{coll} the collision count and dist_i the distance to the next gate; the participant identifier is the deterministic final tie-break. Finished vehicles are therefore ordered by penalized time, and unfinished vehicles by progress, then by fewer collisions, fewer penalties and proximity to the next gate. This ordering preserves the information contained in partial runs instead of discarding them.

7.6. Event Logging

Every run emits a line-delimited event log and a machine-readable summary (requirement R6). Events are timestamped and attributed to the corresponding participant, while the summary records the final status, official and penalized times, course progress and the main penalty-related quantities used for evaluation. In fleet mode it additionally contains the team-level summary of Section 8. The experimental results of Section 10 are obtained by aggregating these outputs.

Table 9: Controller-local and referee course progression.

Circuit	Gate advancements			Mismatches		Local delay (s)	
	referee	local	matched	false	missed	median	p95
Horseshoe Bay	1120	1120	1116	10	4	0.924	1.353
Vertical Serpent	159	160	159	1	0	0.627	1.155
Mixed Endurance	195	192	192	0	3	0.462	1.287
All	1474	1472	1467	11	7	0.858	1.353

A *false* local advancement is a local increment that precedes its same-ordinal referee crossing by more than one control step, or that has no corresponding referee crossing; a *missed* local advancement is a referee crossing with no same-ordinal local increment. Across all runs there were 0 event-consistency errors, 0 finish-order inversions over 61 pairwise comparisons, and 0 cases of a controller declaring itself finished before the referee agreed; the referee recorded the finish before the local tracker in 103 cases. The mean delay is 0.722 s with a sample standard deviation of 2.83 s; the spread is dominated by a single false advancement at -90.8 s, which is retained rather than filtered.

7.7. Measuring the Boundary

Because the controller and the referee maintain independent estimates of course progress, their agreement can be evaluated offline without affecting control. The framework records both quantities separately, allowing discrepancies between local progression and official referee decisions to be inspected after each run.

Table 9 summarizes this comparison. The results support two observations. First, the information boundary is effective: controller-local progress is genuinely independent of the referee and occasional disagreements remain visible rather than being corrected by the framework. Second, this separation does not prevent practical course tracking: the onboard-only local estimate agrees with the official gate sequence in more than 99% of recorded advancements.

8. Fleet Execution, Coordination and Team-Level Evaluation

Fleet mode runs several vehicles as a single cooperative team through the same circuit, so the benchmark measures how quickly and how safely the team completes the course rather than pitting the vehicles against each other. The benchmark models one team: all vehicles share a single simulated world, each controller keeps an independent local course estimate, and the referee keeps one scoring state per vehicle. Enabling fleet mode does not alter single-vehicle behaviour.

8.1. Staggered Release and Spawn Geometry

Each participant i is assigned a release delay $\tau_i = (i - 1)\Delta\tau$ for a start gap $\Delta\tau$, so vehicles enter the course at $0, \Delta\tau, 2\Delta\tau, \dots$. A waiting vehicle receives a zero command, its controller step is *not* called, and neither stuck nor gate-timeout timers accumulate. At release the referee logs a `participant_released` event and resets the vehicle’s progress timers, so a delayed start is never misclassified as stuck or timed out. Spawn poses are laterally separated about the base spawn to reduce launch interference: using zero-based $k = 0, \dots, N - 1$, the k -th vehicle is offset by

$$o_k = (-1)^k \left\lceil \frac{k}{2} \right\rceil s \quad (23)$$

along the right axis $(-\sin\psi_0, \cos\psi_0)$ of the base yaw ψ_0 , giving the alternating pattern $0, -s, +s, -2s, +2s, \dots$ for spacing s ; a candidate that would fall outside the world bounds $\mathcal{B}$ is scaled inward until admissible.

8.2. Inter-Vehicle Proximity Diagnostics

Inter-vehicle proximity events are detected on the referee side and are never exposed to controllers, which is a direct consequence of (15): a vehicle that could read the referee’s proximity flag would be receiving privileged information about a teammate’s exact position. For two running vehicles a, b the detector applies a separable cylinder test

$$\sqrt{(x_a - x_b)^2 + (y_a - y_b)^2} \leq r_{xy} \wedge |z_a - z_b| \leq r_z, \quad (24)$$

with hysteresis — a pair is released only when the horizontal distance exceeds a release threshold r_{rel} or the vertical distance exceeds r_z — and a refractory cooldown to debounce sustained proximity. A vehicle that wants to avoid its teammates must therefore infer their presence from its own sensing and from messages they choose to send.

8.3. Inter-Vehicle Acoustic Communication

Because the fleet is a single cooperative team, Marine Race Arena provides an optional channel so that vehicles can share information while running the course. It is off by default and changes nothing for single-vehicle runs. A controller transmits by returning a small message payload alongside its command and receives an inbox of teammates' messages in its observation; what a message contains and how it is used are entirely the participant's responsibility, so the framework provides only the transport.

To stay faithful to the medium rather than model a perfect radio link, the channel reproduces the constraints of the underwater acoustic channel discussed in Section 2. A message broadcast by vehicle i reaches a teammate j only if the two are within a maximum range $R_{\max}$, after a range-dependent delay

$$\tau_{ij} = \frac{d_{ij}}{c} + \tau_{\text{proc}}, \quad (25)$$

where d_{ij} is the inter-vehicle distance, c the speed of sound and τ_{proc} a fixed processing delay. Each delivery is dropped with probability p_{loss} . Payloads are capped to a small size to model the low bandwidth, and a vehicle may not transmit faster than a configured per-sender interval. Packet loss is drawn from a seeded generator, so a run remains reproducible.

We are explicit about the status of this component: it is an acoustic-inspired transport, not a characterized channel model. It provides the substrate for the coordination policy below and is exercised as part of that validation, but its packet loss, maximum range, latency and freshness window are not independently swept, and no claim about acoustic-communication performance is made in this paper.

The resulting multi-vehicle information flow and independent team referee are shown in Fig. 13.

Multi-Vehicle Fleet Architecture in Marine Race Arena




Figure 13: Conceptual fleet architecture and team referee.

8.4. Distributed Leader–Follower Coordination

Motivation. A staggered team of *identical* controllers stays separated in time, because each vehicle takes roughly as long as the one ahead. A *heterogeneous* team does not: a faster follower released behind a slower leader closes the gap, and the two then converge on the same aperture centre at the same gate, because both are servoing to the same visual target. The resulting contacts are a property of the convoy, not of either controller in isolation, and they are severe — Section 10.5 reports a mean of 127.3 gate and world collisions per run for an uncoordinated three-vehicle team at an 8 s gap.

Module design. We provide a coordination controller that maintains a spaced convoy using only legal information. It *wraps* rather than replaces a gate-passing controller that exposes a `LocalCourseTracker`, which keeps coordination and gate navigation separate concerns; both official camera-assisted controllers satisfy this interface. Each vehicle is assigned a predecessor from the static release order carried in its legal mission information; the frontmost vehicle has no predecessor and races freely. Within the channel’s transmit-rate limit, every vehicle periodically broadcasts exactly three fields — `local_beacon_index`, `local_lap` and `local_status` — read from its own base controller’s tracker. The receiver treats these as estimates, retains the most recently received predecessor report and ignores it after a 2.5 s freshness window.

Let $\hat{g}_i$ denote cumulative locally estimated progress reconstructed from the reported lap and beacon index. A follower yields while its predecessor is fewer than Δg locally estimated gates ahead,

$$\text{yield}_j \iff \hat{g}_{\text{pred}(j)} - \hat{g}_j < \Delta g, \quad (26)$$

and otherwise runs its wrapped base controller unchanged. While yielding, the wrapper still updates the base tracker from onboard observations but commands zero motion. A predecessor locally reported as finished, a missing report or a report older than the freshness window all stop blocking. With the channel disabled no report is delivered and the wrapper reduces exactly to the uncoordinated base behaviour.

Technical advantage and its cost. Every input to (26) is legal: the vehicle’s own tracker, a static predecessor assignment and a controller-authored message that itself contains only the sender’s own estimate. Referee progress can disagree with either local estimate but never changes a coordination decision. The cost is that a deliberate wait is indistinguishable, to the referee, from being stuck: yielding can trigger the independent stuck penalty of (20), and Section 10.5 shows this happening for the more conservative two-gate margin. We regard that as correct behaviour for a benchmark — a coordination policy that buys safety by standing still should pay for the time it spends standing still — and it is one reason the recommended default margin is one gate rather than two.

8.5. Team-Level Score

Only the `team_summary` is the official fleet result; per-vehicle rows are diagnostics. For a team $\mathcal{P}$ of N vehicles with G_{exp} expected gates each,

$$G_{\text{team}}^{\text{exp}} = N G_{\text{exp}}, \quad G_{\text{team}}^{\text{done}} = \sum_{i \in \mathcal{P}} G_i^{\text{done}}. \quad (27)$$

The team finishes only if every vehicle finishes,

$$\text{finished}_{\text{team}} = \bigwedge_{i \in \mathcal{P}} \text{finished}_i, \quad (28)$$

in which case the team elapsed time spans the first release to the last finish,

$$T_{\text{team}} = \max_i T_i^{\text{finish}} - \min_i T_i^{\text{release}}, \quad (29)$$

and the penalized team score aggregates every per-vehicle penalty together with the team-level inter-vehicle term,

$$T_{\text{team}}^{\text{pen}} = T_{\text{team}} + \sum_{i \in \mathcal{P}} P_i + P_{\text{ivc}}, \quad (30)$$

where P_i is the per-vehicle penalty total of (21) and P_{ivc} counts each inter-vehicle proximity event once per pair rather than once per vehicle. The conjunction in (28) is what makes the score a *team* score: a fast leader cannot compensate for a follower that never finishes, so a coordination policy is rewarded for bringing the whole convoy home. Fig. 14 illustrates the aggregation.

9. Learning-Based Extension

Learning is included as a supporting demonstration that a learned controller can enter Marine Race Arena through the same participant interface as the reference controllers. It does not replace the benchmark’s central contribution: the official observation boundary, ordered-gate referee and scoring protocol remain unchanged.

Figure 15 places the recurrent policy in that interface and shows the curriculum progression used before the frozen validation.




Figure 14: Team-level scoring for a staggered two-rover fleet.

PPO control loop in Marine Race Arena




Figure 15: Gen-2 recurrent-policy loop and curriculum.

9.1. Final Gen-2 Observation Contract

The final experiment uses the committed onboard_local_transition_27d_v1 contract. Its 27 bounded features are ordered as follows: beacon presence, bearing sine and cosine, elevation and range; visual centre coordinates and area; depth; DVL surge, sway and heave plus presence; IMU yaw rate; the previous surge, sway, heave and yaw actions; five short-term rates for visual centre, acoustic bearing, elevation and range; a recent-target-change flag; and gate-orientation presence with yaw sine and cosine. Missing values follow the explicit contract: missing visual centre and area are zero, invalid rates are zero, and an unavailable orientation contributes zero mask and trigonometric components. The DVL presence bit is retained.

The encoding contains no gate coordinates, global pose, referee progress, future geometry or reward. It is therefore a local transition representation, not a privileged course-state representation. The four continuous body-frame actions are the same surge, sway, heave and yaw channels in $[-1, 1]$ used by the rule controllers; framework clamping and depth limits are shared.

9.2. Recurrent PPO and Training Configuration

The final policy uses Stable-Baselines3 Contrib RecurrentPPO with the committed Gen-2 recurrent-LSTM architecture:

`gen2_recurrent_lstm_27d_v1`

a frozen normalization buffer, a $27 \rightarrow 256 \rightarrow 128$ tanh encoder, separate one-layer actor and critic LSTMs with hidden size 128, 128-unit policy/value heads, and a four-action Gaussian head. The recorded environment is CPU-based; the dependency lock records Stable-Baselines3 and SB3-Contrib 2.7.1 with PyTorch 2.8.0.

The final direct-speed campaign starts from the committed Gen-2 initialization checkpoint and uses full copies of the three official circuits, three workers, current-free conditions and the recorded fog configuration. It uses no track fragments or synthetic circuits in the reported campaign. PPO settings are learning rate 6×10^{-6} , clip range 0.07, target KL 0.008, two epochs, 128 batch size, 256 rollout steps, zero entropy coefficient and gradient norm limit 0.5. The direct-speed reward records gate crossing (+20), completion (+50), time (-0.01 per step), terminal failure (-55), collision (-20), directness (+0.025) and aligned speed (+0.015); energy and jerk terms are zero in this campaign.

The preceding Gen-2 data, behavioural-cloning and DAgger artifacts are retained as initialization and audit history. They are not pooled with the final PPO validation. The official circuits were available to earlier diagnostic runs, so the present circuit results are retrospective validation of the frozen integration, not an unseen-track generalization claim.

9.3. Checkpoint Selection and Matched Evaluation

Checkpoint selection is recorded separately from training in `training_decision_log.json`. The selected campaign is `ppo_27d_direct_speed_campaign_C_20260909_retry1`; its first registered milestone produced the selected `checkpoint_050000.zip` at 50,688 actual environment steps. Its SHA-256 is recorded alongside the checkpoint and validation manifest. Later 100k diagnostics were not selected because they degraded or partially failed robust evaluation.

The final validation contains three valid episodes per official circuit at the learning control rate. Horseshoe Bay finished 3/3 with no collisions or out-of-bounds events; Vertical Serpent finished 3/3 with one collision event; Mixed Endurance finished 3/3 with no collisions or out-of-bounds events. The results are descriptive ($n = 3$ per circuit), not a statistical superiority claim. The current-free rule demonstration is a separate 10-Hz artifact and finishes all three circuits in 9/9 runs without collision, out-of-bounds or wrong-direction events; its completion rate is a useful interface reference, but its per-run times are not pooled with the PPO table.

9.4. Information-Boundary Interpretation

The learned policy operates within the same participant-level onboard information boundary and is scored by the same independent referee as every other participant. Consequently, the main learning result is not that PPO solves underwater racing. It is that the benchmark exposes a meaningful distinction between local gate-transition control and reliable ordered-sequence completion: the selected policy completed all nine validation episodes, while the surrounding audit keeps the contract, checkpoint hash and episode-level evidence independently inspectable.

10. Experimental Evaluation

This evaluation does not seek an optimal underwater racing controller. It asks whether the benchmark, the observation contract, the referee and the team scoring behave as specified, and whether they expose differences between autonomy strategies that a coarser protocol would hide. We organize the experiments around five research questions rather than chronologically, and we report failures of the reference controllers and of the learned policies as results rather than as caveats.

RQ1 Can Marine Race Arena support fully onboard autonomous completion across heterogeneous underwater racing circuits?

RQ2 Does the benchmark expose meaningful differences between autonomous control strategies?

RQ3 How do environmental currents affect benchmark performance?

RQ4 Can the same evaluation framework support multi-vehicle and team-level racing?

RQ5 Can learning-based controllers be integrated and evaluated within the same participant-level information boundary and referee procedure?

10.1. Protocol, Metrics and Provenance

Three independent bodies of evidence are reported, and they are never pooled.

The systematic benchmark matrix. 78 runs generated with the native HoloOcean simulator backend under a frozen source-tree fingerprint recorded in the artifact manifest. It covers both reference controllers on the three clean circuits and under the medium and strong Horseshoe Bay currents, homogeneous two-vehicle fleets, and matched three-vehicle coordination trials. Clean, current and fleet cases use five seeds; coordination uses three seeds under simultaneous and staggered release. Control runs at $\Delta t = 0.033$ s (30 Hz). This matrix answers RQ2, RQ3 and RQ4, and supplies the boundary measurement of Section 7.7.

The current-free completion demonstration. A focused nine-run demonstration that the reference controller completes each official circuit end to end with currents explicitly disabled and verified zero in the manifest, at $\Delta t = 0.1$ s (10 Hz). This answers RQ1.

The final learned-controller validation. The selected Gen-2 27-D RecurrentPPO checkpoint was evaluated in three valid episodes on each official circuit at $\Delta t = 0.1$ s. This answers RQ5 as an integration and diagnostic question; it is not pooled with the 78-run matrix.

All performance trials use the reference native HoloOcean simulator backend and a BlueROV2-class vehicle, official onboard-only observations, no current compensation and no obstacles, with the standard gate apertures and referee margin unchanged throughout. Separate setup checks confirm that both named current profiles and seeded static obstacles instantiate on every circuit; these are construction checks, not obstacle-avoidance trials, and no obstacle result is claimed.

The reported metrics are completion, completed gates, official and penalized time, collisions, out-of-bounds events, wrong-direction crossings and stuck events. Fleet and coordination trials add inter-vehicle proximity events and team times. Time statistics use finished runs only and are undefined where nothing finished; gate and event means use all seeds. Dispersions are sample standard deviations (ddof = 1). Proportions carry Wilson 95% intervals where sample sizes are small enough for the normal approximation to mislead. Paired comparisons use two-sided sign tests on matched seeds.

Every run records the experiment-generation commit, source or model hash, HoloOcean version, track hash, requested and actual adapter, fallback status, current profile, timestep and seed. Runs that used a non-HoloOcean adapter or permitted the engine-free fallback do not appear in the tables.

10.2. RQ1: Onboard-Only Completion Across Heterogeneous Circuits

The first question a racing benchmark must answer about itself is whether its task is solvable at all under its own restrictions. If no controller can complete a circuit using only onboard sensing, the benchmark measures the difficulty of an impossible problem rather than the quality of autonomy.

Table 10 answers this affirmatively and without qualification for the current-free setting. The Center-then-Commit controller completed all three official circuits in every one of nine runs: 3/3 on Horseshoe Bay (12 gates), 3/3 on Vertical Serpent (17 gates) and 3/3 on Mixed Endurance (22 gates), for 51 of 51 gates in total. Across those nine runs the referee recorded zero gate, bar or world collisions, zero out-of-bounds events and zero wrong-direction crossings, and penalized time equalled official time in every run because no penalty was ever charged. Every run manifest independently confirms the conditions under which this was achieved: the native HoloOcean simulator backend was used, the engine-free fallback was disabled, the requested and measured current vector was $[0, 0, 0]$, and the controller received no privileged input.

Completion time scales with course length roughly as expected from the track geometry: 236.0 ± 6.9 s over 93.8 m, 308.6 ± 6.5 s over 118.3 m and 480.4 ± 15.7 s over 206.3 m, i.e. approximately 0.40, 0.38 and 0.43 m/s of course progress. The consistency of that rate across three structurally different circuits — planar, vertically serpentine

Table 10: Current-free completion by the Center-then-Commit controller.

Circuit	Gates	Runs	Completed	Mean gates	Official time (s)	Coll.	OOB / WD
Horseshoe Bay	12	3	3/3	12.0/12	236.0 ± 6.9	0	0 / 0
Vertical Serpent	17	3	3/3	17.0/17	308.6 ± 6.5	0	0 / 0
Mixed Endurance	22	3	3/3	22.0/22	480.4 ± 15.7	0	0 / 0
Total	51	9	9/9	51.0/51	—	0	0 / 0

Times are mean ± sample standard deviation (ddof= 1) over the three runs; penalized time equals official time in every run because no penalty was charged. *Coll.*: gate, bar and world collision events. *OOB / WD*: out-of-bounds events and wrong-direction crossings. Every run manifest records `adapter_actual=holocean`, `fallback_allowed=false`, `current_profile=none` and `currents_actual=[0,0,0]`, together with the track SHA-256 and the source commit. Mixed Endurance is a `current_gate` task whose validation requires a current, so a current-free run additionally overrides the validation mode to `clean_gate`; the track geometry, gate order, laps and referee configuration are unchanged. These runs use a 10 Hz control rate and are therefore not directly comparable in time with Table 11, which uses 30 Hz.

Table 11: Systematic clean-track evaluation of the reference controllers.

Track	Controller	Seeds	Finished	Gates ↑	Official time (s) ↓	Collisions ↓
Horseshoe Bay	Continuous Servo	5	5/5	12.0/12	194.5 ± 4.8	0.0
	Center-then-Commit	5	5/5	12.0/12	197.7 ± 7.1	0.0
Vertical Serpent	Continuous Servo	5	5/5	17.0/17	236.9 ± 6.4	0.0
	Center-then-Commit	5	4/5	14.8/17	241.1 ± 1.8	0.0
Mixed Endurance	Continuous Servo	5	2/5	17.0/22	394.7 ± 3.9	0.0
	Center-then-Commit	5	5/5	22.0/22	410.7 ± 8.8	0.0

Times are mean ± sample standard deviation (ddof= 1) over finished runs; gates and collisions are five-seed means, and collisions combine gate and world events. No clean run recorded an out-of-bounds or stuck event, and penalized time equals official time throughout. Bold marks the better controller where the two differ in completion. Regenerated from the raw per-run summaries by `article/regenerate_tables.py`.

and long mixed — indicates that the limiting factor is the controller’s gate-by-gate acquire–align–commit cycle rather than any circuit-specific difficulty.

We are deliberate about the scope of this result and about its relationship to the systematic matrix. This is a *focused completion demonstration*: three seeds per circuit, one controller, no disturbance, at a 10 Hz control rate. It establishes that the task is achievable end to end under the full observation contract, and it does so with a per-run manifest that makes the claim auditable. It is not a performance characterization, and it does not supersede the broader evaluation of Sections 10.3–10.5, which uses five seeds, both controllers, three disturbance conditions and a 30 Hz control rate. The two are complementary and are reported separately for that reason; in particular, the times in Table 10 are systematically longer than those in Table 11 because the control rate differs by a factor of three, and the two must not be compared.

10.3. RQ2: Do Control Strategies Separate?

A benchmark that ranks every reasonable controller identically is not measuring anything. The two reference controllers of Section 6 are an unusually clean test of this, because they differ in exactly one respect — what happens during COMMIT and VERIFY_EXIT — while sharing observations, guidance, tracker and confirmation logic.

Table 11 and Fig. 16 show that the answer depends on the circuit, and that neither controller dominates. On the short, largely planar Horseshoe Bay circuit the two are effectively tied: both complete all 12 gates in all five seeds with no collision, out-of-bounds or stuck event, and the staged controller is 3.2 s slower (197.7 ± 7.1 s versus 194.5 ± 4.8 s), a 1.7% penalty with no measured safety benefit. Constant re-centring is cheap here, and paying nothing for it is the right trade.

The controllers separate on the harder circuits, and they separate in opposite directions. On Vertical Serpent, Continuous Servo is more reliable, finishing 5/5 against 4/5 and completing 17.0 against 14.8 mean gates. On Mixed Endurance, which is more than twice as long, the ordering reverses far more sharply: Center-then-Commit finishes all five seeds and completes every gate (22.0/22), while Continuous Servo finishes only 2/5 and averages 17.0/22.




Figure 16: Reference-controller completion fractions over the frozen evaluation conditions.

Table 12: Reference-controller robustness to environmental currents.

Profile	Controller	Seeds	Finished	Gates $\uparrow$	Official (s)	Penalized (s)	Collisions $\downarrow$
None	Continuous Servo	5	5/5	12.0/12	194.5 ± 4.8	194.5 ± 4.8	0.0
	Center-then-Commit	5	5/5	12.0/12	197.7 ± 7.1	197.7 ± 7.1	0.0
medium	Continuous Servo	5	2/5	8.4/12	337.9 ± 19.2	692.9 ± 75.7	43.4
	Center-then-Commit	5	3/5	8.8/12	217.3 ± 1.5	223.9 ± 1.8	25.2
strong	Continuous Servo	5	0/5	3.0/12	—	—	0.0
	Center-then-Commit	5	0/5	3.2/12	—	—	0.8

Times are mean $\pm$ sample standard deviation over finished runs only and are undefined where no run finished. Gates and collisions are five-seed means, and collisions combine gate and world events. The *medium* profile has a constant set-flow component of 0.65 m/s and *strong* of 1.29 m/s, against a mean clean course speed of ≈ 0.48 m/s for the same vehicle on this circuit. The near-zero collision counts under *strong* reflect less progress, not better control: both controllers stall after about three gates and rarely reach the gate-dense part of the course. Run metadata confirms physical current coupling in every run.

The mechanism is consistent with the design argument of Section 6.5. Continuous visual correction is harmless when a passage is short and the approach is well conditioned, and it becomes a liability when partial near-field contours make the image centroid move quickly at exactly the moment the aperture margin is smallest. A long course compounds this: each passage is an independent opportunity for a late correction to fail, so a small per-gate risk becomes a large per-course risk. That Continuous Servo nonetheless wins on Vertical Serpent shows the trade is not one-directional — repeated vertical transitions appear to benefit from continuous re-centring more than they suffer from late corrections — which is precisely the kind of condition-dependent result a single-circuit benchmark would have concealed. The comparison is limited to the three official circuits, two current profiles and the evaluated seeds.

10.4. RQ3: Behaviour Under Environmental Currents

Performance degrades sharply and non-uniformly as current strength increases (Table 12). Under the *medium* Horseshoe Bay profile, Continuous Servo finishes 2/5 runs, averages 8.4/12 completed gates and 43.4 gate and world collisions per run, with finished-only official and penalized times of 337.9 ± 19.2 s and 692.9 ± 75.7 s. Center-then-Commit finishes 3/5, averages 8.8/12 gates and 25.2 collisions, with finished-only times of 217.3 ± 1.5 s and 223.9 ± 1.8 s. Neither controller finishes a single *strong* run: Continuous Servo averages 3.0/12 gates with no gate or world collision, and Center-then-Commit 3.2/12 gates with 0.8.

Two features of these numbers require comment, because both are easy to misread.

First, the collision counts are *non-monotonic* in current strength: the *medium* profile records many collisions (43.4 and 25.2) while the *strong* profile records almost none (0.0 and 0.8). This is not evidence of better control under the stronger disturbance. Under *strong* the vehicles complete only 3.0–3.2 of 12 gates and stall early, so they rarely

Table 13: Multi-vehicle evaluation on clean Horseshoe Bay.

Setting	Policy	Seeds	Team finish	Team gates $\uparrow$	Team time (s) $\downarrow$	GW / IV / S $\downarrow$
<i>Two vehicles, homogeneous, gap 90 s</i>						
	Continuous Servo	5	5/5	24.0/24	303.1 ± 3.2	0.0 / 0.0 / 0.0
	Center-then-Commit	5	5/5	24.0/24	308.0 ± 5.8	0.0 / 0.0 / 0.0
<i>Three vehicles, heterogeneous convoy, gap 0 s</i>						
	No coordination	3	3/3	36.0/36	219.7 ± 2.5	9.3 / 2.3 / 0.0
	LF(1)	3	3/3	36.0/36	262.3 ± 7.7	0.0 / 0.0 / 0.0
	LF(2)	3	3/3	36.0/36	288.3 ± 1.3	0.0 / 0.0 / 1.0
<i>Three vehicles, heterogeneous convoy, gap 8 s</i>						
	No coordination	3	2/3	34.3/36	254.0 ± 24.5	127.3 / 2.0 / 0.0
	LF(1)	3	3/3	36.0/36	266.0 ± 8.3	0.0 / 0.0 / 0.0
	LF(2)	3	3/3	36.0/36	285.7 ± 3.8	0.0 / 0.0 / 0.0

Team times are raw team elapsed mean $\pm$ sample standard deviation over full-team finishes, computed by (29); gates and events are per-condition means. *GW / IV / S*: gate and world collisions, inter-vehicle proximity events and stuck events. No run recorded an out-of-bounds event. LF(Δg) is the leader-follower policy of (26) with a margin of Δg locally estimated gates. Penalized team time differs from raw team time only where events were charged: 266.4 ± 64.1 s and 564.0 ± 463.0 s for the uncoordinated convoy at gaps 0 and 8 s, and 303.3 ± 1.3 s for LF(2) at gap 0 s, whose longer hold trips the independent stuck timer of (20). In the three-vehicle condition the leader runs Continuous Servo and both followers run Center-then-Commit; no controller or course parameter was retuned between conditions.

reach the gate-dense region of the circuit where contacts accumulate. Fewer collisions here reflect less progress, not more skill, and a benchmark that reported collision counts without completion would invert the ranking.

Second, the strong profile is best read as a saturation regime rather than a graded difficulty. Its constant set-flow component alone is $[0.75, 1.05, 0]$ m/s, a magnitude of 1.29 m/s, about $2.7\times$ the mean course speed the same vehicle sustains on clean Horseshoe Bay (93.8 m in 194.5 s, i.e. ≈ 0.48 m/s). No controller restricted to this vehicle’s actuation can make sustained forward progress against it, so the condition functions as a stress case that confirms the disturbance is physically coupled rather than as an informative comparison point. The medium profile (0.65 m/s constant component) is the informative regime, and there the staged controller is clearly the more robust: more completions, roughly half the collisions, and finished runs 120.6 s faster.

The headline finding for benchmark design is the gap between Table 11 and Table 12. Both controllers score a perfect 5/5 with zero events on clean Horseshoe Bay, and drop to 2/5 and 3/5 with tens of collisions on the same circuit under a moderate current. Clean-track success is therefore not a proxy for robustness, and any evaluation protocol that reports only clean completion will rank controllers on a dimension that disappears the moment the water moves. The quantitative current sweep is limited to Horseshoe Bay (Section 12).

10.5. RQ4: Multi-Vehicle and Team-Level Evaluation

The upper block of Table 13 validates that team execution works. With either controller, a homogeneous two-vehicle fleet on clean Horseshoe Bay at a 90 s release gap finishes 24/24 team gates in all five seeds, with no gate or world collision, no inter-vehicle proximity event, no out-of-bounds and no stuck event. Mean team elapsed and penalized times coincide at 303.1 ± 3.2 s for Continuous Servo and 308.0 ± 5.8 s for Center-then-Commit. This is deliberately a low-contact case: it exercises multi-vehicle spawning, staggered release, independent per-vehicle referee state, ranking and team aggregation without the vehicles interacting, since the per-vehicle behaviour matches the single-vehicle case.

The interacting case is the interesting one. A heterogeneous three-vehicle convoy — a Continuous Servo leader followed by two Center-then-Commit vehicles — converges on the same apertures, and the consequences are severe. At an 8 s gap the uncoordinated convoy finishes only 2/3 runs, averages 34.3/36 team gates, 127.3 gate and world collisions and 2.0 inter-vehicle proximity events per run. At simultaneous release it still finishes 3/3, but averages 9.3 collisions and 2.3 proximity events.

The distributed leader-follower policy of (26) removes this entirely. At both gaps, LF(1) finishes 3/3 with 36.0/36 team gates and zero gate, world, proximity, out-of-bounds and stuck events, at 266.0 ± 8.3 s (gap 8 s) and 262.3 ± 7.7 s (gap 0 s). The more conservative LF(2) also eliminates contact but is consistently slower (285.7 ± 3.8 s and 288.3 ± 1.3 s)

Table 14: Final Gen-2 RecurrentPPO validation results.

Circuit	Gates	Episodes	Finished	Mean time (s)	Collisions	OOB
Horseshoe Bay	12	3	3/3	192.5	0	0
Vertical Serpent	17	3	3/3	239.1	1	0
Mixed Endurance	22	3	3/3	365.6	0	0
All circuits	—	9	9/9	—	1	0

The selected artifact is `artifacts_gen2/ppo_27d_direct_speed_campaign_C_20260909_retry1/final_validation_50k`; the selected checkpoint contains 50,688 actual environment steps. The separate current-free rule demonstration also finishes 9/9 episodes, but its timing records are not pooled here because it is a different artifact and controller.

and, at simultaneous release, its longer hold trips the independent stuck timer once per run, raising its penalized team time to 303.3 ± 1.3 s. A larger nominal gate separation therefore bought no additional safety in these conditions while costing both time and a penalty, which is why LF(1) is the recommended default and LF(2) is retained only as a conservative comparison.

Two properties of this result matter beyond the numbers. The coordination policy consumes only the vehicle’s own tracker state, a static predecessor assignment and teammate-authored messages delivered over a channel with range-dependent latency and loss; it never reads referee progress or a teammate’s true position, so the improvement is achieved inside the same information boundary that constrains everything else. And the cost of coordination is visible in the score rather than hidden: LF(2)’s stuck penalty is charged by the same independent referee that charges a collision, so a policy that buys safety by waiting pays for the waiting. This evidence is limited to one clean circuit, three vehicles, two release gaps and three seeds.

10.6. RQ5: Learned Controllers Under the Same Contract

Table 14 reports the final selected Gen-2 checkpoint within the same participant-level information boundary and independent referee used by the rule controllers. The policy completed all nine valid episodes: 3/3 on each circuit. There was one collision event in Vertical Serpent and no out-of-bounds events. These results are descriptive at $n = 3$ per circuit and do not support a statistical ranking.

The rule comparison is intentionally limited. A separate current-free 10-Hz reference artifact also completed 9/9 episodes with no collision, out-of-bounds or wrong-direction events. Its timing records are not pooled with the PPO timings because the campaigns have different controllers and manifests. The larger systematic rule matrix runs at 30 Hz and is reported separately; 30-Hz and 10-Hz timings are not directly comparable and are never combined into one performance estimate.

The learning result is therefore an integration and diagnosis result, not a claim that PPO outperforms the reference controller. It confirms that the Gen-2 27-dimensional onboard-derived observation encoding can drive a recurrent policy through the full evaluation harness, while the independent ordered-sequence referee keeps circuit completion distinct from isolated gate crossings. The official circuits were available to earlier diagnostic runs, so this is retrospective validation rather than unseen-track generalization.

11. Discussion

The experiments support a view of underwater gate racing as a benchmark for ordered-sequence autonomy rather than speed alone. A vehicle must commit to the correct target, reacquire a new target after each passage and compose those transitions over a long course. Small per-transition defects therefore amplify with course length, while a single-gate task can make a weak sequence policy look successful.

11.1. What the Evaluation Contract Reveals

The central value of Marine Race Arena is the conjunction of a fixed information boundary, an ordered referee, configurable race management and machine-readable scoring. The boundary makes controller comparisons meaningful because rule-based and learned controllers operate within the same participant-level information boundary.

The referee makes local progress auditable: controller estimates can be compared with privileged adjudication without allowing the latter into the control loop. The resulting residual disagreements are useful evidence, not hidden implementation details.

The configurable race manager makes this separation operational. Track geometry, participants, environmental conditions, starts, timing and penalties can be changed through scenario configuration while the participant interface and referee logic remain fixed. This permits the benchmark to vary the task and the operating condition independently of the autonomy implementation, which is important when interpreting differences between controllers.

The referee also makes failure interpretable. A missed next gate is recorded as a sequence failure rather than being blurred into a low aggregate gate count, and safety costs remain visible alongside progress. This is why the current sweep cannot be summarized by collision counts alone: a stronger disturbance can reduce collisions simply by stopping a vehicle before the contact-dense part of the course. Reporting progress, completion and events together is a benchmark requirement, not a presentation choice.

11.2. Reference Controllers, Currents and Fleets

The reference-controller comparison is a controlled ablation of visual feedback during passage. Neither controller dominates: continuous correction is useful on repeated vertical transitions, whereas suspending late visual corrections is more robust on the long mixed course and under the informative current profile. The result is evidence that the benchmark can expose a design decision that a single clean circuit would conceal, not evidence that either controller is optimal.

The fleet experiments extend the same logic from individual to team autonomy. Vehicles that are individually competent can still converge on the same aperture and collide when their interaction is not scored. A distributed leader-follower policy removes those events within the legal observation boundary, while the referee charges waiting through the stuck penalty. Thus coordination and its cost are evaluated in the same score rather than in a separate engineering trace.

11.3. Learning as a Demonstration of Generality

The final Gen-2 result supports a more precise interpretation of learning in this benchmark. A recurrent PPO controller using its 27-dimensional onboard-derived observation encoding completed all nine valid validation episodes, with three episodes on each official circuit. One collision event was recorded on Vertical Serpent and no out-of-bounds events were recorded. Thus, a learned controller can be integrated within the same participant-level information boundary and action interface, and can complete full ordered circuits under the reported current-free protocol.

This result remains an integration and diagnostic result rather than a claim that PPO is superior to the reference controllers or that it generalizes to unseen circuits. The validation uses three episodes per circuit, the official circuits had been accessible to earlier diagnostic runs, and the learned campaign uses a different control rate and provenance from the systematic rule matrix. Its timing should therefore not be ranked directly against the rule controller tables. More broadly, the experiment demonstrates that the controller-agnostic interface accommodates recurrent learning without changing race execution or adjudication, but it does not isolate the effects of the architecture, training budget, reward design or observation encoding.

The next learning experiments should extend this matched protocol to currents, multi-vehicle interaction and sealed circuits, while comparing memory, training and curriculum choices under the same independent referee. The benchmark makes these comparisons possible without changing the information boundary or the official scoring mechanism.

11.4. Implications for Benchmark Design

Taken together, the results argue for benchmarks that define both race configuration and evaluation above the simulator. Solvability, disturbance robustness, team interaction and learning generality are different questions and require separate conditions, metrics and provenance. Marine Race Arena provides a controller-agnostic interface, a configurable race manager and an independent referee for underwater ordered-gate racing; its limitations and the scope of the current evidence are given next.

12. Limitations

We state the boundaries of this work explicitly, both because they define what the reported numbers may be used for and because several of them are the natural next steps.

L1. No physical validation, and an unquantified simulation-to-reality gap. Every result in this paper was produced in simulation. No experiment was run on a physical BlueROV2 or in a water tank, and we make no claim that any reported completion rate, time or collision count would be reproduced by a real vehicle. The gap is unquantified rather than small: closing it would require a physical gate course, a calibrated acoustic positioning setup and a characterization of the optical conditions, none of which is attempted here.

L2. Backend fidelity bounds the physical claims. The benchmark inherits the hydrodynamics, rendering and sensor models of its backend. HoloOcean provides a usable and widely adopted approximation, but it is not a validated computational-fluid-dynamics model of a BlueROV2, and the current field of (8) is an analytic superposition of primitives rather than a simulated flow. Conclusions about *relative* controller behaviour are therefore better supported than conclusions about absolute achievable speeds or about the specific current magnitude at which a controller fails.

L3. Optical perception assumptions. The detector of Section 5.1 assumes gates that project to closed, near-rectangular contours of bounded aspect ratio against a background whose edge structure is less rectangular. It was developed for, and validated on, the rendering conditions of these circuits. Turbidity, marine snow, backscatter from an active light source, biofouling on the gate frames, strong caustics and low-light operation are not modelled, and the detector’s geometric filters would require re-derivation under any of them.

L4. Gate appearance assumptions. All gates in this study are square frames of identical $1.5\text{ m} \times 1.5\text{ m}$ inner aperture, and the detector’s size, aspect and near-field area gates are calibrated to that geometry. The known-aperture assumption is also what makes the pose extension’s known-size distance estimate possible. Heterogeneous gate sizes, non-square apertures, partially occluded gates or gates that must be distinguished by appearance rather than by acoustic association are outside the current scope.

L5. Acoustic sensing is an approximation, not a channel model. The beacon model of (5) applies independent Gaussian perturbations to range, bearing and elevation with Bernoulli dropout and a range cut-off. It does not model multipath, ray bending, Doppler on the acoustic carrier, frequency-dependent absorption or the correlated errors these produce. The inter-vehicle channel of (25) is likewise an acoustic-inspired transport with range-dependent latency and independent packet loss; its loss probability, range, latency and freshness window were not independently swept, and no claim about acoustic communication performance is made.

L6. The current sweep is narrow. Quantitative current results are restricted to two named profiles on one circuit (Horseshoe Bay). The strong profile turned out to be a saturation condition rather than a graded difficulty, so the informative disturbance evidence rests effectively on a single profile and five seeds. A finer sweep, and the same sweep on the other two circuits, would be needed to characterize the degradation curve rather than two points on it.

L7. Limited controller diversity in the systematic matrix. The rule-based evidence covers exactly two controllers that differ in one design decision. This is deliberate — it makes the comparison a clean ablation — but the 78-run matrix has not yet been exercised against structurally different autonomy strategies, such as a model-predictive controller, a cascaded design with an inner DVL velocity loop, or a planner that reasons about several gates ahead. The separate Gen-2 recurrent-PPO validation demonstrates interface compatibility, but it does not provide a factorial comparison across controller families, circuits and disturbances. We therefore cannot claim that the difficulty ordering across the three circuits is a property of the circuits rather than of the evaluated controllers.

L8. Limited fleet scale and diversity. Fleet evidence covers two homogeneous vehicles at a 90 s gap and three heterogeneous vehicles at two gaps, on one clean circuit, with three or five seeds. Larger fleets, fleets under current or with obstacles, formation-keeping objectives, cooperative overtaking, dynamic role assignment and competitive multi-team racing are all unaddressed. The coordination result in particular rests on three seeds per condition.

L9. The pose-aware perception extension is not validated. As reported in Section 5.3, the exploratory pose front end recovers a pose in only 11.7% of frames on a dedicated oblique-view rig, with a median plane-yaw error of 48.2° and orientation-class and sign accuracies of 0.00. It is accurate in the near-frontal regime the controller actually operates in (Table 6) and unreliable outside it. It is implemented and released, but it supports no claim in this paper and should not be used as a pose estimator.

L10. The learned-control evidence is a limited validation snapshot. The final learned result is a single selected Gen-2 recurrent-PPO checkpoint using its 27-dimensional onboard-derived observation encoding. It was evaluated in nine valid current-free episodes, three on each official circuit, and completed all nine with one collision event. This supports the claim that a recurrent learned controller can be integrated and evaluated through the common interface; it does not establish superiority, robustness under currents, fleet-level learning, sample efficiency or transfer to unseen circuits. The validation sample is small, the checkpoint was selected from a campaign with prior diagnostic access to the official circuits, and the learning protocol is not factorially matched to the 78-run rule-controller matrix. Broader comparisons of recurrent architectures, training procedures, disturbances and sealed circuits remain open.

L11. The official circuits are no longer unseen. An exploratory holdout evaluation opened Horseshoe Bay, Vertical Serpent and Mixed Endurance, and its access ledger records all sixty accesses. They must therefore not be treated as held-out circuits for any policy trained after that point, and every learned-controller result on them in this paper is labelled a retrospective diagnostic benchmark. Future generalization claims require a sealed circuit set that has not been opened.

L12. Three circuits do not span underwater autonomy. The official track set was designed to vary length, verticality and sequencing, and it does so — 12 to 22 gates, 93.8 to 206.3 m, planar to serpentine. It nonetheless covers a narrow slice of the underwater autonomy problem. Docking, manipulation, inspection, long-horizon navigation with drift, adversarial or dynamic environments and mission-level planning are all outside the benchmark’s task definition, and results here should not be read as evidence about them.

13. Reproducibility

Every number in this paper is traceable to a released artifact through a recorded chain of identifiers. This section documents that chain so a reader can verify a value, re-derive a table, or run an equivalent experiment.

13.1. Repository and Environment

The implementation, the track configurations, the controllers, the evaluation harness and the aggregated result artifacts are available in a public repository [25]. The benchmark environment is HoloOcean 2.3.0 with its Ocean world, Python 3.9.25 and the pinned dependency set in `requirements.txt`; the learning extension adds a separate pinned set in `requirements-rl.txt` (Gymnasium 1.0.0, PyTorch 2.8.0, Stable-Baselines3 2.7.1) that is deliberately installed into a distinct environment so the benchmark environment used for the frozen matrix remains unchanged. Runs reported here were executed on Windows 10 build 26200.

13.2. Configuration-Driven Experiments

A race is fully specified by one JSON file (Section 4): world bounds and environment, the ordered gate sequence with centres, passage directions and aperture sizes, the beacon network and its noise parameters, the current profiles, the obstacle generation rules, the participants with their vehicles, sensors, controllers and release delays, and the referee’s validation margins, penalties and scoring rules. A single entry point runs any of them, and a controller is loaded by alias, by dotted module path or by file path plus class name without modifying the package (Section 6). Reproducing a reported condition therefore requires the track file, the controller identifier and the seed, all of which are recorded per run.

13.3. Seeds and Determinism

Beacon packet randomness is seeded by the run seed, the transmitter, the receiver and the transmission index, so reception is independent of participant count and of the order in which controllers are polled; a fleet run is reproducible for this reason. Inter-vehicle packet loss and procedural course generation draw from seeded generators. For the learning extension, seed bands are registered by role — training, validation, test and final holdout — in a seed registry committed to the repository, and each evaluation records both requested and completed seeds. The matched benchmark of Section 10.6 evaluated every controller on the same seed set and the same generated geometries, which is what makes its paired statistics valid.

13.4. Manifests, Fingerprints and Hashes

Three levels of provenance are recorded.

Per-run. Each run emits a line-delimited event log and a machine-readable summary (R6), plus metadata giving the track, controller, seed, command line, requested and actual adapter, fallback-disabled status, control timestep, requested and measured current vector, and the observation and action contract versions.

Per-experiment. The 78-run matrix carries a complete experiment manifest recording the experiment-generation commit (df098ef4...), a source-tree fingerprint over the package’s Python and JSON files (e7d31077...), the HoloOcean version, the Python version, the operating system, the fallback status and the simulated current-coupling status for every run. It also asserts matrix completeness: 78 of 78 expected runs, no missing and no unexpected metadata.

Per-artifact. Learned-controller evaluations additionally record the model path and its SHA-256, the track file SHA-256 and the campaign commit. The final Gen-2 campaign records the observation contract, checkpoint selection decision, actual timestep, adapter status and per-episode result files; the selected checkpoint SHA-256 is stored beside the validation manifest.

13.5. Regenerating Tables and Figures

All result tables and figures in this paper are *post-processing* of existing artifacts: no command listed below launches the simulator or modifies a raw artifact.

```
# Result tables from the frozen 78-run matrix (also runs the penalty-identity check)
python article/regenerate_tables.py --check

# Track layouts and the controller-comparison plot, from track JSON and artifacts
python article/figures/generate_figures.py

# Journal figures: perception panels from committed artifacts
python article_journal/scripts/make_perception_figure.py

# Verify the compact immutable 78-run release package
python article_journal/scripts/package_78_matrix.py verify

# Manuscript
cd article_journal && latexmk -pdf -interaction=nonstopmode -halt-on-error main.tex
```

Running `regenerate_tables.py` against the frozen matrix reproduces the clean-track, current, fleet and coordination values of Tables 11, 12 and 13 exactly, and also verifies the penalty identity $T^{\text{pen}} = T^{\text{official}} + P$ on every finished run. The two journal figure scripts write a provenance JSON alongside each output recording their input files and asserting that no simulator was launched.

13.6. Re-Running Experiments

Reproducing a run rather than a table requires HoloOcean and the recorded seed. Every experimental command is stored verbatim in the per-run metadata; a clean single-vehicle run has the form

```
python -m marine_race_arena.scripts.run_marine_race \
--track marine_race_arena/tracks/marine_race_horseshoe_bay.json \
--benchmark-task clean_gate \
--controller rule_gate_center_then_commit \
--adapter holoocean --seed 0 --dt 0.033 \
--duration 560.0 --current-profile none --official
```

and the current-free demonstration of Section 10.2 is driven by committed scripts that pass `-current-profile none` and assert `currents_actual = [0, 0, 0]` in the manifest before any episode runs. Two conventions are worth restating for anyone attempting a comparison. First, the engine-free fallback adapter exists for deterministic unit testing and must be disabled for any reported result; the manifest field `adapter_actual` records which adapter actually ran. Second, the control timestep must match the condition being reproduced: $\Delta t = 0.033$ s for the 78-run matrix and $\Delta t = 0.1$ s for the current-free demonstration and final learned-controller validation. These rates define different experimental conditions; their timings must not be pooled.

13.7. Updating the Learning Results

The learning section is deliberately structured so that its validation can be refreshed without changing the benchmark protocol. A new frozen checkpoint is added by recording its SHA-256 and selection decision, running the same three seeds per official circuit, and regenerating Table 14 from the per-episode result files. The methodology of Section 9—the 27-D observation encoding, recurrent architecture, reward and evaluation separation—is independent of any one checkpoint.

14. Conclusion and Future Work

This paper presented Marine Race Arena, a configurable and reproducible benchmark and evaluation framework for onboard-only autonomous underwater gate racing. Its central contribution is a controller-agnostic observation-action interface combined with an explicit race manager and an enforced separation between autonomy and evaluation: controllers use only allow-listed onboard sensing and messages delivered through simulated channels, while an independent referee uses privileged state to adjudicate ordered crossings, penalties and completion.

The experiments show that the common contract is usable across heterogeneous circuits, currents, fleets and controller types. The reference controller completes the current-free demonstration, while the systematic matrix exposes condition-dependent controller robustness and coordination failures that clean single-vehicle scores would miss. The final Gen-2 recurrent-PPO controller also completes all nine valid current-free validation episodes using its 27-dimensional onboard-derived observation encoding and the common four-action interface, with one collision event and no out-of-bounds events. This validates the integration of a learned controller into the benchmark and referee, without turning the result into a claim of superiority or unseen-circuit generalization.

The evidence remains simulation-based and bounded by one backend, a narrow current sweep, limited controller diversity in the systematic matrix, a small fleet and an exploratory pose extension that is not used by the main controllers. The learned validation is based on one selected recurrent-PPO checkpoint, three episodes per circuit and circuits that are not sealed against prior diagnostic access. Future work should extend the disturbance and fleet axes, compare broader sequence-aware learning strategies under matched protocols, seal new evaluation circuits and transfer the contract to a physical BlueROV2. In each case, Marine Race Arena provides the instrument for making the comparison reproducible and meaningful.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

The benchmark implementation, the track configurations, the evaluation harness and the aggregated result artifacts are publicly available at <https://github.com/AndreaBedei/HoloDroneCompetition>. Section 13 describes the artifact layout and regeneration commands. A compact, machine-readable release of the 78-run matrix is included under `article_journal/scientific_release/matrix_78_20260715`; its deterministic checker is run with `python article_journal/scripts/package_78_matrix.py verify`. No archival DOI has been assigned yet (TODO: deposit a release snapshot and add the DOI here).

References

- [1] J. Betz, H. Zheng, A. Liniger, U. Rosolia, P. Karle, M. Behl, V. Krovi, R. Mangharam, Autonomous vehicles on the edge: A survey on autonomous vehicle racing, *IEEE Open Journal of Intelligent Transportation Systems* 3 (2022) 458–488. doi:10.1109/OJITS.2022.3181510.
- [2] M. O’Kelly, H. Zheng, D. Karthik, R. Mangharam, F1TENTH: An open-source evaluation environment for continuous control and reinforcement learning, in: *Proceedings of the NeurIPS 2019 Competition and Demonstration Track*, Vol. 123 of *Proceedings of Machine Learning Research*, 2020, pp. 77–89. URL <https://proceedings.mlr.press/v123/o-kelly20a.html>
- [3] S. Hoffmann, S. Sagmeister, T. Betz, J. Bongard, S. Büttner, D. Ebner, D. Esser, G. Jank, S. Goblirsch, A. Langmann, M. Leitenstern, L. Ögretmen, P. Pitschi, A.-K. Schwehn, C. Schröder, M. Weinmann, F. Werner, B. Lohmann, J. Betz, M. Lienkamp, Head-to-head autonomous racing at the limits of handling in the A2RL challenge, *arXiv preprint arXiv:2602.08571* Abu Dhabi Autonomous Racing League; preprint, not peer reviewed (2026). arXiv:2602.08571. URL <https://arxiv.org/abs/2602.08571>
- [4] P. Foehn, D. Brescianini, E. Kaufmann, T. Cieslewski, M. Gehrig, M. Muglikar, D. Scaramuzza, AlphaPilot: Autonomous drone racing, in: *Robotics: Science and Systems (RSS)*, 2020. doi:10.15607/RSS.2020.XVI.081. URL <https://www.roboticsproceedings.org/rss16/p081.html>
- [5] E. Kaufmann, L. Bauersfeld, A. Loquercio, M. Müller, V. Koltun, D. Scaramuzza, Champion-level drone racing using deep reinforcement learning, *Nature* 620 (7976) (2023) 982–987. doi:10.1038/s41586-023-06419-4.
- [6] A. Manzanilla, S. Reyes, M. A. Garcia, D. A. Mercado, R. Lozano, Autonomous navigation for unmanned underwater vehicles: Real-time experiments using computer vision, *IEEE Robotics and Automation Letters* 4 (2) (2019) 1351–1356. doi:10.1109/LRA.2019.2895272.
- [7] M. Stojanovic, J. Preisig, Underwater acoustic communication channels: Propagation models and statistical characterization, *IEEE Communications Magazine* 47 (1) (2009) 84–89. doi:10.1109/MCOM.2009.4752682.
- [8] M. M. M. Manhães, S. A. Scherer, M. Voss, L. R. Douat, T. Rauschenbach, UUV simulator: A Gazebo-based package for underwater intervention and multi-robot simulation, in: *OCEANS 2016 MTS/IEEE Monterey*, 2016, pp. 1–8. doi:10.1109/OCEANS.2016.7761080.
- [9] E. Potokar, S. Ashford, M. Kaess, J. G. Mangelson, HoloOcean: An underwater robotics simulator, in: *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, 2022, pp. 3040–3046. doi:10.1109/ICRA46639.2022.9812353. URL <https://www.ri.cmu.edu/publications/holocean-an-underwater-robotics-simulator/>
- [10] S. Chu, Z. Huang, Y. Li, M. Lin, D. Li, I. Carlucho, Y. R. Petillot, C. Yang, MarineGym: A high-performance reinforcement learning platform for underwater robotics, in: *2025 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2025, pp. 17146–17153. doi:10.1109/IROS60139.2025.11246146. URL <https://ieeexplore.ieee.org/document/11246146>

- [11] M. von Benzon, F. F. Sørensen, E. Uth, J. Jouffroy, J. Liniger, S. Pedersen, An open-source benchmark simulator: Control of a BlueROV2 underwater robot, *Journal of Marine Science and Engineering* 10 (12) (2022) 1898. doi:10.3390/jmse10121898.
- [12] P. Cieślak, Stonefish: An advanced open-source simulation tool designed for marine robotics, with a ROS interface, in: *OCEANS 2019 - Marseille*, 2019, pp. 1–6. doi:10.1109/OCEANSE.2019.8867434.
- [13] M. M. Zhang, W.-S. Choi, J. Herman, D. Davis, C. Vogt, M. McCarrin, Y. Vijay, D. Dutia, W. Lew, S. Peters, B. Bingham, DAVE aquatic virtual environment: Toward a general underwater robotics simulator, in: *Proceedings of the IEEE/OES Autonomous Underwater Vehicles Symposium (AUV)*, 2022, pp. 1–8. doi:10.1109/AUV53081.2022.9965808.
URL <https://arxiv.org/abs/2209.02862>
- [14] E. Potokar, S. Ashford, M. Kaess, J. G. Mangelson, HoloOcean: A full-featured marine robotics simulator for perception and autonomy, *IEEE Journal of Oceanic Engineering* 49 (4) (2024) 1322–1336. doi:10.1109/JOE.2024.3410290.
- [15] A. Amer, O. Álvarez-Tuñón, H. I. Ugurlu, J. L. F. Sejersen, Y. Brodskiy, E. Kayacan, UNav-Sim: A visually realistic underwater robotics simulator and synthetic data-generation framework, in: *2023 21st International Conference on Advanced Robotics (ICAR)*, IEEE, 2023, pp. 570–576. doi:10.1109/ICAR58858.2023.10406819.
- [16] Z. Huang, M. Buchholz, M. Grimaldi, H. Yu, I. Carlucho, Y. R. Petillot, URoBench: Comparative analyses of underwater robotics simulators from reinforcement learning perspective, in: *OCEANS 2024 - Singapore*, IEEE, 2024. doi:10.1109/OCEANS51537.2024.10682284.
- [17] A. Dosovitskiy, G. Ros, F. Codevilla, A. Lopez, V. Koltun, CARLA: An open urban driving simulator, in: *Proceedings of the Conference on Robot Learning (CoRL)*, Vol. 78 of *Proceedings of Machine Learning Research*, 2017, pp. 1–16.
URL <https://proceedings.mlr.press/v78/dosovitskiy17a.html>
- [18] CARLA Autonomous Driving Leaderboard, CARLA Autonomous Driving Leaderboard: Evaluation criteria 2.1, accessed: 2026-09-10 (2026).
URL https://leaderboard.carla.org/evaluation_v2_1/
- [19] M. Franchi, et al., Underwater robotics competitions: The european robotics league emergency robots experience with feelhippo auv, *Frontiers in Robotics and AI* 7 (2020) 3. doi:10.3389/frobt.2020.00003.
- [20] K. Harada, R. Fukuda, Y. Mizoguchi, Y. Yamamoto, K. Mishima, Y. Tanaka, Y. Nishida, K. Ishii, Autonomous underwater vehicle with vision-based navigation system for underwater robot competition, in: *Proceedings of the International Conference on Artificial Life and Robotics (ICAROB2022)*, Vol. 27, 2022, pp. 354–359. doi:10.5954/ICAROB.2022.OS28-2.
URL https://kyutech.repo.nii.ac.jp/record/7489/files/ICAROB2022_OS28-2.pdf
- [21] M. Xanthidis, M. Kalaitzakis, N. Karapetyan, J. Johnson, N. Vitzilaos, J. M. O’Kane, I. Rekleitis, AquaVis: A perception-aware autonomous navigation framework for underwater vehicles, in: *2021 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2021. doi:10.1109/IROS51168.2021.9636124.
- [22] Y. Yang, Y. Xiao, T. Li, A survey of autonomous underwater vehicle formation: Performance, formation control, and communication capability, *IEEE Communications Surveys & Tutorials* 23 (2) (2021) 815–841. doi:10.1109/COMST.2021.3059998.
- [23] T. Yan, Z. Xu, S. A. Gadsden, Formation control of multiple autonomous underwater vehicles: A review, *Intelligence & Robotics* 3 (1) (2023) 1–22. doi:10.20517/ir.2023.01.
- [24] G. Brockman, V. Cheung, L. Pettersson, J. Schneider, J. Schulman, J. Tang, W. Zaremba, OpenAI gym, *arXiv preprint arXiv:1606.01540* (2016). arXiv:1606.01540.
URL <https://arxiv.org/abs/1606.01540>

- [25] A. Bedei, HoloDroneCompetition: Marine Race Arena repository, <https://github.com/AndreaBedei1/HoloDroneCompetition> (2026).