

Marine Race Arena: A Configurable HoloOcean Benchmark for Underwater Gate Racing and Team-Level Fleet Evaluation

Andrea Bedei^a, Lorenzo Bacchiani^a, Giovanni Pau^a, Roberto Girau^a

^a*Department of Computer Science and Engineering, University of Bologna, Italy*

Abstract

We present Marine Race Arena, a configurable and reproducible framework for developing and evaluating autonomous systems in underwater gate-racing tasks. Complete race scenarios can be specified through declarative configurations, including tracks, vehicles, onboard sensors, acoustic beacons, environmental currents, obstacles, starting conditions, race rules and penalties. A central design principle is the strict separation between autonomy and evaluation: controllers operate exclusively from onboard observations, acoustic measurements and optional inter-vehicle communication, while an independent referee accesses privileged simulator state only to validate ordered gate crossings, enforce the rules and compute the official results. This common interface enables different autonomy strategies to be evaluated under the same sensing and scoring conditions and supports both single-vehicle trials and cooperative multi-vehicle scenarios. The reference implementation includes onboard visual-acoustic gate perception, two interpretable rule-based controllers, leader-follower coordination mechanisms and reinforcement-learning controllers integrated through the same participant-level information boundary and action interface. Three ready-to-use race circuits with different geometric characteristics are provided together with structured logging and automated validation tools for reproducible experimentation. Experiments conducted in HoloOcean with a BlueROV2-class vehicle assess course completion, gate progression, timing, collisions, robustness to environmental currents and fleet coordination. The results show that Marine Race Arena can expose meaningful differences between autonomy strategies and operating conditions while maintaining a common evaluation protocol, providing a unified testbed for studying underwater perception, navigation, control, learning and cooperation under repeatable and progressively challenging racing scenarios.

Keywords: underwater robotics, autonomous racing, robotics benchmark, reproducible evaluation, onboard perception, multi-robot systems, reinforcement learning, HoloOcean

1. Introduction

Autonomous racing has emerged as an effective research format for evaluating perception, planning, control and real-time decision-making within a single quantitative task under repeatable conditions [1]. In ground robotics, F1TENTH provides standardized vehicles, circuits, metrics and baselines for continuous control and reinforcement learning [2], while full-scale competitions such as the Abu Dhabi Autonomous Racing League (A2RL) extend the same principle to head-to-head racing near the limits of vehicle handling [3]. In aerial robotics, gate-based competitions have driven the development of high-speed perception and control pipelines [4], and learning-based systems have more recently demonstrated autonomous drone-racing performance at champion level [5]. Across these domains, the value of racing is not limited to speed: a fixed task, explicit course geometry, common rules and measurable outcomes provide a practical basis for comparing, reproducing and improving complete autonomous systems. Underwater robotics provides a natural but substantially different setting for this type of evaluation. Satellite-based global positioning is unavailable underwater, inertial and Doppler-based navigation accumulates uncertainty over time, and visual navigation is affected by limited visibility, illumination variations and water conditions [6]. Acoustic sensing and communication can extend the observable range but introduce limited bandwidth, propagation delay and measurement uncertainty [7], while persistent currents and strongly damped vehicle dynamics further affect how a vehicle approaches, aligns with and crosses a gate. Underwater racing is therefore not simply a ground or aerial racing task transferred to a marine simulator.




Figure 1: Underwater gate racing in Marine Race Arena. A BlueROV2-class vehicle approaches the ordered, beacon-marked gate line of an official circuit in HoloOcean. The controller sees this scene only through its forward camera, its depth, inertial and Doppler sensors and the acoustic packets it physically receives; the gate positions visible to the reader are never available to it.

A growing ecosystem of underwater simulators already provides the physical and sensor models required to develop autonomous systems. UUV Simulator introduced a Gazebo-based environment for underwater vehicle dynamics and multi-robot experimentation [8], while HoloOcean provides Unreal Engine environments, simulated underwater sensing and multi-agent support [9]. More recently, MarineGym has emphasized high-throughput simulation for learning-based underwater control [10]. These platforms provide increasingly capable simulation backends, but a simulator alone does not define a racing benchmark. It does not necessarily specify the ordered task to be completed, which information a controller is allowed to use, how gate crossings are validated, how penalties and timeouts are applied, or how different participants are evaluated under a common set of rules. This distinction is particularly important in simulation, where global vehicle pose, exact gate coordinates, the true current field or official course progress can simplify navigation substantially despite being unavailable to a deployed underwater vehicle. At the same time, relying only on the controller’s internal estimate to determine task completion can hide perception, localization or sequencing failures. A meaningful benchmark should therefore separate the information available to autonomy from the information required for evaluation: controllers should operate from observations representative of onboard sensing, whereas privileged simulator state should remain confined to an independent evaluation mechanism.

This paper presents Marine Race Arena, a configurable HoloOcean benchmark for autonomous underwater gate racing and team-level fleet evaluation (Fig. 1). Marine Race Arena is implemented as a benchmark and race-management layer above the simulation backend. Complete race scenarios are defined through declarative configurations specifying the environment, ordered gates, vehicles, sensor profiles, acoustic beacons, environmental currents, obstacles, participants, starting conditions, timing rules and penalties. The reference implementation and the experiments reported in this work use HoloOcean with a BlueROV2-class vehicle [11], while simulator-specific sensing and actuation are isolated behind an adapter layer. A central feature of the framework is its controller-agnostic autonomy interface: the benchmark does not prescribe how a vehicle should navigate the course, but only defines the observations available to the participant and the commands it must provide. A controller can therefore be rule-based, optimization-based, learning-based or follow another autonomy paradigm without requiring changes to the race runner or evaluation logic. In parallel, a configurable automated referee uses privileged simulator state exclusively to validate ordered gate crossings, detect collisions and other rule violations, apply penalties and timeouts, determine race termination and compute the official results. Referee-side state and official course progress are never returned

to the controller. The same architecture extends to multi-vehicle experiments, where each vehicle retains independent autonomy and local state while the framework provides simultaneous or staggered starts, optional inter-vehicle communication, proximity monitoring, distributed leader–follower coordination and team-level scoring.

The release includes three ready-to-use gate-racing circuits with different geometric characteristics, together with configurable current profiles, obstacles and participant configurations. The common controller interface is demonstrated with two interpretable rule-based controllers and a recurrent reinforcement-learning controller, showing that substantially different autonomy approaches can be integrated without modifying the race-management or scoring logic. Structured event logs and machine-readable race summaries record both participant execution and official referee outcomes. The experimental evaluation exercises the framework across different circuits, environmental conditions and multi-vehicle configurations, with the purpose of assessing whether heterogeneous autonomy strategies and race scenarios can be executed and compared within the same participant-level information boundary, action interface and independent evaluation mechanism.

The main contributions of this work are as follows:

1. **A configurable benchmark for autonomous underwater gate racing.** Marine Race Arena defines complete race scenarios through declarative configurations covering tracks, ordered gates, vehicles, onboard sensors, acoustic beacons, environmental currents, obstacles, participants, starting conditions, timing, penalties and scoring. Three ready-to-use circuits are provided, while the same configuration model supports custom racing scenarios.
2. **A controller-agnostic autonomy interface.** Controllers interact with the benchmark through a common participant-level information boundary and action interface and can implement rule-based, optimization-based, learning-based or other autonomy strategies without modifying the race execution or evaluation logic. The release provides rule-based and reinforcement-learning controllers as reference implementations of this interface.
3. **A configurable race manager with an independent automated referee.** Race conditions and rules are defined independently from the controller, while privileged simulator state is restricted to the evaluation layer. The referee validates ordered gate crossings, detects violations, applies penalties and timeouts and computes official race results without exposing privileged progress information to autonomy.
4. **Single-vehicle and team-level multi-vehicle evaluation within the same framework.** Marine Race Arena supports independent per-vehicle autonomy, staggered or simultaneous starts, optional acoustic-inspired communication, proximity monitoring, distributed leader–follower coordination and team-level scoring while maintaining independent official progress for each participant.
5. **An experimental validation of the common benchmark interface.** The framework is exercised in HoloOcean across different race circuits, current conditions and multi-vehicle configurations, including both rule-based reference controllers and a reinforcement-learning controller integrated through the same participant-level information boundary and action interface. Structured outputs and automated validation procedures support repeatable evaluation and extension with new controllers and racing scenarios.

The remainder of the paper is organized as follows. Section 2 reviews underwater simulation, autonomous-racing benchmarks, underwater perception, multi-vehicle cooperation and learning-based marine control. Sections 3 and 4 introduce the benchmark model, runtime architecture, HoloOcean environment and observation interface. Section 5 describes onboard gate perception, followed by the reference controllers in Section 6 and the independent referee in Section 7. Section 8 presents multi-vehicle execution and coordination, while Section 9 describes the integration of learning-based control. Section 10 reports the experimental evaluation. The final sections discuss the results, limitations and reproducibility before concluding the paper.

2. Related Work

Underwater robotics has benefited from increasingly capable simulation environments. UUV Simulator extends Gazebo with underwater vehicle dynamics, sensors and multi-robot support [8]. Stonefish provides a marine-oriented simulator with hydrodynamics, underwater rendering and ROS integration [12], while DAVE combines underwater vehicle simulation with optical and acoustic sensing [13]. HoloOcean uses Unreal Engine to provide visually rich

environments, onboard sensor models and multi-agent execution [14]. UNav-Sim focuses on vision-based underwater navigation in photorealistic environments [15], whereas MarineGym targets high-throughput reinforcement-learning experiments using GPU-accelerated underwater simulation [10]. These platforms differ in physical fidelity, sensing, throughput and intended tasks, but primarily provide the simulation substrate on which autonomy is developed. UroBench addresses reproducible comparison at a different level by evaluating underwater simulators through common navigation and learning tasks [16]. Marine Race Arena is complementary to both directions: the underlying simulator determines how vehicles, sensors and the environment evolve, while the benchmark defines the race, the participant interface and how the result is officially evaluated.

Autonomous racing provides the closest methodological reference. F1TENTH combines a standardized scaled vehicle with racing environments, interchangeable controllers and support for both conventional and learning-based methods [2]. Its simulator can also execute multiple vehicles in the same environment. At full scale, A2RL extends autonomous racing to head-to-head operation near the handling limits of the vehicle [3]. In aerial robotics, AlphaPilot established ordered gate racing as a system-level task for perception, planning and control [4], while Swift demonstrated the potential of deep reinforcement learning in high-performance drone racing [5]. CARLA offers a related example from autonomous driving: the simulator supports multiple simultaneously active vehicles [17], and its Autonomous Driving Leaderboard defines a participant interface, restricts access to privileged simulator information and evaluates route completion and infractions outside the submitted agent [18]. Although these systems address different domains and tasks, they show the value of separating the autonomy method from a common task definition and evaluation procedure.

Underwater gate traversal has also been investigated directly, although mainly as part of specific missions rather than as a reusable racing protocol. FeelHippo demonstrates underwater gate navigation through integrated perception and control [19], while Harada et al. consider vision-based navigation for an underwater Gate Pass task [20]. The broader underwater perception problem is shaped by the absence of satellite positioning, uncertain localization and limited optical visibility. Vision-based localization and control have therefore been studied on real underwater vehicles [6], and perception-aware navigation has been used to explicitly account for the visibility of relevant targets during motion [21]. Acoustic sensing provides complementary range and bearing information but introduces limited bandwidth, latency, multipath and time-varying uncertainty [7]. In Marine Race Arena these sensing problems remain part of the participant autonomy stack: visual detections and acoustic measurements may be combined by the controller, but privileged gate geometry is not used to resolve the target.

The same distinction becomes important in multi-vehicle operation. Cooperation between AUVs depends strongly on communication availability and relative-state estimation [22], while formation-control literature highlights the close coupling between localization, communication and coordinated motion [23]. Marine Race Arena does not propose a new formation or acoustic communication method. Instead, it provides the race-level mechanisms needed to test such strategies under a common protocol. Each vehicle retains its own controller and local state, optional information is exchanged through an acoustic-inspired communication channel and the independent race manager tracks per-vehicle progress, proximity events and team-level results. The included leader-follower strategy is used only as a reference implementation for exercising these capabilities.

Learning-based control provides a further case in which a common interface is useful. MarineGym exposes observation and action spaces designed for underwater reinforcement-learning tasks [10], while general environment interfaces such as OpenAI Gym demonstrate the broader value of separating an algorithm from the environment with which it interacts [24]. Marine Race Arena follows the same interface principle but is intended primarily for race execution and evaluation rather than high-throughput training. A participant is required only to consume the allowed observations and produce the expected vehicle command. Its internal implementation may therefore be rule-based, optimization-based, learned or hybrid. The recurrent reinforcement-learning controller used in this work serves as one example of this extensibility and is evaluated within the same participant-level information boundary and action interface, using the same referee as the rule-based controllers.

Table 1 summarizes the resulting positioning. The comparison considers only properties relevant to the benchmark contribution: an explicit racing task, a participant-facing controller contract, restriction of privileged simulator state, evaluation independent from the participant, multiple controllable vehicles in the same scenario and an aggregate team result. These criteria do not rank simulator quality, and multi-vehicle support alone does not imply team-level evaluation. Marine Race Arena combines these elements around a configurable underwater ordered-gate task, allowing heterogeneous controllers to be evaluated by the same configurable race manager and independent referee.

Table 1: Benchmark-level capabilities of representative systems.

System	Domain	Race task	Agent contract	Privileged-state restriction	Independent evaluator	Multi-vehicle	Team score
UUV Simulator [8]	Underwater	–	–	–	–	✓	–
Stonefish [12]	Underwater	–	–	–	–	✓	–
DAVE [13]	Underwater	–	–	–	–	✓	–
HoloOcean [14]	Underwater	–	–	–	–	✓	–
UNav-Sim [15]	Underwater	–	–	–	–	✓	–
URoBench [16]	Underwater	–	✓	–	–	–	–
MarineGym [10]	Underwater	–	✓	–	–	–	–
F1TENTH [2]	Ground	✓	✓	–	–	✓	–
CARLA [17] / Leaderboard [18]	Ground	–	✓	✓	✓	✓	–
Marine Race Arena (this work)	Underwater	✓	✓	✓	✓	✓	✓

✓: explicitly supported; –: not established as a benchmark-level property in the cited system.

3. Benchmark Model and System Architecture

Marine Race Arena combines four architectural elements: a declarative race configuration, a set of participants and their autonomy interfaces, a race-management and evaluation layer, and a simulator adapter. This separation defines the benchmark independently of both the controller implementation and the simulation backend. A race can therefore be reconfigured without changing the autonomy code, while a new controller or vehicle can be connected without changing the race definition or adjudication logic.

3.1. Benchmark Instance

Definition 1 (Benchmark instance). A Marine Race Arena instance is defined by the tuple

$$\mathcal{E} = (\mathcal{W}, \mathcal{G}, \mathcal{S}, \mathcal{F}, \mathcal{C}, \mathcal{O}, \mathcal{P}, \mathcal{R}), \quad (1)$$

where:

- $\mathcal{W}$ is the simulation world and navigable volume;
- $\mathcal{G} = (g_1, \dots, g_M)$ is the ordered gate sequence;
- $\mathcal{S}$ defines the starting conditions, including participant spawn and release settings;
- $\mathcal{F}$ defines the race finish condition;
- $\mathcal{C}$ represents the environmental current field;
- $\mathcal{O}$ is the set of obstacles in the scenario;
- $\mathcal{P} = \{1, \dots, N\}$ is the participant set, with each participant associated with a vehicle, sensor profile and controller;
- $\mathcal{R}$ contains the race and referee configuration, including validation criteria, penalties, timeouts and scoring rules.

The tuple is a conceptual description of the benchmark instance. Its concrete values are supplied by the declarative track and race configuration and are interpreted by the configuration loader, arena builder and simulator adapter at run time.

3.2. Architectural Requirements

The requirements in Table 2 express the properties that make the benchmark configurable, controller-agnostic and independently evaluable. They describe architectural constraints rather than a particular autonomy algorithm or vehicle implementation.

Table 2: Design requirements.

ID	Requirement
R1	The race is defined through a declarative configuration covering the scenario, gates, participants, environmental conditions and evaluation rules.
R2	Autonomy uses only the information exposed through the participant interface and receives no information produced by the official evaluation system.
R3	Progression, violations and the final result are determined independently of the controller implementation.
R4	Controllers based on different paradigms can be connected through the same participant/controller interface.
R5	Simulator- and vehicle-specific operations are isolated through an adapter.
R6	Every execution produces structured events and machine-readable results, including per-participant and, when applicable, team-level results.

3.3. Runtime Architecture

The complete execution flow is shown in Fig. 2. A race starts from the declarative configuration, which is checked by the configuration loader and used by the arena builder to instantiate the scenario through the selected simulator adapter. Once the arena has been created, the race runner becomes the central runtime component: it advances the simulation, coordinates the participants and mediates their interaction with the vehicle backend. At each control step, the runner obtains the information made available to each participant, passes it to the corresponding controller and forwards the returned vehicle command to the adapter for execution. This cycle is repeated throughout the race and applies independently to every participating vehicle.




Figure 2: Marine Race Arena runtime architecture.

At the same time, the referee receives the information required to determine official course progression, rule violations, penalties and completion. Its decisions are used for adjudication and reporting but are not returned to the controllers, so the evaluation path does not become part of the autonomy loop. The framework records the resulting events and per-participant results and, in multi-vehicle configurations, can additionally produce the configured team-level result. Fig. 3 provides the complementary conceptual view of this architecture, emphasizing the two information flows: one connects participant observations to vehicle commands, while the other supports race evaluation and scoring.

4. Vehicle, Simulator, Sensing and the Observation Contract

4.1. Simulator and Vehicle

The reference adapter used in the HoloOcean experiments reported here drives HoloOcean [9] with a BlueROV2-class vehicle [11]. The simulator advances at 30 physics ticks per second.




Figure 3: Controller and referee information flows.

Table 3: Onboard sensor suite.

Sensor	Rate	Output and configuration
Depth	30 Hz	Scalar depth (m); noise $\sigma = 0$.
IMU	30 Hz	Linear acceleration and angular rate, with bias terms.
DVL	15 Hz	Body-frame velocity (m/s); beam elevation 22.5° , max range 50 m.
Front camera	30 Hz	RGB image, 640×480 , 90° field of view.
Collision	optional	Boolean onboard contact flag when configured.

The experiments use two control rates over this common physics rate: 30 Hz, with one control step per physics tick, and 10 Hz. The evaluation protocol identifies the rate used by each campaign; the two rates are not pooled.

The portable participant command uses four normalized body-frame channels,

$$u_{i,t} = [u^{\text{surge}}, u^{\text{sway}}, u^{\text{heave}}, u^{\text{yaw}}]^T \in [-1, 1]^4. \quad (2)$$

The reference HoloOcean adapter maps this abstract command to the eight BlueROV2 thrusters by a fixed mixing matrix. The four vertical thrusters receive u^{heave} ; the four horizontal thrusters receive

$$\tau_{1-4} = s_\tau \begin{bmatrix} u^{\text{surge}} + u^{\text{sway}} + \kappa u^{\text{yaw}} \\ u^{\text{surge}} - u^{\text{sway}} - \kappa u^{\text{yaw}} \\ u^{\text{surge}} + u^{\text{sway}} - \kappa u^{\text{yaw}} \\ u^{\text{surge}} - u^{\text{sway}} + \kappa u^{\text{yaw}} \end{bmatrix}, \quad (3)$$

where $\kappa = 0.35$ is the yaw-mixing gain and s_τ the thruster scale; each component is clamped to the thruster limit. A direct-thruster interface remains an adapter-specific capability, but all results reported here use the abstract four-degree-of-freedom command.

4.2. Sensor Suite

Table 3 lists the onboard sensors available to a participant. The official profile combines a forward RGB camera, a depth sensor, an IMU and a DVL; an onboard contact flag is also permitted when configured.

Before specifying beacon placement, we define the logical frame associated with each gate. For gate g with passage direction n_g , let

$$\hat{n} = \frac{n_g}{\|n_g\|}, \quad \hat{r}_g = \frac{\hat{z} \times \hat{n}}{\|\hat{z} \times \hat{n}\|}, \quad \hat{s}_g = \hat{n} \times \hat{r}_g, \quad (4)$$

with world-up $\hat{z} = (0, 0, 1)$. When the passage direction is parallel to the vertical axis, $\hat{r}_g$ is defined with respect to the global $\hat{x}$ axis. This passage direction defines a right-handed gate frame, reused by both the beacon model and the referee crossing test.

4.3. Acoustic Beacon Model

Using the right-handed frame of (4), each gate carries one beacon, placed at $b_g = c_g + \delta_n \hat{n} + \delta_r \hat{r}_g + \delta_s \hat{s}_g$ relative to the gate centre c_g (default offset $(0, 0, 0.35)$ m along $\hat{s}_g$). For a receiver at position p_r with heading ψ , let $\delta = b_g - p_r$, let $\rho = \|\delta\|$ be the geometric range and let $\tilde{\rho}$ be its noisy delivered measurement. The beacon reports $\tilde{\rho}$, a yaw-relative bearing, an elevation and a signal strength:

$$\rho = \|\delta\|, \quad \tilde{\rho} = \max(0, \rho + \eta_\rho), \quad \eta_\rho \sim \mathcal{N}(0, \sigma_\rho^2), \quad (5)$$

$$\beta = \text{wrap}(\text{atan2}(\delta_y, \delta_x) - \psi) + \eta_\beta, \quad (6)$$

$$\varepsilon = \text{atan2}(\delta_z, \sqrt{\delta_x^2 + \delta_y^2}) + \eta_\varepsilon, \quad (7)$$

$$\text{SS} = \max\left(0, 1 - \frac{\tilde{\rho}}{\rho_{\max}}\right),$$

where $\rho_{\max}$ is the beacon range and $\text{wrap}(\cdot)$ maps to $(-180^\circ, 180^\circ]$. Bearing and elevation have independent angular perturbations $\eta_\beta, \eta_\varepsilon \sim \mathcal{N}(0, \sigma_\theta^2)$, while range has perturbation η_ρ with standard deviation σ_ρ . The signal strength uses the same noisy range. Packets are absent beyond $\rho_{\max}$ or when a Bernoulli dropout occurs. The official tracks use a common per-track value in $\{0.2, 0.45, 0.6\}$ for the angular and range noise standard deviations (degrees and metres, respectively), with dropout probabilities up to 0.04.

Every ordered gate has one independent periodic transmitter with a unique contiguous identifier $B01, \dots, BM$. All beacons transmit independently of referee state, and each vehicle receives every in-range, non-dropped packet. The received packets do not indicate which beacon is currently relevant; the controller maintains its expected beacon from its own local course state and selects the corresponding packet (Sections 5 and 6). Packet randomness is seeded by the run, transmitter, receiver and transmission index, keeping fleet reception reproducible and independent of participant count and call order.

4.4. Current Field Model

The current field is the sum of analytic primitives evaluated at vehicle position p and time t ,

$$\mathbf{v}_c(p, t) = \sum_j \mathbf{v}_j(p, t), \quad (8)$$

The available primitives are a constant flow, a localized jet, a sinusoidal component and a vortex. A localized jet centred at c_j with radius R_j contributes, for $\|p - c_j\| \leq R_j$,

$$\mathbf{v}_j = \mathbf{v}_0 \omega(d_j), \quad \omega(s) = \begin{cases} \max(0, 1 - \frac{s}{R_j}) & \text{linear,} \\ \exp(-\frac{s^2}{2\sigma_j^2}) & \text{Gaussian,} \end{cases} \quad (9)$$

where $d_j = \|p - c_j\|$ is the jet distance and $\sigma_j = \max(R_j/2, 10^{-6})$ the jet width; the contribution is zero outside R_j . The constant term contributes a fixed velocity $\mathbf{v}_0$, while a sinusoidal term modulates one axis, $v_{j,\text{axis}}(t) = v_{0,\text{axis}} + A \sin(2\pi f t + \phi)$. A *vortex* of radius R_j , tangential speed V_t and vertical speed V_z contributes, with $r = \sqrt{\Delta x^2 + \Delta y^2}$ and tangent $\hat{\theta} = (-\Delta y, \Delta x)/r$,

$$\mathbf{v}_j = (\hat{\theta}_x V_t \omega(r), \hat{\theta}_y V_t \omega(r), V_z \omega(r)). \quad (10)$$

The current magnitude and direction can be configured as part of the race scenario. The selected field is defined relative to the gate layout and is available on every official circuit. The adapter applies the field physically before each control step but does not report it directly to the controller, which must infer the disturbance from its onboard measurements, including body-frame DVL velocity.

Table 4: Official track configurations.

Track file	Track name	Gates	Laps	Total gates	Length (m)	Intended use
marine_race_horseshoe_bay.json	Marine Race Horseshoe Bay	12	1	12	93.8	Clean-gate horseshoe baseline
marine_race_vertical_serpent.json	Marine Race Vertical Serpent	17	1	17	118.3	Vertical and slalom gate sequencing
marine_race_mixed_endurance.json	Marine Race Mixed Endurance	22	1	22	206.3	Longer endurance and current-oriented layout

4.5. Declarative Track Configuration

A track is a self-contained JSON file. Its sections define the race timing, benchmark task, world, ordered gates, sensors, beacons, currents, obstacles, participants and referee. The selected task supplies the corresponding validation contract; Listing 1 shows a representative excerpt.

Listing 1: Representative excerpt of a track configuration.

```

"track": {
  "gate_inner_size_m": [1.5, 1.5],
  "gate_sequence": ["G01", "G02", "..."]
},
"gates": [
  { "id": "G01", "position": [4.0, -2.0, -3.5],
    "rotation_rpy_deg": [0.0, 0.0, 0.0],
    "passage_direction": [1.0, 0.0, 0.0] }
],
"beacon": {
  "range_m": 90.0,
  "angular_noise_std_deg": 0.2,
  "range_noise_std_m": 0.2,
  "dropout_probability": 0.0,
  "update_rate_hz": 10.0
},
"participants": [
  { "id": "bluerov2_01", "vehicle": "BlueROV2",
    "controller": "rule_gate_baseline",
    "start_delay_s": 0.0 }
],
"referee": {
  "penalties": { "minor_collision_s": 5.0 },
  "gate_validation": {
    "vehicle_clearance_margin_m": 0.1 }
}

```

The three official circuits are summarized in Table 4 and their layouts shown in Fig. 4. They are deliberately heterogeneous: Horseshoe Bay is a short, largely planar baseline; Vertical Serpent forces repeated vertical transitions and slalom sequencing; Mixed Endurance is more than twice as long as Horseshoe Bay and combines both. All three use $1.5 \text{ m} \times 1.5 \text{ m}$ apertures, and neither the aperture size nor the referee margin was changed at any point of this study.

The obstacle_gate task accepts explicit fixed obstacles or deterministic random boxes generated between consecutive gates with a configured density and minimum clearance. The HoloOcean adapter records the requested and physically spawned counts. Obstacle avoidance is supported by the task configuration but is not evaluated in the results reported here.

4.6. Official Observation Contract

At each control step, the participant receives an observation dictionary whose fields are listed in Table 5. It contains participant-local time, the configured onboard sensors, packets delivered by the acoustic beacon channel and, when




Figure 4: Official race-circuit layouts.

Table 5: Participant observation fields.

Field	Meaning
<code>local_time_s</code>	Elapsed time since this vehicle’s release; it is not official race time.
<code>sensors</code>	Filtered onboard dictionary containing only the configured camera, depth, IMU, DVL and optional collision measurements.
<code>beacons</code>	List of packets delivered by the simulated acoustic channel from independent transmitters; it may be empty or contain several identifiers, and no field marks which one is the next gate.
<code>comms</code>	Present only when the optional inter-vehicle acoustic channel is enabled: an <code>inbox</code> of messages received from teammates, each with the sender identifier, the controller-authored payload and a receiver-local reception timestamp (Section 8).

enabled, optional inter-vehicle communication. The observation builder is separated from the referee and uses only the participant state and allow-listed sensing interfaces; evaluation data therefore cannot enter the observation through this path. Communication payloads are authored by the corresponding controllers from their own legal information. World-frame velocity, pose-derived heading and depth, and environmental-current telemetry may exist internally for simulation or adjudication but are excluded from the official observation.

The absence of three classes of information defines the task. The controller never receives the global ground-truth vehicle pose; any localization estimate must therefore be derived from participant-accessible onboard information. It does not receive ground-truth gate positions or orientations from the simulator, so the spatial relation to the course must be inferred from onboard sensing and acoustic measurements. It also receives neither the referee’s notion of the next gate, nor whether a crossing was valid, nor the number of remaining gates; course progression is consequently a controller-side estimate that can be wrong.

5. Onboard Gate Perception and Target Association

A racing benchmark is only meaningful if target selection remains part of the participant’s autonomy problem. The observation contract of Section 4.6 does not expose ground-truth gate geometry or referee progress, while the forward camera may contain several gates and the acoustic channel may simultaneously deliver packets from several transmitters. The received packets do not indicate which beacon is currently relevant; that state is maintained locally by the controller. Selecting which visible gate corresponds to the locally expected beacon must therefore be performed onboard rather than by the benchmark.

The reference perception front end combines geometric detection in the forward camera with acoustic target association, as summarized in Fig. 5. The visual stage produces multiple candidate gates rather than a single preselected target, while the acoustic stage provides bearing and range information for the beacon expected by the controller. Their combination produces either one visual target or no visual target for the current frame.




Figure 5: Onboard gate-perception pipeline.

5.1. Gate Detection from the Forward Camera

The gates are rendered as square frames formed by light-coloured bars. Direct colour or brightness segmentation is unreliable because underwater rendering can substantially alter their apparent intensity relative to the water and seabed. The detector therefore relies primarily on geometric structure: edges are extracted from the grayscale image, closed contours are converted into bounding boxes, and candidates are filtered according to their apparent size, aspect ratio and consistency with a rectangular frame. These criteria are fixed for the evaluation and are applied unchanged across the three official circuits.

For an image of width W and height H , the centre of a surviving candidate is expressed in normalized image coordinates as

$$c_x = \frac{(x + \frac{1}{2}w_{px}) - \frac{1}{2}W}{\frac{1}{2}W}, \quad c_y = \frac{(y + \frac{1}{2}h_{py}) - \frac{1}{2}H}{\frac{1}{2}H}. \quad (11)$$

Here (c_x, c_y) is the displacement of the candidate centre from the optical axis, while its apparent size is represented by the image-area fraction α . Each candidate is also assigned a confidence γ that combines shape consistency, aspect-ratio plausibility, apparent size and image centredness. A single fixed parameter set is used for all official circuits; the exact thresholds and weights are provided with the released implementation.

The detector retains multiple plausible candidates and removes near-duplicate detections before passing the resulting ranked list to the association stage. Retaining several candidates is important in the racing setting because consecutive gates can be simultaneously visible; a visual detector that selected a single rectangle by itself would implicitly solve part of the course-association problem.

5.2. Acoustic–Visual Target Association

The controller maintains its own locally expected beacon (Section 6). When a packet from that beacon is available, it provides a noisy bearing β and measured range $\tilde{\rho}$ according to the acoustic model of Section 4. The perception problem is then to determine which, if any, of the visual candidates is geometrically consistent with that acoustic observation.

The conceptual association path is shown in Fig. 6.

Association first rejects candidates that are inconsistent with the camera field of view and the acoustic bearing. The horizontal image position of each candidate provides an approximate visual bearing, which can be compared with β . Apparent target size is also used as a consistency check at short range: a nearby expected gate should occupy a substantial portion of the image, whereas a small complete rectangle is more likely to belong to a farther gate.

When the acoustic bearing provides a clear directional constraint, candidates that strongly disagree with it are rejected. The remaining detections are ranked using visual confidence, image centredness, relative apparent size and acoustic–visual bearing consistency. Relative rather than purely absolute size is used so that a farther but visually centred gate does not systematically dominate a closer candidate. The same association parameters are used throughout the official circuits.

If no packet from the locally expected beacon is currently available, the association falls back to visual evidence and favours large, confident and well-centred candidates. If no candidate remains plausible, no visual target is returned and the controller continues using the other onboard measurements.




Figure 6: Conceptual onboard visual-acoustic perception path.

Table 6: Perception audit on the official circuits.

Circuit	Detection	4-corner	Pose avail.	Multi-candidate association
Horseshoe Bay	90.0%	81.7%	51.7%	7/7
Vertical Serpent	88.3%	78.3%	48.3%	4/4
Mixed Endurance	91.7%	83.3%	53.3%	11/11

This design preserves the benchmark information boundary. Target association depends only on the controller’s local course state, camera detections and received acoustic measurements. It does not use simulator gate poses, the global vehicle pose or the referee’s expected-gate index. An incorrect association is therefore an autonomy error rather than an error hidden by benchmark-side target selection.

5.3. Exploratory Pose-Aware Extension

The centre-based representation does not describe the orientation of the gate plane. We therefore also implemented a pose-aware variant that uses a detected aperture quadrilateral together with the known nominal gate aperture to estimate relative plane orientation and distance.

Its behaviour is strongly dependent on viewing geometry. In a dedicated set of oblique views, a valid pose estimate was obtained in only 11.7% of frames and the median plane-yaw error was 48.2° . The estimator is therefore unreliable under large oblique viewing angles, motivating the separate characterization under the viewing conditions encountered during normal racing.

5.4. Perception Audit on the Official Circuits

To characterize perception under the conditions encountered by the reference controller, we recorded 60 frames per circuit along its trajectories on the three official circuits. Ground truth was used only offline to evaluate the perception output and never entered the control loop. Table 6 summarizes detection availability, quadrilateral recovery, pose availability and multi-candidate target association, while Fig. 7 shows representative states.

Detection was available in approximately 88–92% of the recorded frames, while a validated four-corner aperture was recovered in approximately 78–83%. A metric pose estimate was available in roughly half of the frames, confirming that pose recovery is substantially less reliable than centre-based gate detection. The reference controllers are consequently designed to tolerate intermittent visual measurements rather than assume that a complete gate geometry is available at every control step.




Figure 7: Representative onboard gate-perception states.

Target association was considerably more stable. In every recorded frame containing multiple visual candidates, the online association selected the same gate as the offline reference: 7/7 cases on Horseshoe Bay, 4/4 on Vertical Serpent and 11/11 on Mixed Endurance. No unintended online target switch was observed. This supports the use of acoustic bearing primarily as an identity cue when several geometrically plausible gates are visible.

The pose results also clarify the difference between the dedicated oblique-view test and normal racing. During the reference trajectories, the controller tends to approach the gate approximately frontally. Under these conditions, the recovered plane orientation was stable, with median yaw error below 1.2° on all three circuits and no gross orientation or sign errors observed. Under strongly oblique views, by contrast, pose recovery became sparse and inaccurate. The relevant conclusion is therefore not that the pose estimator is generally reliable, but that its accuracy depends strongly on viewing geometry.

For the racing task considered here, reliable gate-centre detection and acoustic–visual association are sufficient for the reference controllers, while full pose recovery is substantially more sensitive to viewpoint. The perception audit characterizes these components independently from the course-completion results reported in Section 10.

6. Controller Interface and Rule-Based Reference Controllers

A central usability goal of Marine Race Arena is that an autonomy stack can be integrated through a small interface without modifying the runner, the referee or the adapter. This section specifies that interface, the dynamic loading mechanism, the controller-local course tracker and the two rule-based reference controllers whose comparison anchors the evaluation.

6.1. Controller Interface

A controller implements the participant policy

$$\pi_i : o_{i,t} \mapsto u_{i,t}, \quad (12)$$

mapping the official observation $o_{i,t}$ to the vehicle command $u_{i,t}$. The software contract that realizes this mapping consists of three methods: `reset(mission_info)`, called once before the run with a minimal static assignment; `step(observation)`, called at every control tick to receive the official observation of (12) and return the command of (2); and `close()`, called at teardown. The contract is duck-typed, so any object exposing the three callables is accepted and a participant need not depend on framework base classes. A manual controller may raise a dedicated stop exception from `step` to end a run cleanly; any other exception is recorded as a controller error and terminates the participant. Listing 2 shows a minimal example.

Listing 2: Minimal custom controller.

Table 7: High-level controller command fields of (2).

Field	Interpretation
surge	Longitudinal body-frame thrust; positive moves forward.
sway	Lateral body-frame thrust; sign follows the adapter mixing of (3).
heave	Vertical thrust; additionally constrained by depth-safety logic near the world bounds.
yaw	Yaw-rate command; sign follows the adapter mixing of (3).

```

class MyController:
    def reset(self, mission_info):
        self.tracker = LocalCourseTracker(
            initial_beacon_id=mission_info["initial_beacon_id"],
            total_beacons=mission_info["total_beacons"],
            laps=mission_info["laps"])

    def step(self, observation):
        sensors = observation["sensors"]
        state = self.tracker.update(
            local_time_s=observation["local_time_s"],
            beacons=observation["beacons"],
            camera_image=sensors.get("FrontCamera"),
            dvl_velocity=sensors.get("DVLSensor"))
        if state.finished:
            return {"surge": 0.0, "sway": 0.0, "heave": 0.0, "yaw": 0.0}
        # ... map state.beacon and camera to a command ...
        return {"surge": 0.3, "sway": 0.0, "heave": 0.0, "yaw": 0.0}

    def close(self):
        pass

```

The core `mission_info` passed to `reset` is deliberately minimal: the initial beacon identifier, the number of beacons and the lap count; fleet runs may additionally provide static team-assignment metadata. The high-level command fields are listed in Table 7; values are normalized to $[-1, 1]$ and clamped by the framework before the action mapping of (3). In fleet mode a controller may additionally return a small message payload alongside its command and read incoming teammate messages from the observation.

6.2. Dynamic Loading

A controller is selected at run time by a built-in alias, a dotted module path with an explicit class, or a file path and class name imported dynamically. The runner also provides manual controllers for inspection and data collection. An external policy is therefore evaluated without modifying the package:

```

python -m marine_race_arena.scripts.run_marine_race \
  --track ../marine_race_horseshoe_bay.json \
  --benchmark-task clean_gate \
  --controller path/to/my_controller.py \
  --controller-class MyController \
  --adapter holoocean --official --headless

```

The same mechanism loads the learned policies of Section 9: a thin wrapper exposes the three interface methods, encodes the official observation into the policy’s input vector and returns the policy’s output as the command of (2). No part of the runner, referee or adapter is aware that the controller is neural.




Figure 8: The LOCALCOURSETRACKER state machine shared by both reference controllers. ADVANCE loops back to SEARCH for the next gate and reaches FINISHED after the last locally confirmed passage. No transition consumes referee state.

6.3. Controller-Local Course Progression

Because the official observation does not contain referee progress (Section 4.6), each controller maintains its own estimate of which gate it is currently approaching. The reference controllers use the reusable LocalCourseTracker, initialized from the mission assignment and updated exclusively from participant-accessible onboard information.

The tracker maintains the locally expected beacon and progresses through the states

$$\text{SEARCH} \rightarrow \text{APPROACH} \rightarrow \text{VISUAL_ALIGN} \rightarrow \text{COMMIT} \rightarrow \text{VERIFY_EXIT} \rightarrow \text{ADVANCE},$$

before returning to SEARCH for the next gate or entering FINISHED after the final locally confirmed passage.

At each control step, the tracker selects the packet corresponding to its locally expected beacon and combines acoustic bearing and range with the associated visual target and body-frame DVL motion. A passage is not inferred from a single measurement. The vehicle first establishes a stable visual-acoustic alignment, then commits to the aperture, and finally checks whether its subsequent onboard measurements are consistent with having crossed and moved away from the expected gate. Forward DVL displacement, the evolution of beacon range and bearing, and the persistence or disappearance of the visual target jointly provide this evidence.

Temporary loss of a camera detection or an acoustic packet is therefore insufficient to advance the local course state, as is an isolated change in range without supporting motion evidence. When the available measurements do not consistently support a completed passage, the tracker remains on, or returns to, the current gate rather than advancing to the next one. This conservative progression logic keeps gate sequencing inside the participant while avoiding any dependence on referee feedback.

Fig. 8 summarizes this controller-local state machine. Controller-local progression and referee decisions are logged separately for evaluation and are never fed back into the controller.

The two controllers described below use the same tracker and the same onboard information. They are intended as interpretable reference baselines rather than optimized control policies.

6.4. Continuous Servo

Both reference controllers combine acoustic guidance, visual centring, depth hold and DVL-based passage verification through the shared local tracker of Fig. 8. They differ only in how visual feedback is used during the final gate passage.

Continuous Servo steers toward the locally expected beacon using its bearing and elevation, and refines that alignment with the associated visual target as the gate becomes visible. Depth feedback maintains the vertical reference,




Figure 9: Center-then-commit passage and local recovery.




Figure 10: LocalCourseTracker states and passage styles.

while DVL motion contributes to passage verification. During COMMIT and VERIFY_EXIT, the controller continues applying bounded visual corrections while moving through the aperture, continuously servoing toward the observed gate centre.

6.5. Center-then-Commit

Continuous visual correction is useful during approach but can become undesirable close to the gate, when parts of the frame leave the camera field of view and the detected geometry becomes partial. In this regime, the estimated image centre can shift even after the vehicle is correctly aligned, producing unnecessary late corrections when the available clearance is smallest. Center-then-Commit addresses this behaviour while retaining the same observations, acoustic-visual guidance, local tracker and passage-verification logic as Continuous Servo. The controller uses visual feedback normally during approach and alignment; once a stable alignment has been established, it stops reacting to subsequent image-centre changes during COMMIT and VERIFY_EXIT and continues through the aperture while maintaining depth and attitude from onboard feedback. The two reference controllers therefore differ only in the final traversal strategy, providing a matched comparison between continuous visual servoing and a stabilized committed passage. The corresponding state transitions are shown in Fig. 9, while Fig. 10 highlights the behavioural difference between the two strategies; their experimental comparison is reported in Section 10.3.

7. Independent Referee and the Information Boundary

The referee is the component that makes results comparable, and it is the component a participant must not be able to influence. This section formalizes the two-sided information boundary, the gate-crossing test, the referee state machine and arbitration order, the penalty model and the scoring rules, and it closes by treating the boundary as something to be measured rather than asserted.

7.1. The Two-Sided Boundary

The benchmark partitions the run state into two disjoint views. Writing x_t for the full simulator state at step t , the *controller view* of participant i is

$$o_{i,t} = \Phi_i(x_t) = (t_i^{\text{loc}}, \mathcal{Z}_{i,t}, \mathcal{A}_{i,t}, \mathcal{M}_{i,t}), \quad (13)$$

where t_i^{loc} is time since this participant's release, $\mathcal{Z}_{i,t}$ the allow-listed onboard measurements, $\mathcal{A}_{i,t}$ the set of acoustic packets delivered by the simulated channel at step t , and $\mathcal{M}_{i,t}$ the optional teammate inbox. The *evaluator view* is

$$e_t = \Psi(x_t) = (\{p_{i,t}, \dot{p}_{i,t}\}_{i \in \mathcal{P}}, \mathcal{G}, \mathcal{O}, \mathcal{B}, \mathbf{v}_c), \quad (14)$$

comprising exact vehicle poses and velocities, the ordered gate geometry, the obstacle set, the world bounds and the true current field. The evaluator view therefore contains information that is deliberately excluded from the participant interface. In particular, the official expected-gate index, completed-gate count, participant status and accrued penalties are not inputs to the observation map.

In the implementation this separation is structural rather than conventional: the observation builder's signature admits the track configuration, the arena, the adapter and the participant state, and does not admit a referee object. Referee-derived quantities are therefore unreachable from the code path that constructs $o_{i,t}$.

Information flows in one direction only. The referee reads x_t and the controller's issued commands; the controller reads neither e_t nor any function of it. In particular, a controller is never told that a crossing was accepted, never told which gate is expected next, and never told that it has been penalized. Where the controller's own estimate of its progress disagrees with the referee's, the disagreement is logged and left standing (Section 7.7).

7.2. Gate Frame and Crossing Test

The referee reuses the right-handed gate frame $(\hat{n}, \hat{r}_g, \hat{s}_g)$ of (4), defined from the passage direction n_g with the vertical-normal fallback to the world $\hat{x}$ axis. For consecutive referee-side positions p_{t-1}, p_t , the signed distances to the gate plane are

$$d_{t-1} = \hat{n}^\top(p_{t-1} - c_g), \quad d_t = \hat{n}^\top(p_t - c_g). \quad (15)$$

A candidate crossing requires a sign change in the *correct direction*,

$$d_{t-1} \leq 0 \leq d_t \wedge (p_t - p_{t-1})^\top \hat{n} > 0, \quad (16)$$

i.e. travel from the negative to the positive side along the normal. The intersection of the segment with the plane is

$$\lambda = \frac{-d_{t-1}}{d_t - d_{t-1}}, \quad q = p_{t-1} + \lambda(p_t - p_{t-1}), \quad \lambda \in [0, 1], \quad (17)$$

and the crossing is accepted only when q lies inside the aperture, shrunk by the configured safety margin m_g ,

$$|\hat{r}_g^\top(q - c_g)| \leq \frac{w_g}{2} - m_g, \quad |\hat{s}_g^\top(q - c_g)| \leq \frac{h_g}{2} - m_g. \quad (18)$$

Gates must be passed *in sequence*: the referee tests only the expected gate g_ℓ , advancing the index $\ell \leftarrow \ell + 1$ on a valid crossing and incrementing the lap when the sequence completes. Crossing the aperture of a *non-target* gate is a missed-gate event; a sign change in the wrong direction is logged but neither penalized nor counted as progress. Fig. 11 illustrates the test and Fig. 12 the principal automatic lifecycle.

Two design choices in this test are worth naming. Testing only the expected gate, rather than all gates, is what turns the course into an ordered task instead of a collection of independent hoops. Shrinking the aperture by m_g before the containment test makes the pass criterion strictly conservative: a trajectory that clips the frame is not credited.




Figure 11: Gate-crossing geometry and acceptance region.




Figure 12: Automatic participant lifecycle.

7.3. Referee State Machine and Arbitration

Each participant occupies one of the states NOT_STARTED, RUNNING, FINISHED, DNF, TIMEOUT, STUCK, DSQ, MANUAL_STOP or CONTROLLER_ERROR; all but the first two are terminal. A participant is released to RUNNING when the race clock reaches its start delay (Section 8). On each tick of a running participant, events are evaluated in a fixed priority order that defines the arbitration: (1) a controller error transitions to CONTROLLER_ERROR and returns; (2) out-of-bounds; (3) generic collision; (4) obstacle collision; (5) stuck accumulation; (6) duration timeout; (7) gate progression. Steps (2) to (5) accrue penalties without terminating; only (1), (6), a missed gate in (7) or an external gate-timeout produce a terminal state.

The order is not arbitrary. Placing controller errors first prevents a crashed participant from also accruing collision or stuck charges from its final frozen command. Placing gate progression last ensures that a tick in which the vehicle both crosses a gate and touches its frame is charged for the contact and credited for the crossing, rather than either event masking the other. To avoid double counting, an obstacle-collision tick suppresses the generic collision flag. Repeated collision, obstacle and out-of-bounds charges use independent cooldowns; inter-vehicle proximity uses its own team-level cooldown; stuck is latched once per episode; and terminal events are not cooldown-limited.

7.4. Penalties and Termination

Penalties accumulate into a per-participant total P_i ; the applied defaults are listed in Table 8. The non-trivial conditions are formalized as follows. The instantaneous speed is $s_t = \|p_t - p_{t-1}\|/\Delta t$, and a stuck accumulator

Table 8: Default referee events, penalties and terminal conditions.

Event	Penalty	Terminal
Gate, bar or generic collision	5.0 s each	no
Obstacle collision	5.0 s each [†]	no
Out of bounds	10.0 s each	no
Stuck (soft)	15.0 s per episode	no
Missed gate	N/A	yes (DNF)
Timeout	N/A	yes
Gate-timeout stuck	N/A	yes
Inter-vehicle proximity	0 or 5.0 s per pair	no (team-level)

[†] Per-obstacle value if specified; otherwise the default collision penalty is used.

integrates low-speed time,

$$T^{\text{stuck}} \leftarrow \begin{cases} T^{\text{stuck}} + \Delta t, & s_t < v_{\text{stuck}} \\ 0, & \text{otherwise,} \end{cases} \quad (19)$$

with a soft penalty charged once per continuous low-speed episode when $T^{\text{stuck}} \geq T_{\text{max}}^{\text{stuck}}$ and $v_{\text{stuck}} = 0.02$ m/s. The configured $T_{\text{max}}^{\text{stuck}}$ is 45 s on Horseshoe Bay and 65 s on Vertical Serpent and Mixed Endurance. An out-of-bounds event is $p_t \notin \mathcal{B}$. A duration timeout, disabled by default, fires when $t - t_{\text{green}} > T_{\text{max}}$.

A missed gate sets DNF when `missed_gate_dnf` is enabled, which is the default, but carries no time penalty. The combination is deliberate: making a disqualifying error terminal but free of accumulated time ensures that losing the course is never cheaper than completing a slow but legal lap, which would otherwise be an exploitable scoring artefact.

7.5. Single-Vehicle Score and Ranking

For a finished participant the official time T_i^{official} is measured between the first and last gate (or green-to-finish, according to the timing mode) and the score is the penalized time

$$T_i^{\text{pen}} = T_i^{\text{official}} + P_i. \quad (20)$$

Ranking places finished participants before unfinished ones and is defined by the ascending lexicographic key

$$\kappa_i = \begin{cases} (0, T_i^{\text{pen}}, 0, 0, 0, \text{id}_i), & f_i = 1, \\ (1, -G_i^{\text{done}}, n_i^{\text{coll}}, P_i, \text{dist}_i, \text{id}_i), & f_i = 0, \end{cases} \quad (21)$$

where $f_i = 1$ denotes a finished vehicle, G_i^{done} the number of completed gates, n_i^{coll} the collision count and dist_i the distance to the next gate; the participant identifier is the deterministic final tie-break. Finished vehicles are therefore ordered by penalized time, and unfinished vehicles by progress, then by fewer collisions, fewer penalties and proximity to the next gate. This ordering preserves the information contained in partial runs instead of discarding them.

7.6. Event Logging

Every run emits a line-delimited event log and a machine-readable summary (requirement R6). Events are timestamped and attributed to the corresponding participant, while the summary records the final status, official and penalized times, course progress and the main penalty-related quantities used for evaluation. In fleet mode it additionally contains the team-level summary of Section 8. The experimental results of Section 10 are obtained by aggregating these outputs.

Table 9: Controller-local and referee course progression.

Circuit	Gate advancements			Mismatches		Local delay (s)	
	referee	local	matched	false	missed	median	p95
Horseshoe Bay	1120	1120	1116	10	4	0.924	1.353
Vertical Serpent	159	160	159	1	0	0.627	1.155
Mixed Endurance	195	192	192	0	3	0.462	1.287
All	1474	1472	1467	11	7	0.858	1.353

False denotes a local advancement not matched by the corresponding referee crossing; missed denotes a referee crossing without the corresponding local advancement.

7.7. Measuring the Boundary

Because the controller and the referee maintain independent estimates of course progress, their agreement can be evaluated offline without affecting control. The framework records both quantities separately, allowing discrepancies between local progression and official referee decisions to be inspected after each run.

Table 9 summarizes this comparison. The results support two observations. First, the information boundary is effective: controller-local progress is genuinely independent of the referee and occasional disagreements remain visible rather than being corrected by the framework. Second, this separation does not prevent practical course tracking: the onboard-only local estimate agrees with the official gate sequence in more than 99% of recorded advancements.

8. Fleet Execution, Coordination and Team-Level Evaluation

Fleet mode runs several vehicles as a single cooperative team through the same circuit, so the benchmark measures how quickly and how safely the team completes the course rather than pitting the vehicles against each other. The benchmark models one team: all vehicles share a single simulated world, each controller keeps an independent local course estimate, and the referee keeps one scoring state per vehicle.

8.1. Staggered Release and Spawn Geometry

Each participant i is assigned a release delay $\tau_i = (i - 1)\Delta\tau$ for a start gap $\Delta\tau$, so vehicles enter the course at $0, \Delta\tau, 2\Delta\tau, \dots$. A waiting vehicle receives a zero command, its controller `step` is *not* called, and neither stuck nor gate-timeout timers accumulate. At release the referee logs a `participant_released` event and resets the vehicle’s progress timers, so a delayed start is never misclassified as stuck or timed out. Spawn poses are distributed laterally around the base spawn to reduce launch interference while remaining inside the configured world bounds.

8.2. Inter-Vehicle Proximity Diagnostics

Inter-vehicle proximity events are detected on the referee side and are never exposed to controllers, which follows directly from this separation: a vehicle that could read the referee’s proximity flag would be receiving privileged information about a teammate’s exact position. For two running vehicles a, b the detector applies a separable cylinder test

$$\sqrt{(x_a - x_b)^2 + (y_a - y_b)^2} \leq r_{xy} \wedge |z_a - z_b| \leq r_z, \quad (22)$$

with hysteresis — a pair is released only when the horizontal distance exceeds a release threshold r_{rel} or the vertical distance exceeds r_z — and a refractory cooldown to debounce sustained proximity. A vehicle that wants to avoid its teammates must therefore infer their presence from its own sensing and from messages they choose to send.

8.3. Inter-Vehicle Acoustic Communication

Because the fleet is a single cooperative team, Marine Race Arena provides an optional channel so that vehicles can share information while running the course. A controller transmits by returning a small message payload alongside its command and receives an inbox of teammates' messages in its observation; what a message contains and how it is used are entirely the participant's responsibility, so the framework provides only the transport.

Rather than assuming an ideal communication link, the channel introduces acoustic-inspired range, propagation-delay, packet-loss and bandwidth constraints. A message broadcast by vehicle i reaches a teammate j only if the two are within a maximum range $R_{\max}$, after a range-dependent delay

$$\tau_{ij} = \frac{d_{ij}}{c} + \tau_{\text{proc}}, \quad (23)$$

where d_{ij} is the inter-vehicle distance, c the speed of sound and τ_{proc} a fixed processing delay. Each delivery is dropped with probability p_{loss} . Payloads are capped to a small size to model the low bandwidth, and a vehicle may not transmit faster than a configured per-sender interval. Packet loss is drawn from a seeded generator, so a run remains reproducible.

The channel is intended as a reproducible coordination substrate rather than a validated acoustic-channel model. The communication parameters are configurable, but communication performance itself is not evaluated as a separate contribution of this work.

The resulting multi-vehicle information flow and independent team referee are shown in Fig. 13.

Multi-Vehicle Fleet Architecture in Marine Race Arena




Figure 13: Conceptual fleet architecture and team referee.

8.4. Distributed Leader-Follower Coordination

Motivation. A staggered team of *identical* controllers tends to stay separated in time, because each vehicle takes roughly as long as the one ahead. A *heterogeneous* team need not: a faster follower released behind a slower leader can close the gap, and the two can then converge on the same gate. The resulting contacts emerge from the interaction of the convoy and do not necessarily indicate a problem with either controller in isolation.

Module design. We provide a leader–follower coordination controller based on locally estimated course progress and participant-authored teammate messages. It *wraps* rather than replaces a gate-passing controller that exposes a `LocalCourseTracker`, which keeps coordination and gate navigation separate concerns; both official camera-assisted controllers satisfy this interface. Each vehicle is assigned a predecessor from the static release order carried in its legal mission information; the frontmost vehicle has no predecessor and races freely. Each vehicle periodically broadcasts a compact message derived from its own local course tracker. The receiver retains the most recent report from its predecessor and treats it as an estimate of that vehicle’s local progress. Stale predecessor reports are ignored.

Let $\hat{g}_i$ denote locally estimated course progression. A follower yields while its predecessor is fewer than Δg locally estimated gates ahead,

$$\text{yield}_j \iff \hat{g}_{\text{pred}(j)} - \hat{g}_j < \Delta g, \quad (24)$$

While yielding, the wrapper still updates the base tracker from onboard observations but commands zero motion. Yielding is based only on fresh predecessor reports; otherwise the wrapped base controller runs unchanged.

Information boundary. All quantities used for coordination are participant-accessible: the vehicle’s own local tracker, a static predecessor assignment and controller-authored teammate messages. Referee progress may disagree with either local estimate but never enters the coordination decision. While yielding, the wrapper continues updating the local tracker but commands zero motion; any resulting race events or penalties are assessed independently by the referee.

8.5. Team-Level Score

Only the `team_summary` is the official fleet result; per-vehicle rows are diagnostics. For a team $\mathcal{P}$ of N vehicles with G_{exp} expected gates each,

$$G_{\text{team}}^{\text{exp}} = N G_{\text{exp}}, \quad G_{\text{team}}^{\text{done}} = \sum_{i \in \mathcal{P}} G_i^{\text{done}}. \quad (25)$$

The team finishes only if every vehicle finishes,

$$\text{finished}_{\text{team}} = \bigwedge_{i \in \mathcal{P}} \text{finished}_i, \quad (26)$$

in which case the team elapsed time spans the first release to the last finish,

$$T_{\text{team}} = \max_i T_i^{\text{finish}} - \min_i T_i^{\text{release}}, \quad (27)$$

and the penalized team score aggregates every per-vehicle penalty together with the team-level inter-vehicle term,

$$T_{\text{team}}^{\text{pen}} = T_{\text{team}} + \sum_{i \in \mathcal{P}} P_i + P_{\text{ivc}}, \quad (28)$$

where P_i is the per-vehicle penalty total of (20) and P_{ivc} counts each inter-vehicle proximity event once per pair rather than once per vehicle. The conjunction in (26) is what makes the score a *team* score: a fast leader cannot compensate for a follower that never finishes, so a coordination policy is rewarded for bringing the whole convoy home. Fig. 14 illustrates the aggregation.

9. Learning-Based Controller

A learning-based controller is included to demonstrate that policies obtained through reinforcement learning can be integrated into Marine Race Arena without modifying the race-management or evaluation framework. The controller receives the same participant-side information available during normal benchmark operation and is evaluated by the independent referee described in Section 7. Rather than treating learning as a separate benchmark mode, the policy is exposed through the same controller interface used by the reference approaches.

Figure 15 summarizes the resulting control loop and the progression adopted during training.




Figure 14: Team-level scoring for a staggered two-rover fleet.

PPO control loop in Marine Race Arena




Figure 15: Recurrent learning-based controller and training progression.

9.1. Observation Representation and Recurrent Policy

The policy receives a 27-dimensional observation vector built exclusively from participant-accessible measurements. The representation combines acoustic information associated with the locally tracked beacon, visual information about the corresponding gate target, vehicle motion measurements and short-term temporal information. It includes beacon availability, acoustic bearing, elevation and range, visual target centre and apparent size, depth, body-frame DVL velocity, yaw rate, the previous control action, short-term variations of the visual and acoustic observations, an indicator of recent target transitions and, when available, a local estimate of gate orientation. Availability indicators are used to distinguish valid measurements from missing ones, while unavailable quantities are represented through neutral values. Global vehicle pose, exact gate coordinates, future course geometry and referee-derived progress are not included.

The controller is trained with Proximal Policy Optimization (PPO) [25] using a recurrent actor-critic policy. The recurrent component is implemented with LSTM memory, so the decision at the current control step depends not only on the present observation but also on information retained from the recent observation history. This is useful in gate racing because the problem is only partially observable from a single frame: visual detections may be incomplete or temporarily disappear close to the aperture, acoustic measurements evolve continuously during the approach and the

same instantaneous observation may correspond to different phases of a passage depending on the vehicle’s recent motion.

At each step, the current observation is first mapped to an internal feature representation and then processed together with the recurrent state maintained by the LSTM. The resulting temporal representation is shared by the actor–critic architecture: the actor produces the continuous control action, while the critic estimates the state value used during PPO optimization. The recurrent state is carried forward between consecutive control steps and reset at the beginning of a new episode. In this way, temporal context is learned directly by the policy instead of being supplied through an explicit handcrafted navigation state machine.

9.2. Training Strategy

Training is organized progressively so that the policy first acquires the local behaviours needed for gate traversal before being required to solve complete racing sequences. Initial stages emphasize approaching the target, maintaining useful alignment and crossing the aperture. The training task is then extended to longer gate sequences and finally to complete circuits, while keeping the same observation representation and control interface.

The reward follows the same progression. Successful gate crossings and course completion provide the principal positive task signals, while additional terms encourage progress and useful alignment during the approach. Elapsed time, collisions and terminal failures are penalized. These terms are used only to shape the learning process and are not included in the observation seen by the controller.

PPO training produces a sequence of policy checkpoints as optimization progresses. A single policy is selected and frozen before the final evaluation; the controller is then executed without additional learning or online adaptation. The experimental results of this fixed policy are reported together with the other benchmark evaluations in Section 10.6.

10. Experimental Evaluation

This evaluation asks whether the benchmark, the observation boundary, the referee and the team scoring represent and distinguish autonomous behaviours under common experimental conditions. We organize the experiments around five research questions.

RQ1 Can Marine Race Arena support fully onboard autonomous completion across heterogeneous underwater racing circuits?

RQ2 Does the benchmark expose meaningful differences between autonomous control strategies?

RQ3 How do environmental currents affect benchmark performance?

RQ4 Can the same evaluation framework support multi-vehicle and team-level racing?

RQ5 Can learning-based controllers be integrated and evaluated within the same participant-level information boundary and referee procedure?

10.1. Protocol and Metrics

The systematic benchmark evaluation comprises both reference controllers on the three clean circuits, the medium current condition on Horseshoe Bay, homogeneous two-vehicle fleets and three-vehicle coordination trials. These experiments run at 30 Hz. The current-free completion demonstration comprises nine runs covering the three official circuits, without current, at 10 Hz. The recurrent PPO validation uses three episodes per circuit at 10 Hz.

All experiments use the native HoloOcean backend with a BlueROV2-class vehicle and participant-accessible onboard observations. The reported metrics are completion, completed gates, official and penalized time, collisions, out-of-bounds, wrong-direction and stuck events; fleet trials additionally report inter-vehicle proximity and team time.

Table 10: Current-free completion on the official circuits.

Circuit	Gates	Runs	Completed	Mean gates	Official time (s)	Coll.	OOB / WD
Horseshoe Bay	12	3	3/3	12.0/12	236.0 $\pm$ 6.9	0	0 / 0
Vertical Serpent	17	3	3/3	17.0/17	308.6 $\pm$ 6.5	0	0 / 0
Mixed Endurance	22	3	3/3	22.0/22	480.4 $\pm$ 15.7	0	0 / 0
Total	51	9	9/9	51.0/51	—	0	0 / 0

Times are mean $\pm$ sample standard deviation; Coll. denotes collision events, and OOB / WD denotes out-of-bounds and wrong-direction events.

Table 11: Controller comparison on the three official circuits under clean water.

Track	Controller	Seeds	Finished	Gates $\uparrow$	Official time (s) $\downarrow$	Collisions $\downarrow$
Horseshoe Bay	Continuous Servo	5	5/5	12.0/12	194.5 $\pm$ 4.8	0.0
	Center-then-Commit	5	5/5	12.0/12	197.7 $\pm$ 7.1	0.0
Vertical Serpent	Continuous Servo	5	5/5	17.0/17	236.9 $\pm$ 6.4	0.0
	Center-then-Commit	5	4/5	14.8/17	241.1 $\pm$ 1.8	0.0
Mixed Endurance	Continuous Servo	5	2/5	17.0/22	394.7 $\pm$ 3.9	0.0
	Center-then-Commit	5	5/5	22.0/22	410.7 $\pm$ 8.8	0.0

Collisions combine gate and world events; bold marks the better completion result where the controllers differ.

10.2. RQ1: Onboard-Only Completion Across Heterogeneous Circuits

The first question is whether the task is solvable under its own restrictions. RQ1 therefore tests end-to-end completion using only onboard sensing.

Table 10 answers this affirmatively for the current-free setting. The Center-then-Commit controller completed all three official circuits in every one of nine runs: 3/3 on Horseshoe Bay (12 gates), 3/3 on Vertical Serpent (17 gates) and 3/3 on Mixed Endurance (22 gates), for 51 of 51 gates in total. Across those nine runs the referee recorded zero gate, bar or world collisions, zero out-of-bounds events and zero wrong-direction crossings, and penalized time equalled official time in every run because no penalty was ever charged.

Completion time scales with course length roughly as expected from the track geometry: 236.0 $\pm$ 6.9 s over 93.8 m, 308.6 $\pm$ 6.5 s over 118.3 m and 480.4 $\pm$ 15.7 s over 206.3 m, i.e. approximately 0.40, 0.38 and 0.43 m/s of course progress. The consistency of that rate across three structurally different circuits — planar, vertically serpentine and long mixed — indicates that the limiting factor is the controller’s gate-by-gate acquire–align–commit cycle rather than any circuit-specific difficulty.

These results establish that the task is achievable end to end under the full participant-level observation boundary on three structurally different circuits.

10.3. RQ2: Do Control Strategies Separate?

A benchmark that ranks every reasonable controller identically is not measuring anything. The two reference controllers of Section 6 are an unusually clean test of this, because they differ in exactly one respect — what happens during COMMIT and VERIFY_EXIT — while sharing observations, guidance, tracker and confirmation logic.

Table 11 and Fig. 16 show that the answer depends on the circuit, and that neither controller dominates. On the short, largely planar Horseshoe Bay circuit the two are effectively tied: both complete all 12 gates in all five seeds with no collision, out-of-bounds or stuck event, and the staged controller is 3.2 s slower (197.7 $\pm$ 7.1 s versus 194.5 $\pm$ 4.8 s), a 1.7% penalty with no measured safety benefit. Constant re-centring is cheap here, and paying nothing for it is the right trade.

The controllers separate on the harder circuits, and they separate in opposite directions. On Vertical Serpent, Continuous Servo is more reliable, finishing 5/5 against 4/5 and completing 17.0 against 14.8 mean gates. On Mixed Endurance, which is more than twice as long, the ordering reverses far more sharply: Center-then-Commit finishes all five seeds and completes every gate (22.0/22), while Continuous Servo finishes only 2/5 and averages 17.0/22.




Figure 16: Reference-controller completion on clean and medium-current conditions.

Table 12: Controller performance on Horseshoe Bay with and without current.

Profile	Controller	Seeds	Finished	Gates $\uparrow$	Official (s)	Penalized (s)	Collisions $\downarrow$
Clean	Continuous Servo	5	5/5	12.0/12	194.5 ± 4.8	194.5 ± 4.8	0.0
	Center-then-Commit	5	5/5	12.0/12	197.7 ± 7.1	197.7 ± 7.1	0.0
medium	Continuous Servo	5	2/5	8.4/12	337.9 ± 19.2	692.9 ± 75.7	43.4
	Center-then-Commit	5	3/5	8.8/12	217.3 ± 1.5	223.9 ± 1.8	25.2

Times are mean $\pm$ sample standard deviation; collisions combine gate and world events.

The mechanism is consistent with the design argument of Section 6.5. Continuous visual correction is harmless when a passage is short and the approach is well conditioned, and it becomes a liability when partial near-field contours make the image centroid move quickly at exactly the moment the aperture margin is smallest. A long course compounds this: each passage is an independent opportunity for a late correction to fail, so a small per-gate risk becomes a large per-course risk. That Continuous Servo nonetheless wins on Vertical Serpent shows the trade is not one-directional — repeated vertical transitions appear to benefit from continuous re-centring more than they suffer from late corrections — which is precisely the kind of condition-dependent result a single-circuit benchmark would have concealed.

10.4. RQ3: Behaviour Under Environmental Currents

Table 12 shows the effect of the medium current condition on Horseshoe Bay. Under this condition, Continuous Servo finishes 2/5 runs, averages 8.4/12 completed gates and 43.4 gate and world collisions per run, with official and penalized times of 337.9 ± 19.2 s and 692.9 ± 75.7 s. Center-then-Commit finishes 3/5, averages 8.8/12 gates and 25.2 collisions, with official and penalized times of 217.3 ± 1.5 s and 223.9 ± 1.8 s.

The headline finding for benchmark design is the gap between Table 11 and Table 12. Both controllers score a perfect 5/5 with zero events on clean Horseshoe Bay, and drop to 2/5 and 3/5 with tens of collisions on the same circuit under a moderate current. Clean-track success is therefore not a proxy for robustness, and any evaluation protocol that reports only clean completion will rank controllers on a dimension that disappears the moment the water moves. The medium-current condition thus provides a meaningful robustness test for the benchmark.

10.5. RQ4: Multi-Vehicle and Team-Level Evaluation

The upper block of Table 13 validates that team execution works. With either controller, a homogeneous two-vehicle fleet on clean Horseshoe Bay at a 90 s release gap finishes 24/24 team gates in all five seeds, with no gate or world collision, no inter-vehicle proximity event, no out-of-bounds and no stuck event. Mean team elapsed and penalized times coincide at 303.1 ± 3.2 s for Continuous Servo and 308.0 ± 5.8 s for Center-then-Commit. This is deliberately a low-contact case: it exercises multi-vehicle spawning, staggered release, independent per-vehicle

Table 13: Multi-vehicle evaluation on clean Horseshoe Bay.

Setting	Policy	Seeds	Team finish	Team gates $\uparrow$	Team time (s) $\downarrow$	GW / IV / S $\downarrow$
<i>Two vehicles, homogeneous, gap 90 s</i>						
	Continuous Servo	5	5/5	24.0/24	303.1 ± 3.2	0.0 / 0.0 / 0.0
	Center-then-Commit	5	5/5	24.0/24	308.0 ± 5.8	0.0 / 0.0 / 0.0
<i>Three vehicles, heterogeneous convoy, gap 0 s</i>						
	No coordination	3	3/3	36.0/36	219.7 ± 2.5	9.3 / 2.3 / 0.0
	LF(1)	3	3/3	36.0/36	262.3 ± 7.7	0.0 / 0.0 / 0.0
	LF(2)	3	3/3	36.0/36	288.3 ± 1.3	0.0 / 0.0 / 1.0
<i>Three vehicles, heterogeneous convoy, gap 8 s</i>						
	No coordination	3	2/3	34.3/36	254.0 ± 24.5	127.3 / 2.0 / 0.0
	LF(1)	3	3/3	36.0/36	266.0 ± 8.3	0.0 / 0.0 / 0.0
	LF(2)	3	3/3	36.0/36	285.7 ± 3.8	0.0 / 0.0 / 0.0

GW / IV / S denotes gate and world collisions, inter-vehicle proximity and stuck events. LF(Δg) is the leader-follower policy of (24).

Table 14: Recurrent PPO validation results.

Circuit	Gates	Episodes	Finished	Mean time (s)	Collisions	OOB
Horseshoe Bay	12	3	3/3	192.5	0	0
Vertical Serpent	17	3	3/3	239.1	1	0
Mixed Endurance	22	3	3/3	365.6	0	0
All circuits	—	9	9/9	—	1	0

referee state, ranking and team aggregation without the vehicles interacting, since the per-vehicle behaviour matches the single-vehicle case.

The interacting case is the interesting one. A heterogeneous three-vehicle convoy — a Continuous Servo leader followed by two Center-then-Commit vehicles — converges on the same apertures, and the consequences are severe. At an 8 s gap the uncoordinated convoy finishes only 2/3 runs, averages 34.3/36 team gates, 127.3 gate and world collisions and 2.0 inter-vehicle proximity events per run. At simultaneous release it still finishes 3/3, but averages 9.3 collisions and 2.3 proximity events.

The distributed leader-follower policy of (24) removes this entirely. At both gaps, LF(1) finishes 3/3 with 36.0/36 team gates and zero gate, world, proximity, out-of-bounds and stuck events, at 266.0 ± 8.3 s (gap 8 s) and 262.3 ± 7.7 s (gap 0 s). The more conservative LF(2) also eliminates contact but is consistently slower (285.7 ± 3.8 s and 288.3 ± 1.3 s) and, at simultaneous release, its longer hold trips the independent stuck timer once per run, raising its penalized team time to 303.3 ± 1.3 s. A larger nominal gate separation therefore bought no additional safety in these conditions while costing both time and a penalty, which is why LF(1) is the recommended default and LF(2) is retained only as a conservative comparison.

Two properties of this result matter beyond the numbers. The coordination policy consumes only the vehicle’s own tracker state, a static predecessor assignment and teammate-authored messages delivered over a channel with range-dependent latency and loss; it never reads referee progress or a teammate’s true position, so the improvement is achieved inside the same information boundary that constrains everything else. And the cost of coordination is visible in the score rather than hidden: LF(2)’s stuck penalty is charged by the same independent referee that charges a collision, so a policy that buys safety by waiting pays for the waiting.

10.6. RQ5: Learning-Based Control Under the Same Evaluation Protocol

The recurrent PPO controller was evaluated in three episodes on each of the three official circuits using the same participant-level information boundary and independent referee employed throughout the benchmark. The policy completed all nine evaluation episodes: 3/3 on Horseshoe Bay, 3/3 on Vertical Serpent and 3/3 on Mixed Endurance.

No out-of-bounds event was recorded in any of the nine runs. Horseshoe Bay and Mixed Endurance were completed without collisions, while one collision event was recorded on Vertical Serpent. These results show that the

learned controller can compose local gate-navigation behaviour into complete ordered circuits while remaining entirely within the participant-side observation contract.

The sample is intentionally small, with three episodes per circuit, and is reported descriptively. Together, the rule-based and PPO evaluations illustrate that the benchmark interface is controller-agnostic: autonomy methods with different internal designs can be evaluated through the same participant-level information boundary, required body-frame command interface and independent referee procedure. The same interface can accommodate rule-based control, reinforcement learning, model-predictive control, optimization-based methods and other compatible controllers.

11. Discussion

The experiments support a view of underwater gate racing as a benchmark for ordered-sequence autonomy rather than speed alone. A vehicle must commit to the correct target, reacquire a new target after each passage and compose those transitions over a long course. Small per-transition defects therefore amplify with course length, while a single-gate task can make a weak sequence policy look successful.

11.1. What the Evaluation Contract Reveals

The central value of Marine Race Arena is the conjunction of a fixed information boundary, an ordered referee, configurable race management and machine-readable scoring. The boundary makes controller comparisons meaningful because rule-based and learned controllers operate within the same participant-level information boundary. The referee makes local progress auditable: controller estimates can be compared with privileged adjudication without allowing the latter into the control loop. The resulting residual disagreements are useful evidence, not hidden implementation details.

The configurable race manager makes this separation operational. Track geometry, participants, environmental conditions, starts, timing and penalties can be changed through scenario configuration while the participant interface and referee logic remain fixed. This permits the benchmark to vary the task and the operating condition independently of the autonomy implementation, which is important when interpreting differences between controllers.

The referee also makes failure interpretable. A missed next gate is recorded as a sequence failure rather than being blurred into a low aggregate gate count, and safety costs remain visible alongside progress. The medium-current condition also shows why completion, progress and safety events should be interpreted jointly: both controllers complete fewer runs and accumulate substantially more collisions than under clean conditions. Reporting these quantities together makes the effect of environmental disturbance visible beyond completion rate alone.

11.2. Reference Controllers, Currents and Fleets

The reference-controller comparison is a controlled ablation of visual feedback during passage. Neither controller dominates: continuous correction is useful on repeated vertical transitions, whereas suspending late visual corrections is more robust on the long mixed course and under the medium-current condition. This condition-dependent result shows that the benchmark can expose a design decision that a single clean circuit would conceal.

The fleet experiments extend the same logic from individual to team autonomy. Vehicles that are individually competent can still converge on the same aperture and collide when their interaction is not scored. A distributed leader-follower policy removes those events within the legal observation boundary, while the referee charges waiting through the stuck penalty. Thus coordination and its cost are evaluated in the same score rather than in a separate engineering trace.

11.3. Learning as a Demonstration of Generality

The learning experiment demonstrates the generality of this design. A recurrent PPO controller can be integrated within the same participant-level information boundary, body-frame command interface and independent referee procedure as the reference controllers, and can be assessed on complete ordered circuits under the reported current-free protocol. Because the interface constrains the legal information and action contract rather than the internal controller design, the same benchmark can accommodate reinforcement learning, rule-based control, model-predictive control, optimization-based methods and other compatible autonomy strategies.

12. Limitations

The present evaluation has several limitations that define the scope of the reported results and motivate future extensions of the benchmark.

L1. Simulation-only validation. All experiments reported in this work were conducted in simulation. No physical BlueROV2 experiments were performed, and the reported completion rates, times and collision counts should therefore be interpreted as simulation results rather than estimates of real-world vehicle performance.

L2. Simulator fidelity. The benchmark inherits the hydrodynamic, rendering and sensor models of its simulation back-end. HoloOcean provides the vehicle and sensing environment used in this study, while the environmental current field is represented through analytic primitives rather than computational fluid dynamics. Absolute vehicle performance and disturbance thresholds therefore depend on the fidelity of the underlying simulation.

L3. Optical perception assumptions. The reference gate detector was developed for the visual conditions and gate geometry used in the official circuits. More challenging optical effects such as strong turbidity, marine snow, severe backscatter, biofouling or very low-light conditions are not considered in the present evaluation and may require a different or more robust perception front end.

L4. Acoustic sensing abstraction. The acoustic beacon and inter-vehicle communication models are simplified and acoustic-inspired rather than high-fidelity channel simulations. Effects such as multipath propagation, ray bending, carrier Doppler and frequency-dependent absorption are not explicitly modelled.

L5. Limited current evaluation. Quantitative current results are restricted to one medium-current condition on Horse-shoe Bay. A broader range of current magnitudes and the same evaluation on the other circuits would provide a more complete characterization of disturbance robustness.

L6. Limited fleet scale. The fleet experiments consider two- and three-vehicle configurations on a single clean circuit. Larger fleets, operation under environmental disturbances and more complex cooperative or competitive behaviours remain to be evaluated.

L7. Baseline learning controller. The recurrent PPO controller is intended primarily as a reference demonstration of learning-based integration rather than as an optimized reinforcement-learning solution. More advanced algorithms, alternative recurrent or attention-based architectures, larger training budgets and different curriculum or reward designs could improve learned-control performance and provide a broader assessment of the benchmark for robot learning.

13. Reproducibility

Every number in this paper is traceable to a released artifact through a recorded chain of identifiers. This section documents that chain so a reader can verify a value, re-derive a table, or run an equivalent experiment.

13.1. Repository and Environment

The implementation, the track configurations, the controllers, the evaluation harness and the aggregated result artifacts are available in a public repository [26]. The benchmark environment is HoloOcean 2.3.0 with its Ocean world, Python 3.9.25 and the pinned dependency set in `requirements.txt`; the learning extension uses a separate pinned set in `requirements-rl.txt` installed in a distinct environment. Runs reported here were executed on Windows 10 build 26200.

13.2. Configuration-Driven Experiments

A race is fully specified by one JSON file (Section 4): world bounds and environment, the ordered gate sequence with centres, passage directions and aperture sizes, the beacon network and its noise parameters, the current profiles, the obstacle generation rules, the participants with their vehicles, sensors, controllers and release delays, and the referee’s validation margins, penalties and scoring rules. A single entry point runs any of them, and a controller is loaded by alias, by dotted module path or by file path plus class name without modifying the package (Section 6). Reproducing a reported condition therefore requires the track file, the controller identifier and the seed, all of which are recorded per run.

13.3. Seeds and Determinism

Beacon packet randomness is seeded by the run seed, the transmitter, the receiver and the transmission index, so reception is independent of participant count and of the order in which controllers are polled; a fleet run is reproducible for this reason. Inter-vehicle packet loss and procedural course generation draw from seeded generators. For the learning extension, seed bands are registered by role — training, validation, test and final holdout — in a seed registry committed to the repository, and each evaluation records both requested and completed seeds.

13.4. Recorded Artifacts

Each run emits a line-delimited event log and a machine-readable summary, together with the track, controller, seed and interface contract identifiers needed to locate the corresponding result. Detailed provenance and artifact metadata are retained in the repository release rather than repeated in the manuscript.

13.5. Regenerating Tables and Figures

All result tables and figures in this paper are *post-processing* of existing artifacts: no command listed below launches the simulator or modifies a raw artifact.

```
# Result tables and the penalty-identity check
python article/regenerate_tables.py --check

# Track layouts and the controller-comparison plot, from track JSON and artifacts
python article/figures/generate_figures.py

# Journal figures: perception panels from committed artifacts
python article_journal/scripts/make_perception_figure.py

# Manuscript
cd article_journal && latexmk -pdf -interaction=nonstopmode -halt-on-error main.tex
```

Running `regenerate_tables.py` against the recorded result artifacts reproduces the clean-track, current, fleet and coordination values of Tables 11, 12 and 13 exactly, and also verifies the penalty identity $T^{\text{pen}} = T^{\text{official}} + P$ on every finished run. The two journal figure scripts write a provenance JSON alongside each output recording their input files and asserting that no simulator was launched.

13.6. Re-Running Experiments

Reproducing a run rather than a table requires HoloOcean and the recorded seed. Every experimental command is stored verbatim in the per-run metadata; a clean single-vehicle run has the form

```
python -m marine_race_arena.scripts.run_marine_race \
  --track marine_race_arena/tracks/marine_race_horseshoe_bay.json \
  --benchmark-task clean_gate \
  --controller rule_gate_center_then_commit \
  --adapter holoocean --seed 0 \
  --duration 560.0 --current-profile none --official
```

and the current-free demonstration of Section 10.2 is driven by committed scripts that pass `-current-profile` none. The control rate must match the reported campaign: 30 Hz for the systematic benchmark evaluation and 10 Hz for the current-free demonstration and learned-controller validation.

13.7. Updating the Learning Results

The learning section is deliberately structured so that its validation can be refreshed without changing the benchmark protocol. A new frozen checkpoint is added by recording its selection decision, running the same three seeds per official circuit, and regenerating Table 14 from the per-episode result files. The methodology of Section 9—the 27-D observation encoding, recurrent architecture, reward and evaluation separation—is independent of any one checkpoint.

14. Conclusion and Future Work

This paper presented Marine Race Arena, a configurable and reproducible benchmark and evaluation framework for autonomous underwater gate racing. The framework separates participant autonomy from race adjudication through a controller-agnostic participant interface, a configurable race-management layer and an independent referee. Controllers operate within a common participant-level information boundary and provide body-frame commands, while privileged simulator state is reserved for ordered-crossing validation, penalties and official race results.

The framework was validated with two rule-based reference controllers and a recurrent PPO controller, demonstrating that autonomy methods with substantially different internal designs can be integrated through the same benchmark interface. Experiments across heterogeneous circuits, a medium-current condition and multi-vehicle configurations showed that the benchmark can expose differences in passage strategy, disturbance robustness and vehicle interaction. The fleet experiments further demonstrated that the same evaluation framework can represent both independent multi-vehicle operation and distributed cooperative behaviour while preserving separate referee state for each participant.

Future work will extend these capabilities along several complementary directions. Multi-vehicle evaluation can incorporate more advanced cooperation strategies, particularly under environmental disturbances where vehicle interaction, waiting policies and shared progress information become more challenging. Competitive multi-team scenarios and heterogeneous vehicle fleets with different dynamics and sensing suites can also be evaluated under the same race-management and referee framework. Current-aware control and explicit disturbance compensation provide another natural extension, together with a broader evaluation over different current magnitudes, directions and time-varying flow conditions. On the learning side, the recurrent PPO controller serves as a baseline demonstration and can be replaced or complemented by more advanced reinforcement-learning, recurrent or attention-based approaches. Finally, transferring the benchmark interface and race protocol to physical underwater vehicles will enable direct investigation of sim-to-real behaviour and real-world racing experiments.

Data availability

The benchmark implementation, configurations and experimental artifacts are publicly available in the Marine Race Arena repository [26].

References

- [1] J. Betz, H. Zheng, A. Liniger, U. Rosolia, P. Karle, M. Behl, V. Krovi, R. Mangharam, Autonomous vehicles on the edge: A survey on autonomous vehicle racing, *IEEE Open Journal of Intelligent Transportation Systems* 3 (2022) 458–488. doi:10.1109/OJITS.2022.3181510.
- [2] M. O’Kelly, H. Zheng, D. Karthik, R. Mangharam, F1TENTH: An open-source evaluation environment for continuous control and reinforcement learning, in: *Proceedings of the NeurIPS 2019 Competition and Demonstration Track*, Vol. 123 of *Proceedings of Machine Learning Research*, 2020, pp. 77–89. URL <https://proceedings.mlr.press/v123/o-kelly20a.html>

- [3] S. Hoffmann, S. Sagmeister, T. Betz, J. Bongard, S. Büttner, D. Ebner, D. Esser, G. Jank, S. Goblirsch, A. Langmann, M. Leitenstern, L. Ögretmen, P. Pitschi, A.-K. Schwehn, C. Schröder, M. Weinmann, F. Werner, B. Lohmann, J. Betz, M. Lienkamp, Head-to-head autonomous racing at the limits of handling in the A2RL challenge, arXiv preprint arXiv:2602.08571 Abu Dhabi Autonomous Racing League; preprint, not peer reviewed (2026). arXiv:2602.08571.
URL <https://arxiv.org/abs/2602.08571>
- [4] P. Foehn, D. Brescianini, E. Kaufmann, T. Cieslewski, M. Gehrig, M. Muglikar, D. Scaramuzza, AlphaPilot: Autonomous drone racing, in: *Robotics: Science and Systems (RSS)*, 2020. doi:10.15607/RSS.2020.XVI.081.
URL <https://www.roboticsproceedings.org/rss16/p081.html>
- [5] E. Kaufmann, L. Bauersfeld, A. Loquercio, M. Müller, V. Koltun, D. Scaramuzza, Champion-level drone racing using deep reinforcement learning, *Nature* 620 (7976) (2023) 982–987. doi:10.1038/s41586-023-06419-4.
- [6] A. Manzanilla, S. Reyes, M. A. Garcia, D. A. Mercado, R. Lozano, Autonomous navigation for unmanned underwater vehicles: Real-time experiments using computer vision, *IEEE Robotics and Automation Letters* 4 (2) (2019) 1351–1356. doi:10.1109/LRA.2019.2895272.
- [7] M. Stojanovic, J. Preisig, Underwater acoustic communication channels: Propagation models and statistical characterization, *IEEE Communications Magazine* 47 (1) (2009) 84–89. doi:10.1109/MCOM.2009.4752682.
- [8] M. M. M. Manhães, S. A. Scherer, M. Voss, L. R. Douat, T. Rauschenbach, UUV simulator: A Gazebo-based package for underwater intervention and multi-robot simulation, in: *OCEANS 2016 MTS/IEEE Monterey*, 2016, pp. 1–8. doi:10.1109/OCEANS.2016.7761080.
- [9] E. Potokar, S. Ashford, M. Kaess, J. G. Mangelson, HoloOcean: An underwater robotics simulator, in: *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, 2022, pp. 3040–3046. doi:10.1109/ICRA46639.2022.9812353.
URL <https://www.ri.cmu.edu/publications/holocean-an-underwater-robotics-simulator/>
- [10] S. Chu, Z. Huang, Y. Li, M. Lin, D. Li, I. Carlucho, Y. R. Petillot, C. Yang, MarineGym: A high-performance reinforcement learning platform for underwater robotics, in: *2025 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2025, pp. 17146–17153. doi:10.1109/IROS60139.2025.11246146.
URL <https://ieeexplore.ieee.org/document/11246146>
- [11] M. von Benzon, F. F. Sørensen, E. Uth, J. Jouffroy, J. Liniger, S. Pedersen, An open-source benchmark simulator: Control of a BlueROV2 underwater robot, *Journal of Marine Science and Engineering* 10 (12) (2022) 1898. doi:10.3390/jmse10121898.
- [12] P. Cieślak, Stonefish: An advanced open-source simulation tool designed for marine robotics, with a ROS interface, in: *OCEANS 2019 - Marseille*, 2019, pp. 1–6. doi:10.1109/OCEANSE.2019.8867434.
- [13] M. M. Zhang, W.-S. Choi, J. Herman, D. Davis, C. Vogt, M. McCarrin, Y. Vijay, D. Dutia, W. Lew, S. Peters, B. Bingham, DAVE aquatic virtual environment: Toward a general underwater robotics simulator, in: *Proceedings of the IEEE/OES Autonomous Underwater Vehicles Symposium (AUV)*, 2022, pp. 1–8. doi:10.1109/AUV53081.2022.9965808.
URL <https://arxiv.org/abs/2209.02862>
- [14] E. Potokar, S. Ashford, M. Kaess, J. G. Mangelson, HoloOcean: A full-featured marine robotics simulator for perception and autonomy, *IEEE Journal of Oceanic Engineering* 49 (4) (2024) 1322–1336. doi:10.1109/JOE.2024.3410290.
- [15] A. Amer, O. Álvarez-Tuñón, H. I. Ugurlu, J. L. F. Sejersen, Y. Brodskiy, E. Kayacan, UNav-Sim: A visually realistic underwater robotics simulator and synthetic data-generation framework, in: *2023 21st International Conference on Advanced Robotics (ICAR)*, IEEE, 2023, pp. 570–576. doi:10.1109/ICAR58858.2023.10406819.

- [16] Z. Huang, M. Buchholz, M. Grimaldi, H. Yu, I. Carlucho, Y. R. Petillot, URoBench: Comparative analyses of underwater robotics simulators from reinforcement learning perspective, in: OCEANS 2024 - Singapore, IEEE, 2024. doi:10.1109/OCEANS51537.2024.10682284.
- [17] A. Dosovitskiy, G. Ros, F. Codevilla, A. Lopez, V. Koltun, CARLA: An open urban driving simulator, in: Proceedings of the Conference on Robot Learning (CoRL), Vol. 78 of Proceedings of Machine Learning Research, 2017, pp. 1–16.
URL <https://proceedings.mlr.press/v78/dosovitskiy17a.html>
- [18] CARLA Autonomous Driving Leaderboard, CARLA Autonomous Driving Leaderboard: Evaluation criteria 2.1, accessed: 2026-09-10 (2026).
URL https://leaderboard.carla.org/evaluation_v2_1/
- [19] M. Franchi, et al., Underwater robotics competitions: The european robotics league emergency robots experience with feelhippo auv, *Frontiers in Robotics and AI* 7 (2020) 3. doi:10.3389/frobt.2020.00003.
- [20] K. Harada, R. Fukuda, Y. Mizoguchi, Y. Yamamoto, K. Mishima, Y. Tanaka, Y. Nishida, K. Ishii, Autonomous underwater vehicle with vision-based navigation system for underwater robot competition, in: Proceedings of the International Conference on Artificial Life and Robotics (ICAROB2022), Vol. 27, 2022, pp. 354–359. doi:10.5954/ICAROB.2022.OS28-2.
URL https://kyutech.repo.nii.ac.jp/record/7489/files/ICAROB2022_OS28-2.pdf
- [21] M. Xanthidis, M. Kalaitzakis, N. Karapetyan, J. Johnson, N. Vitzilaos, J. M. O’Kane, I. Rekleitis, AquaVis: A perception-aware autonomous navigation framework for underwater vehicles, in: 2021 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS), 2021. doi:10.1109/IROS51168.2021.9636124.
- [22] Y. Yang, Y. Xiao, T. Li, A survey of autonomous underwater vehicle formation: Performance, formation control, and communication capability, *IEEE Communications Surveys & Tutorials* 23 (2) (2021) 815–841. doi:10.1109/COMST.2021.3059998.
- [23] T. Yan, Z. Xu, S. A. Gadsden, Formation control of multiple autonomous underwater vehicles: A review, *Intelligence & Robotics* 3 (1) (2023) 1–22. doi:10.20517/ir.2023.01.
- [24] G. Brockman, V. Cheung, L. Pettersson, J. Schneider, J. Schulman, J. Tang, W. Zaremba, OpenAI gym, arXiv preprint arXiv:1606.01540 (2016). arXiv:1606.01540.
URL <https://arxiv.org/abs/1606.01540>
- [25] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, O. Klimov, Proximal policy optimization algorithms, arXiv preprint arXiv:1707.06347 (2017). arXiv:1707.06347.
URL <https://arxiv.org/abs/1707.06347>
- [26] A. Bedei, HoloDroneCompetition: Marine Race Arena repository, <https://github.com/AndreaBedei1/HoloDroneCompetition> (2026).